

William Stallings

Computer Organization

and Architecture

Chapter 1

Introduction

Architecture & Organization 1

⌘ Architecture is those attributes visible to the programmer

☑ Instruction set, number of bits used for data representation, I/O mechanisms, addressing techniques.

☑ e.g. Is there a multiply instruction?

⌘ Organization is how features are implemented

☑ Control signals, interfaces, memory technology.

☑ e.g. Is there a hardware multiply unit or is it done by repeated addition?

Architecture & Organization 2

- ⌘ All Intel x86 family share the same basic architecture
- ⌘ The IBM System/370 family share the same basic architecture
- ⌘ This gives code compatibility
 - ☑ At least backwards
- ⌘ Organization differs between different versions

Structure & Function

- ⌘ Structure is the way in which components relate to each other
- ⌘ Function is the operation of individual components as part of the structure

Function

⌘ All computer functions are:

- ☑ Data processing

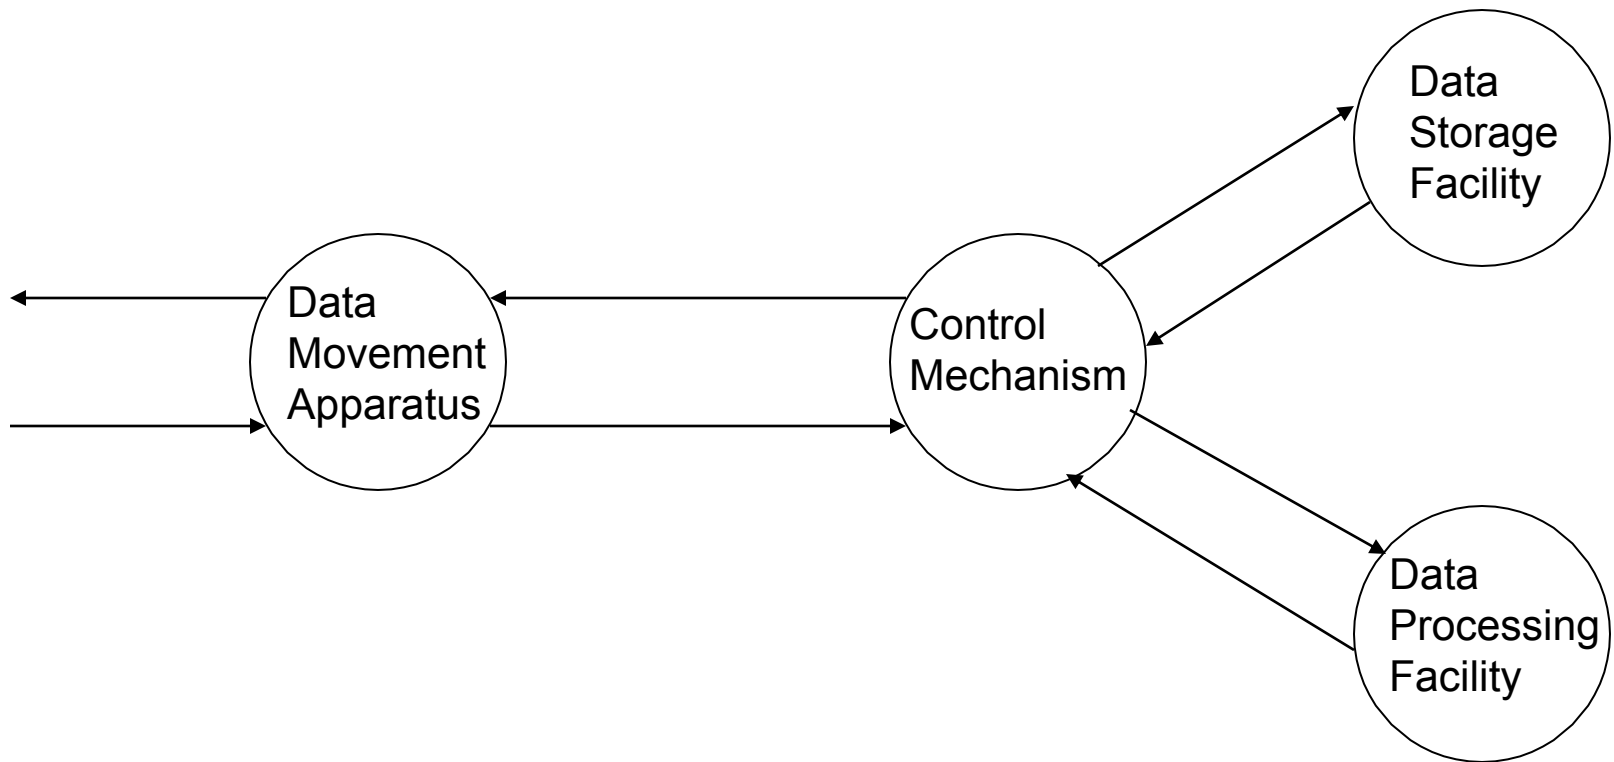
- ☑ Data storage

- ☑ Data movement

- ☑ Control

Functional view

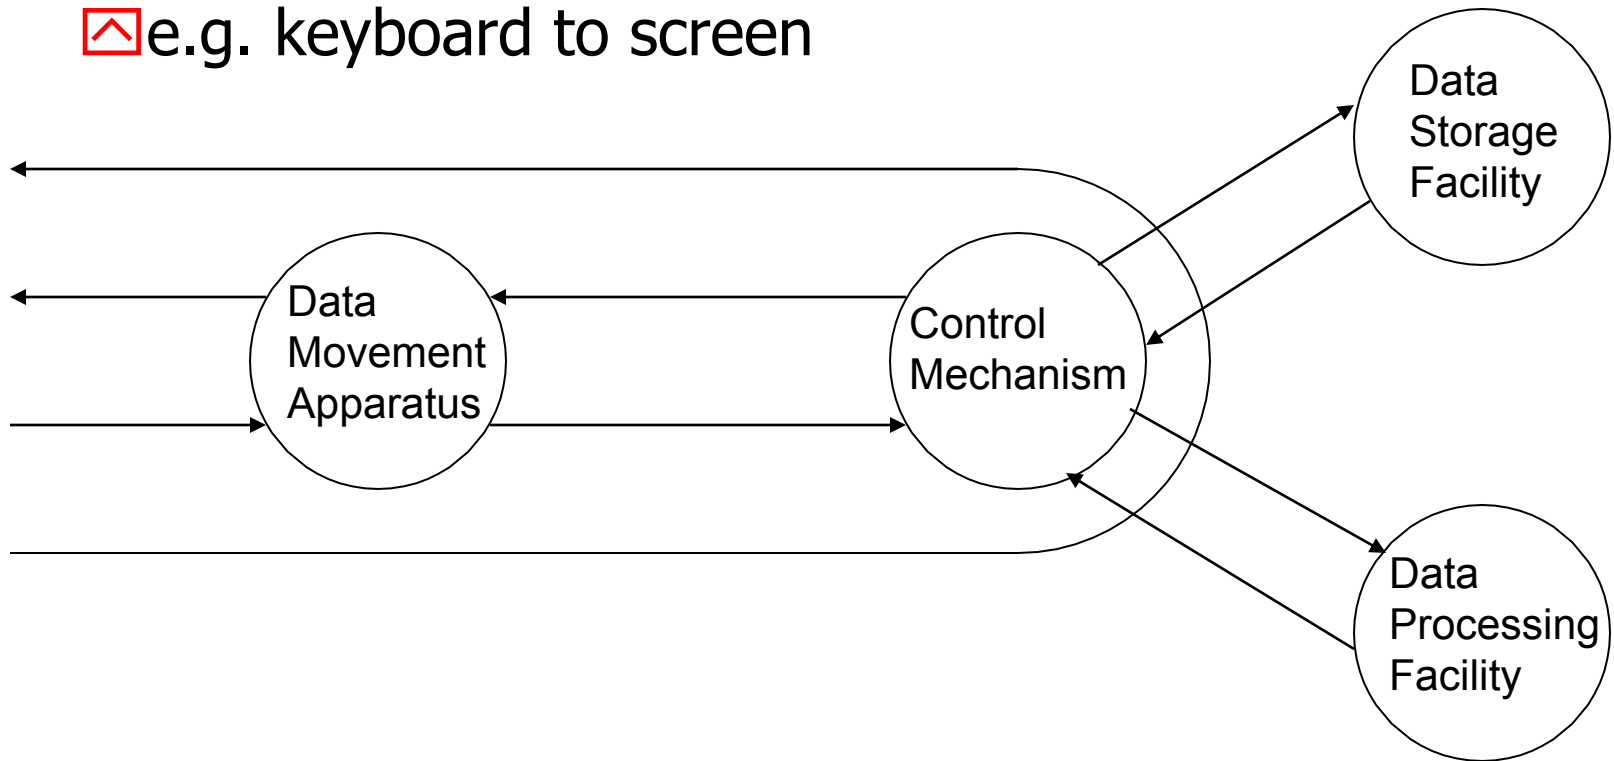
⌘ Functional view of a computer



Operations (1)

⌘ Data movement

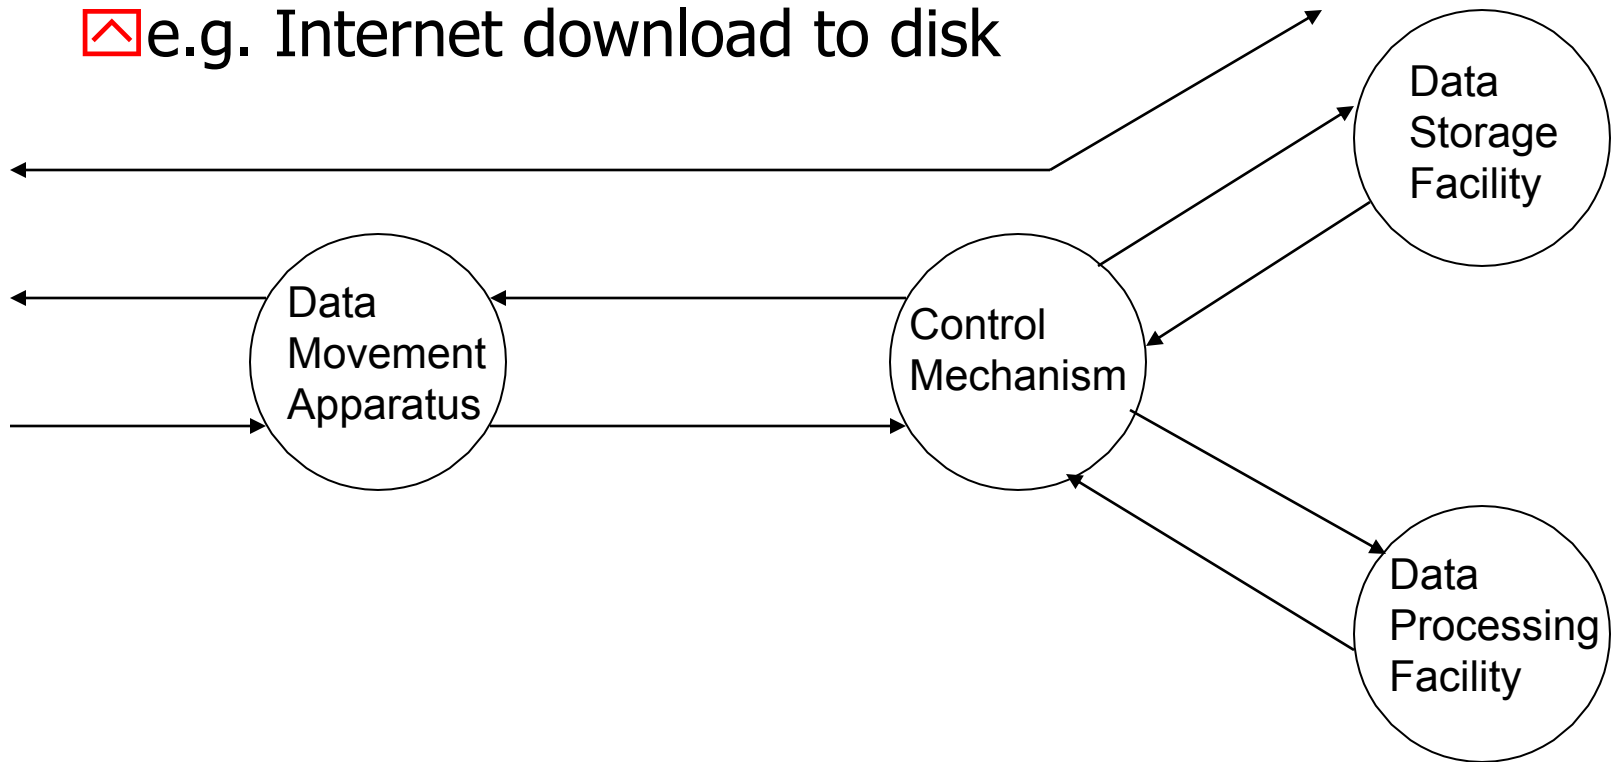
☐ e.g. keyboard to screen



Operations (2)

⌘ Storage

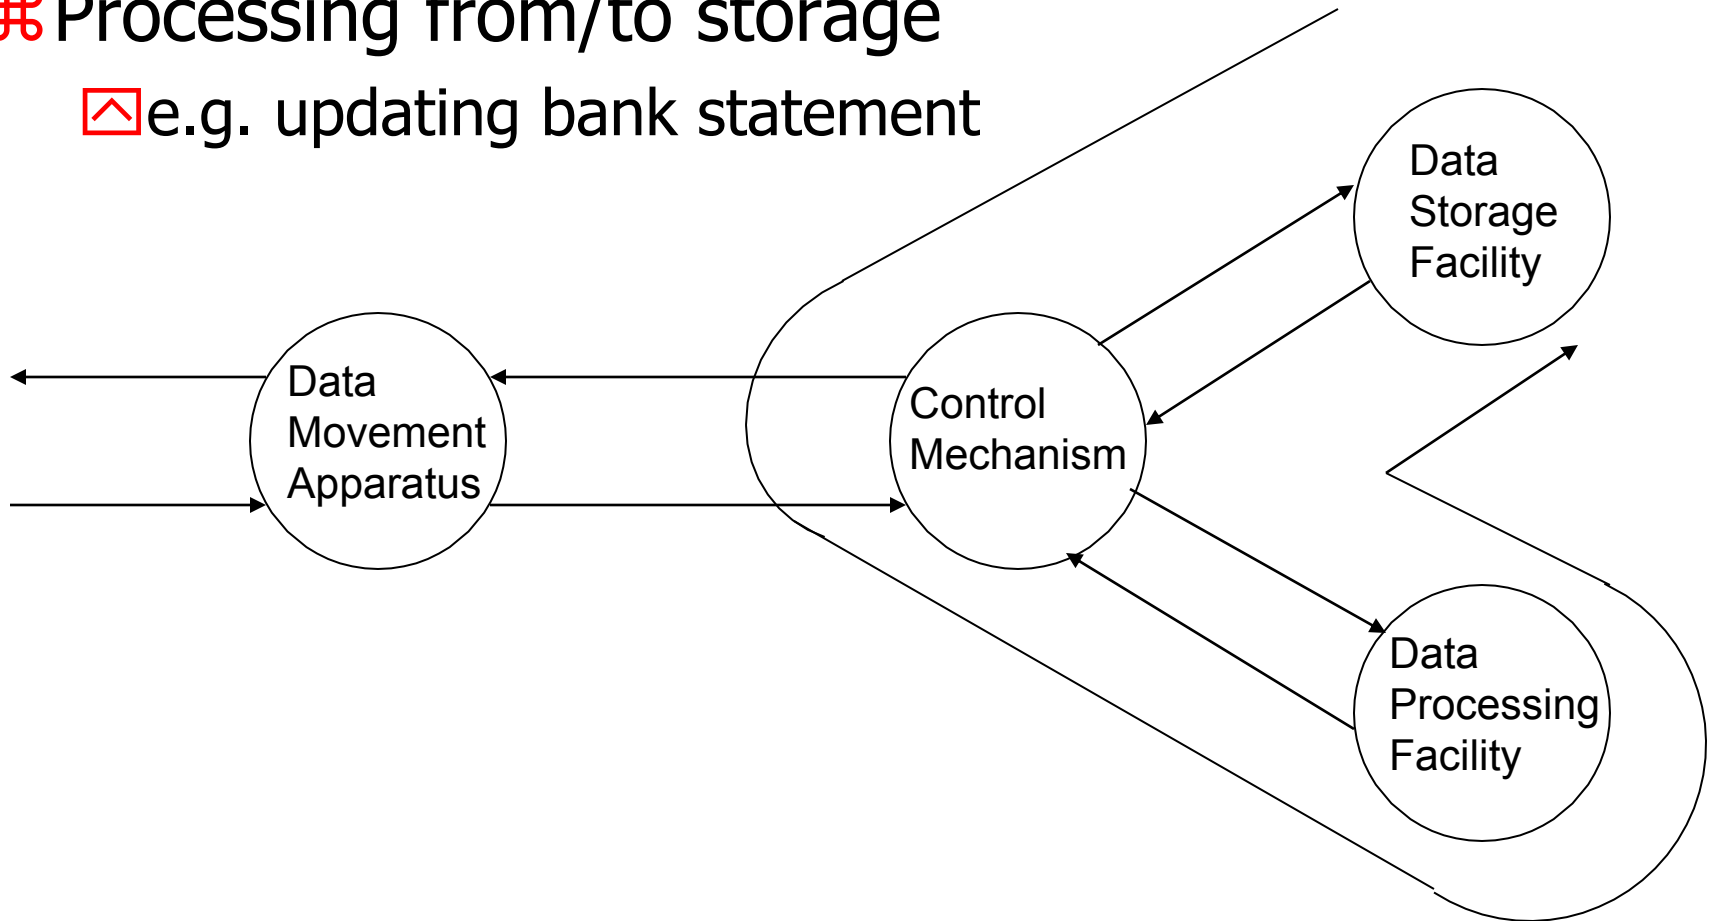
☐ e.g. Internet download to disk



Operation (3)

⌘ Processing from/to storage

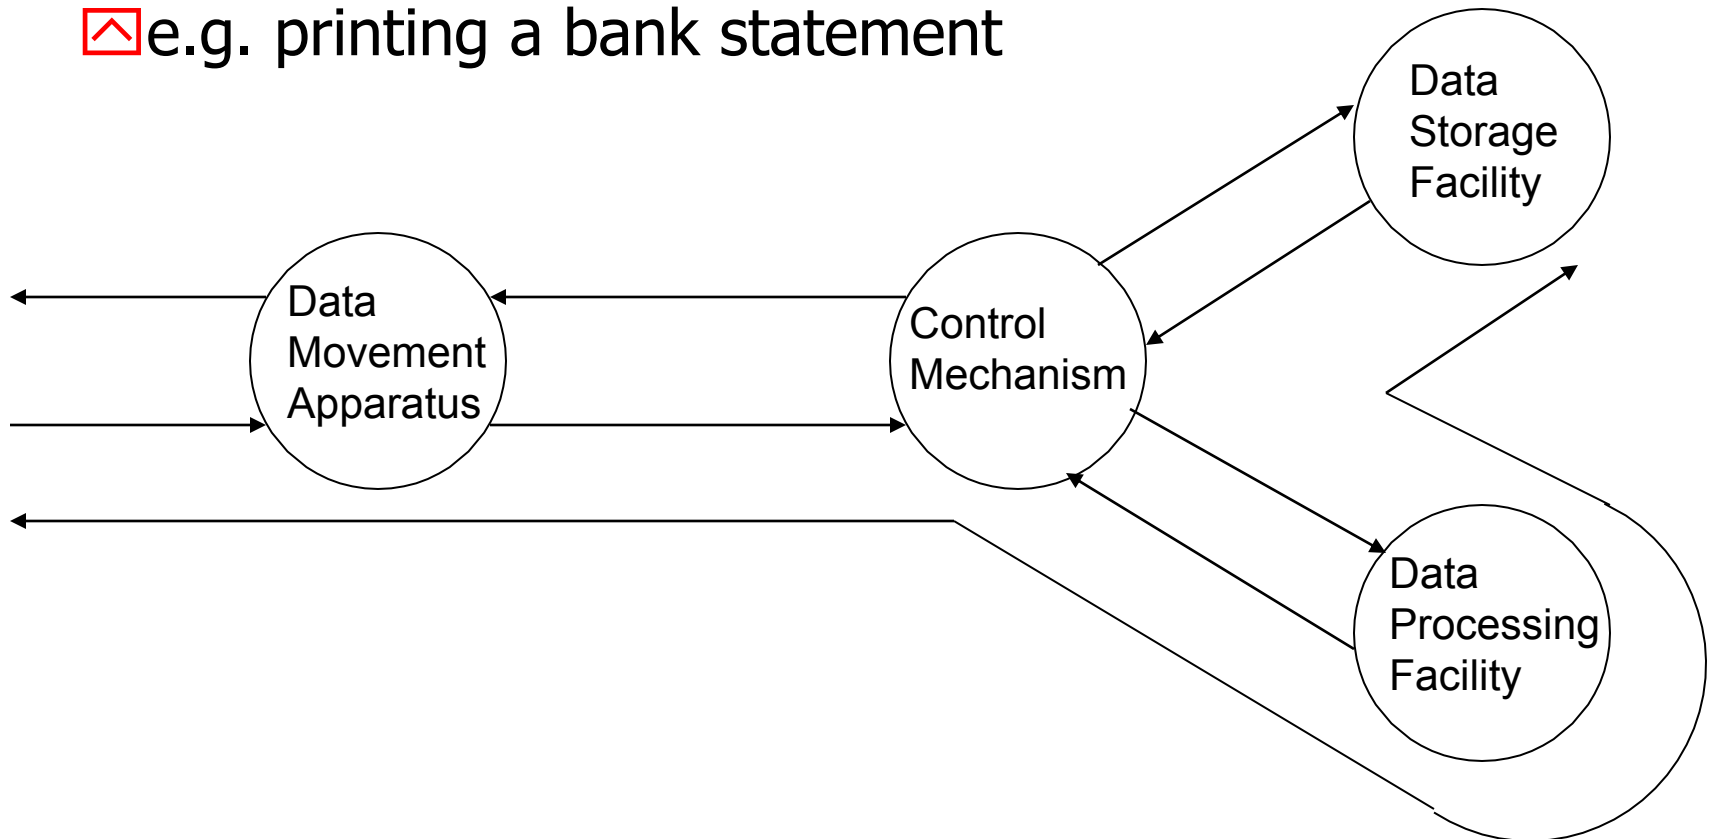
☑ e.g. updating bank statement



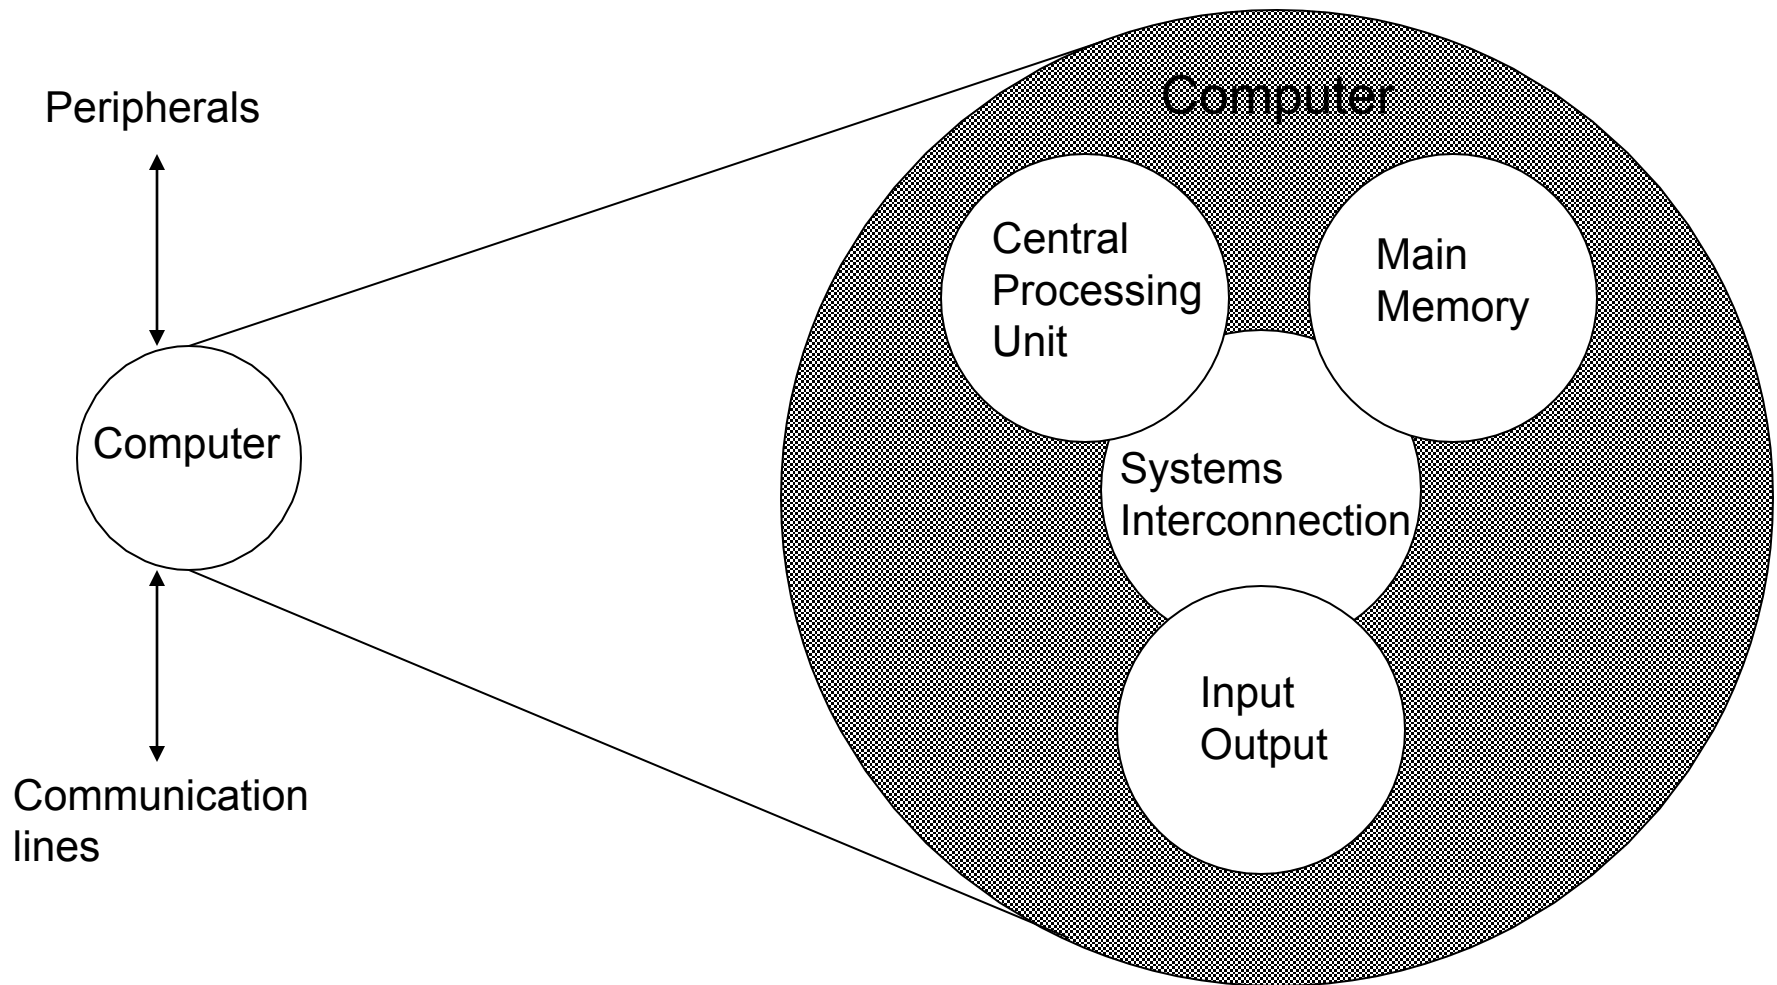
Operation (4)

⌘ Processing from storage to I/O

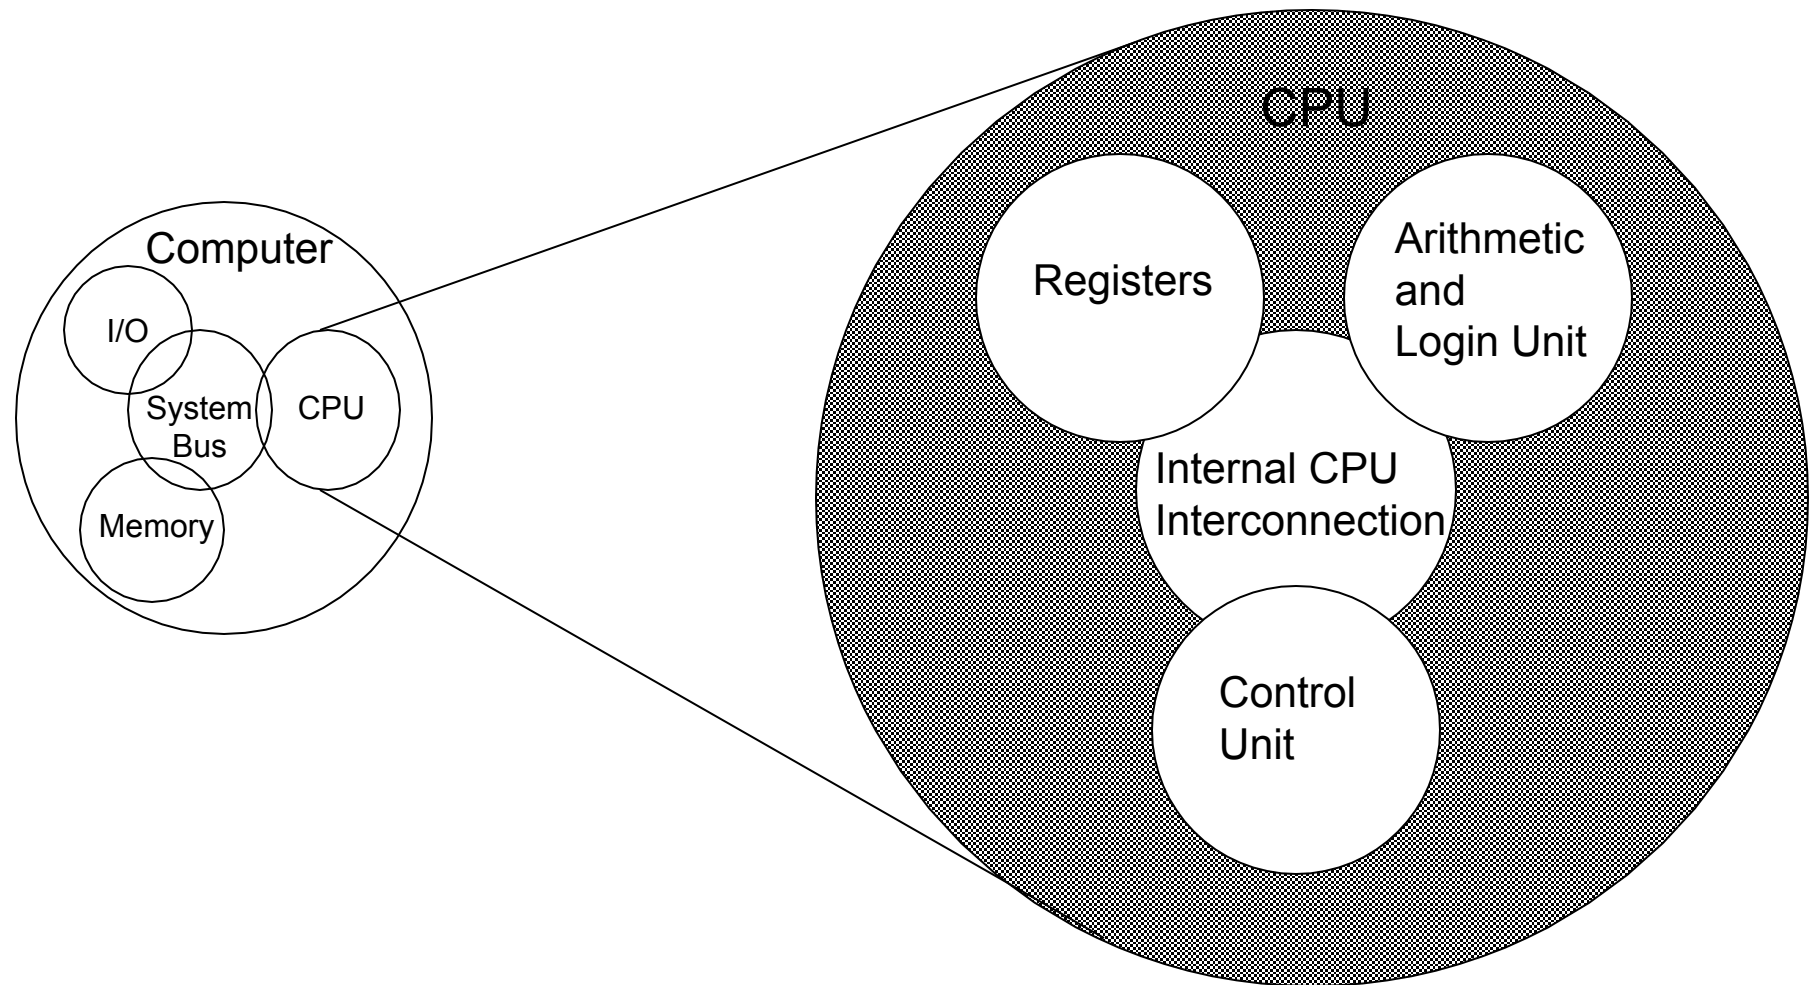
☐ e.g. printing a bank statement



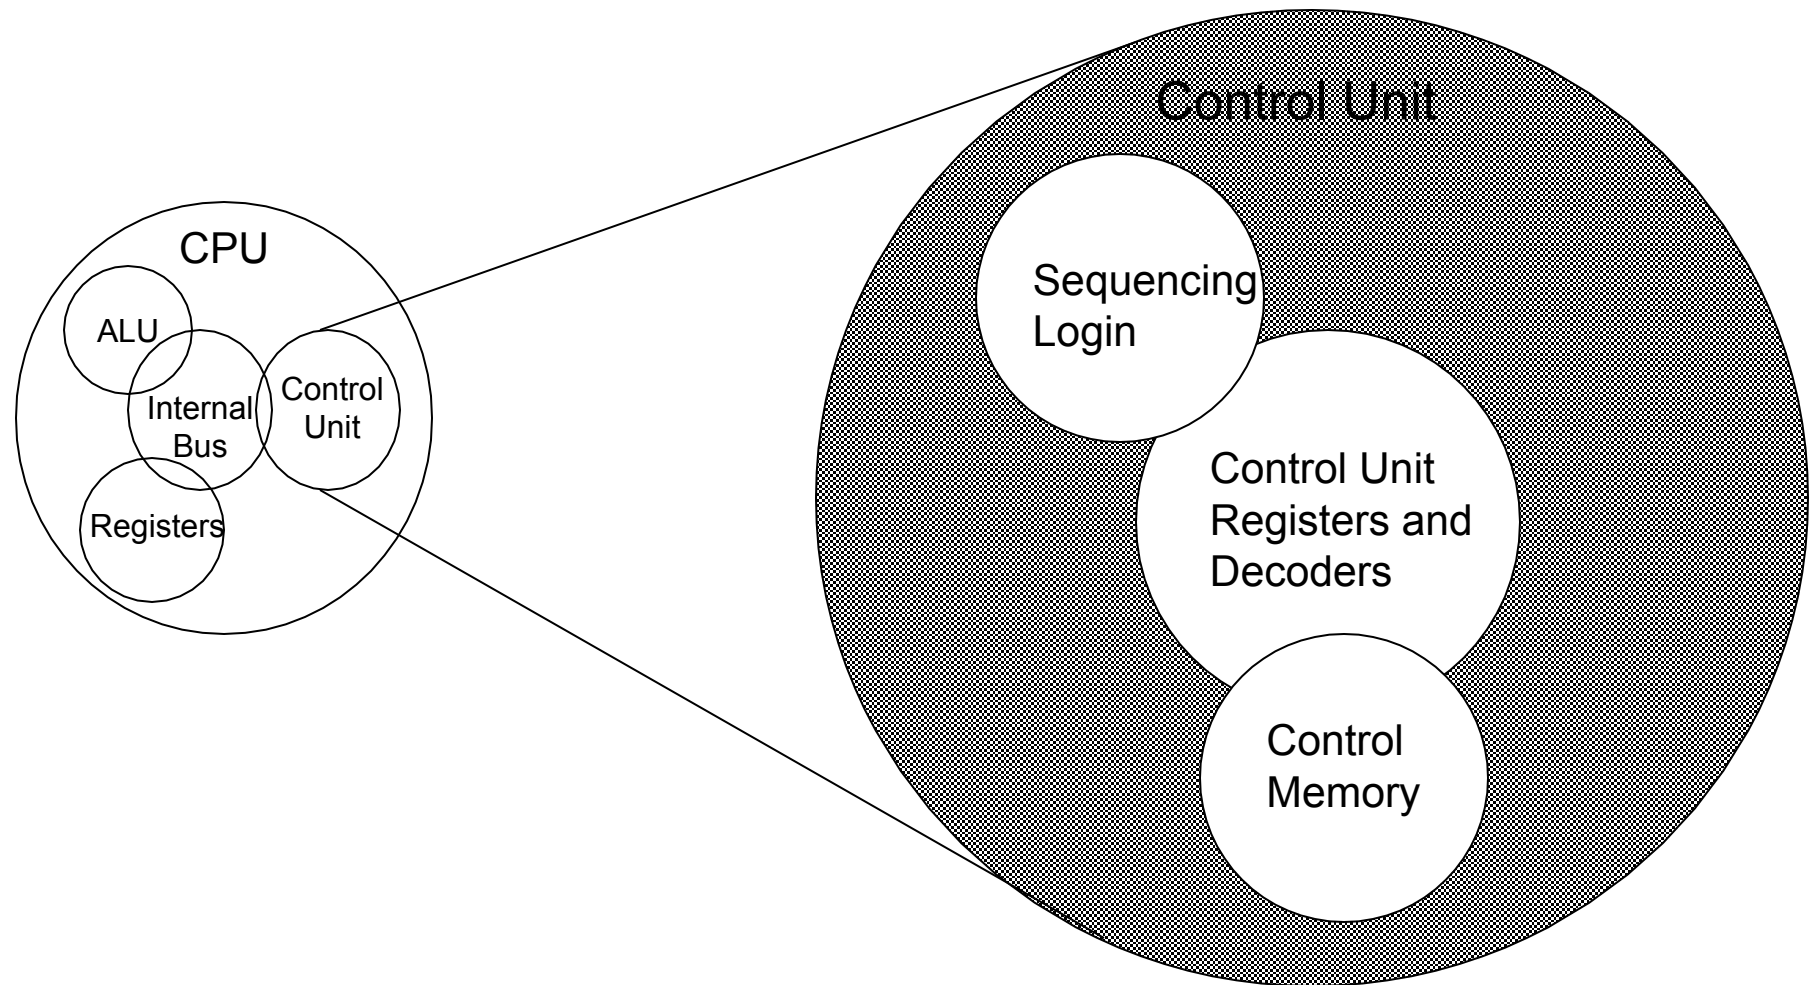
Structure - Top Level



Structure - The CPU



Structure - The Control Unit



Outline of the Book (1)

- ⌘ Computer Evolution and Performance
- ⌘ Computer Interconnection Structures
- ⌘ Internal Memory
- ⌘ External Memory
- ⌘ Input/Output
- ⌘ Operating Systems Support
- ⌘ Computer Arithmetic
- ⌘ Instruction Sets

Outline of the Book (2)

- ⌘ CPU Structure and Function
- ⌘ Reduced Instruction Set Computers
- ⌘ Superscalar Processors
- ⌘ Control Unit Operation
- ⌘ Microprogrammed Control
- ⌘ Multiprocessors and Vector Processing
- ⌘ Digital Logic (Appendix)

Internet Resources

- Web site for book

⌘ <http://www.shore.net/~ws/COA5e.html>

- ☐ links to sites of interest

- ☐ links to sites for courses that use the book

- ☐ errata list for book

- ☐ information on other books by W. Stallings

Internet Resources

- Web sites to look for

- ⌘ WWW Computer Architecture Home Page
- ⌘ CPU Info Center
- ⌘ ACM Special Interest Group on Computer Architecture
- ⌘ IEEE Technical Committee on Computer Architecture
- ⌘ Intel Technology Journal
- ⌘ Manufacturer's sites
 - ☒ Intel, IBM, etc.

Internet Resources

- Usenet News Groups

⌘ comp.arch

⌘ comp.arch.arithmetic

⌘ comp.arch.storage

William Stallings

Computer Organization and Architecture

Chapter 2

Computer Evolution and Performance

ENIAC - background

- ⌘ Electronic Numerical Integrator And Computer
- ⌘ Eckert and Mauchly
- ⌘ University of Pennsylvania
- ⌘ Trajectory tables for weapons
- ⌘ Started 1943
- ⌘ Finished 1946
 - ☒ Too late for war effort
- ⌘ Used until 1955

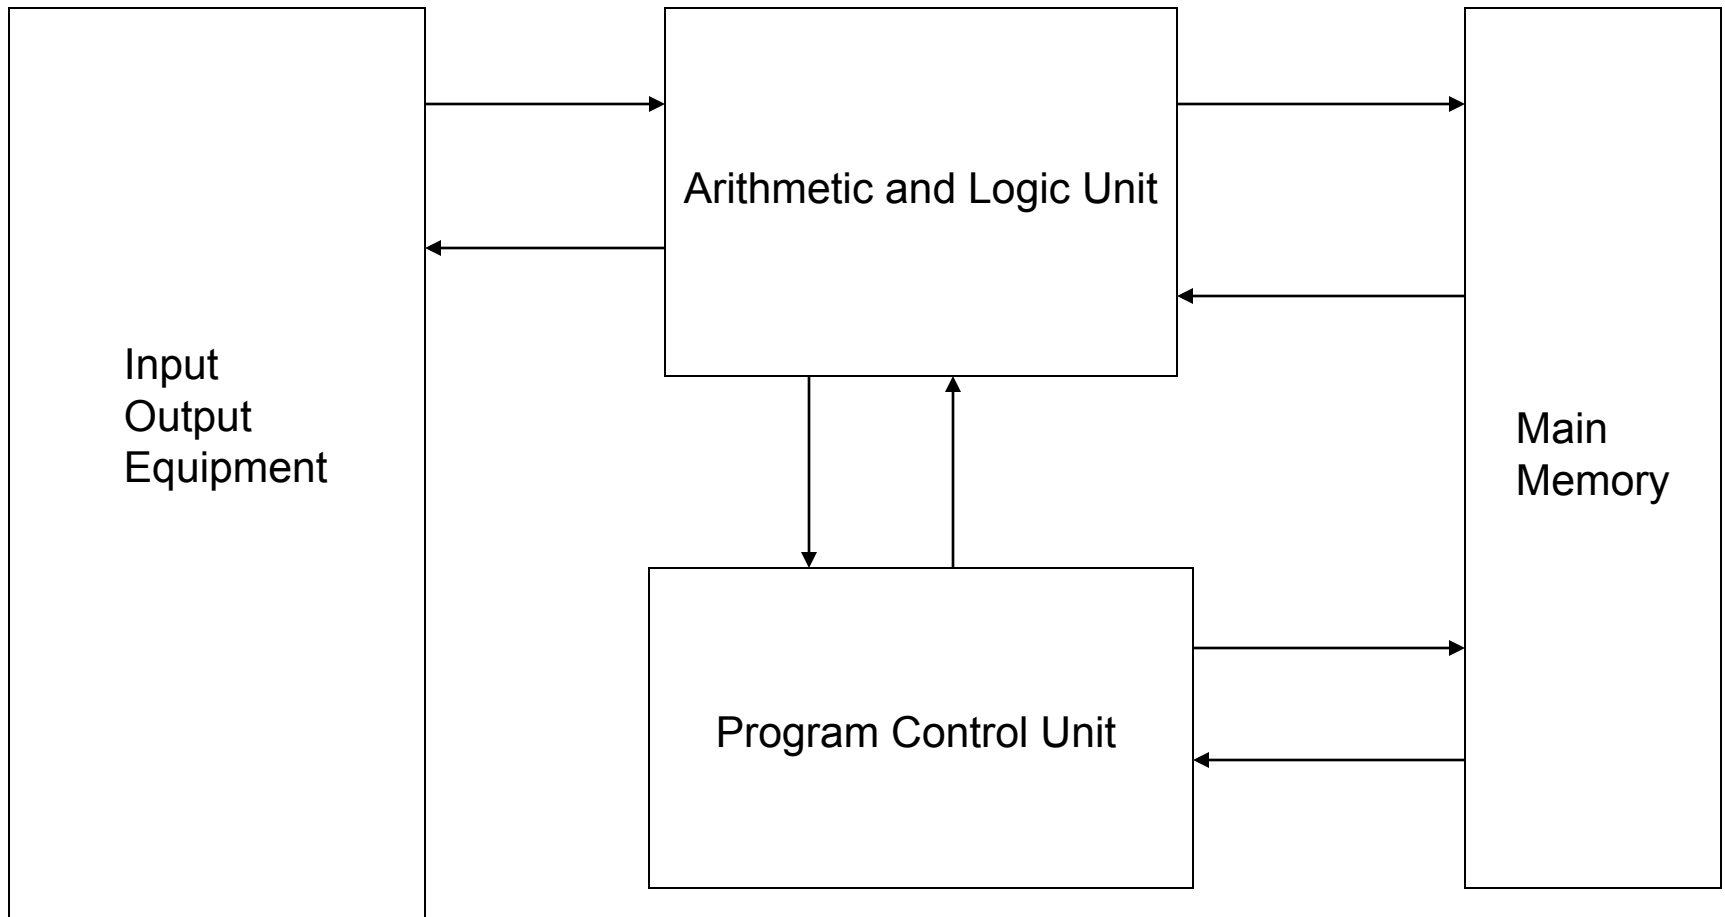
ENIAC - details

- ⌘ Decimal (not binary)
- ⌘ 20 accumulators of 10 digits
- ⌘ Programmed manually by switches
- ⌘ 18,000 vacuum tubes
- ⌘ 30 tons
- ⌘ 15,000 square feet
- ⌘ 140 kW power consumption
- ⌘ 5,000 additions per second

von Neumann/Turing

- ⌘ Stored Program concept
- ⌘ Main memory storing programs and data
- ⌘ ALU operating on binary data
- ⌘ Control unit interpreting instructions from memory and executing
- ⌘ Input and output equipment operated by control unit
- ⌘ Princeton Institute for Advanced Studies
 - ☒ IAS
- ⌘ Completed 1952

Structure of von Nuemann machine



IAS - details

⌘ 1000 x 40 bit words

- ☒ Binary number

- ☒ 2 x 20 bit instructions

⌘ Set of registers (storage in CPU)

- ☒ Memory Buffer Register

- ☒ Memory Address Register

- ☒ Instruction Register

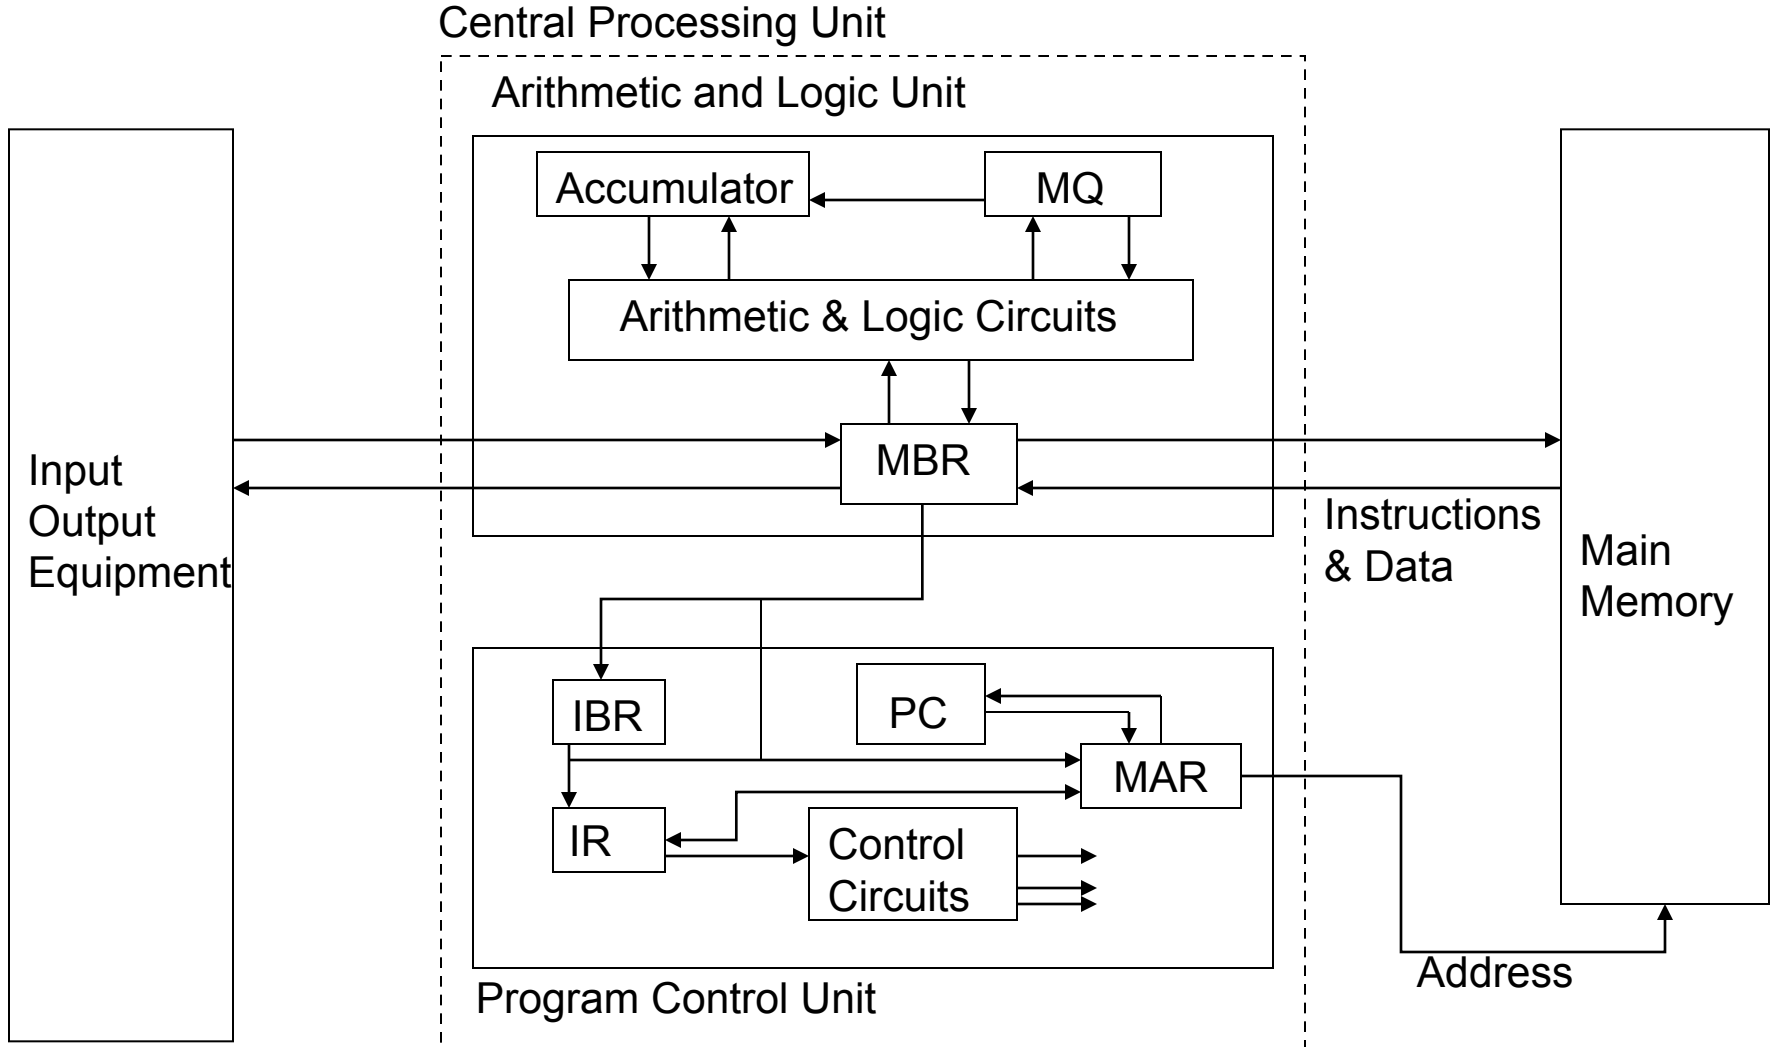
- ☒ Instruction Buffer Register

- ☒ Program Counter

- ☒ Accumulator

- ☒ Multiplier Quotient

Structure of IAS - detail



Commercial Computers

- ⌘ 1947 - Eckert-Mauchly Computer Corporation
- ⌘ UNIVAC I (Universal Automatic Computer)
- ⌘ US Bureau of Census 1950 calculations
- ⌘ Became part of Sperry-Rand Corporation
- ⌘ Late 1950s - UNIVAC II
 - ☒ Faster
 - ☒ More memory

IBM

- ⌘ Punched-card processing equipment

- ⌘ 1953 - the 701

 - ☑ IBM's first stored program computer

 - ☑ Scientific calculations

- ⌘ 1955 - the 702

 - ☑ Business applications

- ⌘ Lead to 700/7000 series

Transistors

- ⌘ Replaced vacuum tubes
- ⌘ Smaller
- ⌘ Cheaper
- ⌘ Less heat dissipation
- ⌘ Solid State device
- ⌘ Made from Silicon (Sand)
- ⌘ Invented 1947 at Bell Labs
- ⌘ William Shockley et al.

Transistor Based Computers

- ⌘ Second generation machines
- ⌘ NCR & RCA produced small transistor machines
- ⌘ IBM 7000
- ⌘ DEC - 1957
 - ☑ Produced PDP-1

Microelectronics

- ⌘ Literally - “small electronics”
- ⌘ A computer is made up of gates, memory cells and interconnections
- ⌘ These can be manufactured on a semiconductor
- ⌘ e.g. silicon wafer


Generations of Computer

- ⌘ Vacuum tube - 1946-1957
- ⌘ Transistor - 1958-1964
- ⌘ Small scale integration - 1965 on
 - ⌘ Up to 100 devices on a chip
- ⌘ Medium scale integration - to 1971
 - ⌘ 100-3,000 devices on a chip
- ⌘ Large scale integration - 1971-1977
 - ⌘ 3,000 - 100,000 devices on a chip
- ⌘ Very large scale integration - 1978 to date
 - ⌘ 100,000 - 100,000,000 devices on a chip
- ⌘ Ultra large scale integration
 - ⌘ Over 100,000,000 devices on a chip

Moore's Law

- ⌘ Increased density of components on chip
- ⌘ Gordon Moore - cofounder of Intel
- ⌘ Number of transistors on a chip will double every year
- ⌘ Since 1970's development has slowed a little
 - ☒ Number of transistors doubles every 18 months
- ⌘ Cost of a chip has remained almost unchanged
- ⌘ Higher packing density means shorter electrical paths, giving higher performance
- ⌘ Smaller size gives increased flexibility
- ⌘ Reduced power and cooling requirements
- ⌘ Fewer interconnections increases reliability

Growth in CPU Transistor Count



IBM 360 series

⌘ 1964

⌘ Replaced (& not compatible with) 7000 series

⌘ First planned “family” of computers

- ☒ Similar or identical instruction sets

- ☒ Similar or identical O/S

- ☒ Increasing speed

- ☒ Increasing number of I/O ports (i.e. more terminals)

- ☒ Increased memory size

- ☒ Increased cost

⌘ Multiplexed switch structure

DEC PDP-8

⌘ 1964

⌘ First minicomputer (after miniskirt!)

⌘ Did not need air conditioned room

⌘ Small enough to sit on a lab bench


⌘ \$16,000

☐ \$100k+ for IBM 360

⌘ Embedded applications & OEM

⌘ BUS STRUCTURE

DEC - PDP-8 Bus Structure



Semiconductor Memory

⌘ 1970

⌘ Fairchild

⌘ Size of a single core

☐ i.e. 1 bit of magnetic core storage

⌘ Holds 256 bits

⌘ Non-destructive read

⌘ Much faster than core

⌘ Capacity approximately doubles each year

Intel

⌘ 1971 - 4004

- ☑ First microprocessor
- ☑ All CPU components on a single chip
- ☑ 4 bit

⌘ Followed in 1972 by 8008

- ☑ 8 bit
- ☑ Both designed for specific applications

⌘ 1974 - 8080

- ☑ Intel's first general purpose microprocessor


Speeding it up

- ⌘ Pipelining
- ⌘ On board cache
- ⌘ On board L1 & L2 cache
- ⌘ Branch prediction
- ⌘ Data flow analysis
- ⌘ Speculative execution


Performance Mismatch

- ⌘ Processor speed increased
- ⌘ Memory capacity increased
- ⌘ Memory speed lags behind processor speed

DRAM and Processor Characteristics



Trends in DRAM use



Solutions

⌘ Increase number of bits retrieved at one time

☑ Make DRAM “wider” rather than “deeper”

⌘ Change DRAM interface

☑ Cache

⌘ Reduce frequency of memory access

☑ More complex cache and cache on chip

⌘ Increase interconnection bandwidth

☑ High speed buses

☑ Hierarchy of buses

Internet Resources

⌘ <http://www.intel.com/>

☐ Search for the Intel Museum

⌘ <http://www.ibm.com>

⌘ <http://www.dec.com>

⌘ Charles Babbage Institute

⌘ PowerPC

⌘ Intel Developer Home

William Stallings

Computer Organization

and Architecture

Chapter 3

**A top level view of computer
function & Interconnection**

Program Concept

- ❑ Hardwired systems are inflexible
- ❑ General purpose hardware can do different tasks, given correct control signals
- ❑ Instead of re-wiring, supply a new set of control signals

What is a program?

- A sequence of steps
- For each step, an arithmetic or logical operation is done
- For each operation, a different set of control signals is needed

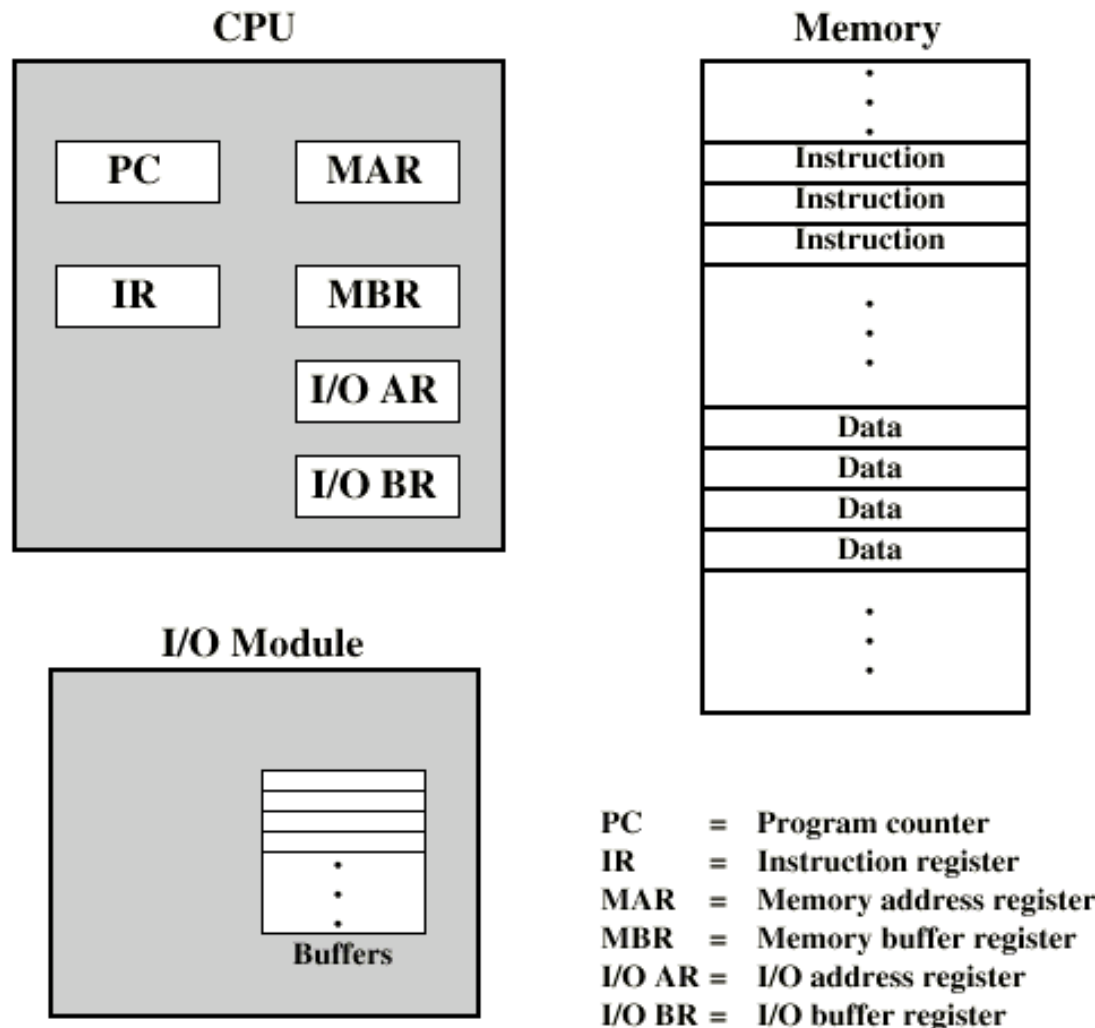
Function of Control Unit

- For each operation a unique code is provided
 - e.g. ADD, MOVE
- A hardware segment accepts the code and issues the control signals
- We have a computer!

Components


- ❑ The Control Unit and the Arithmetic and Logic Unit constitute the Central Processing Unit
- ❑ Data and instructions need to get into the system and results out
 - ❑ Input/output
- ❑ Temporary storage of code and results is needed
 - ❑ Main memory

Computer Components: Top Level View



Instruction Cycle

- Two steps:
 - Fetch
 - Execute



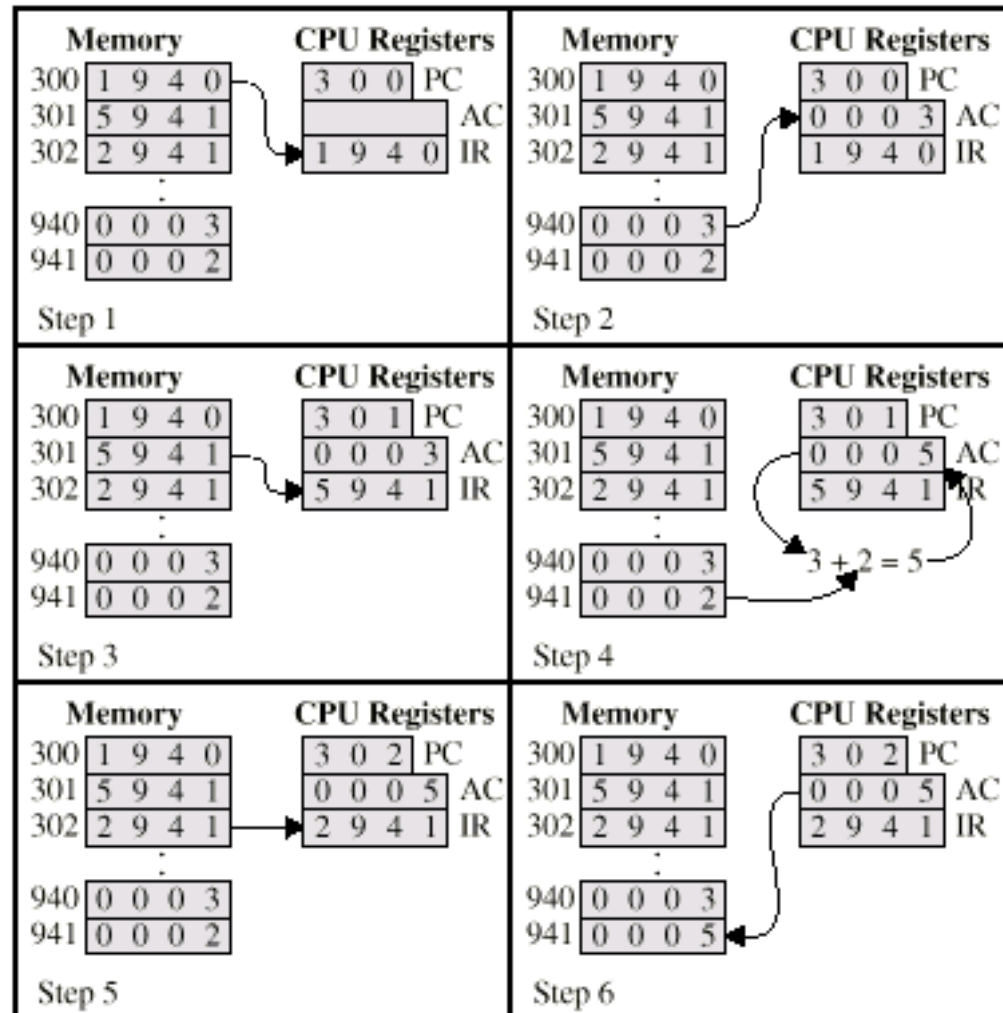
Fetch Cycle

- ❑ Program Counter (PC) holds address of next instruction to fetch
- ❑ Processor fetches instruction from memory location pointed to by PC
- ❑ Increment PC
 - ❑ Unless told otherwise
- ❑ Instruction loaded into Instruction Register (IR)
- ❑ Processor interprets instruction and performs required actions

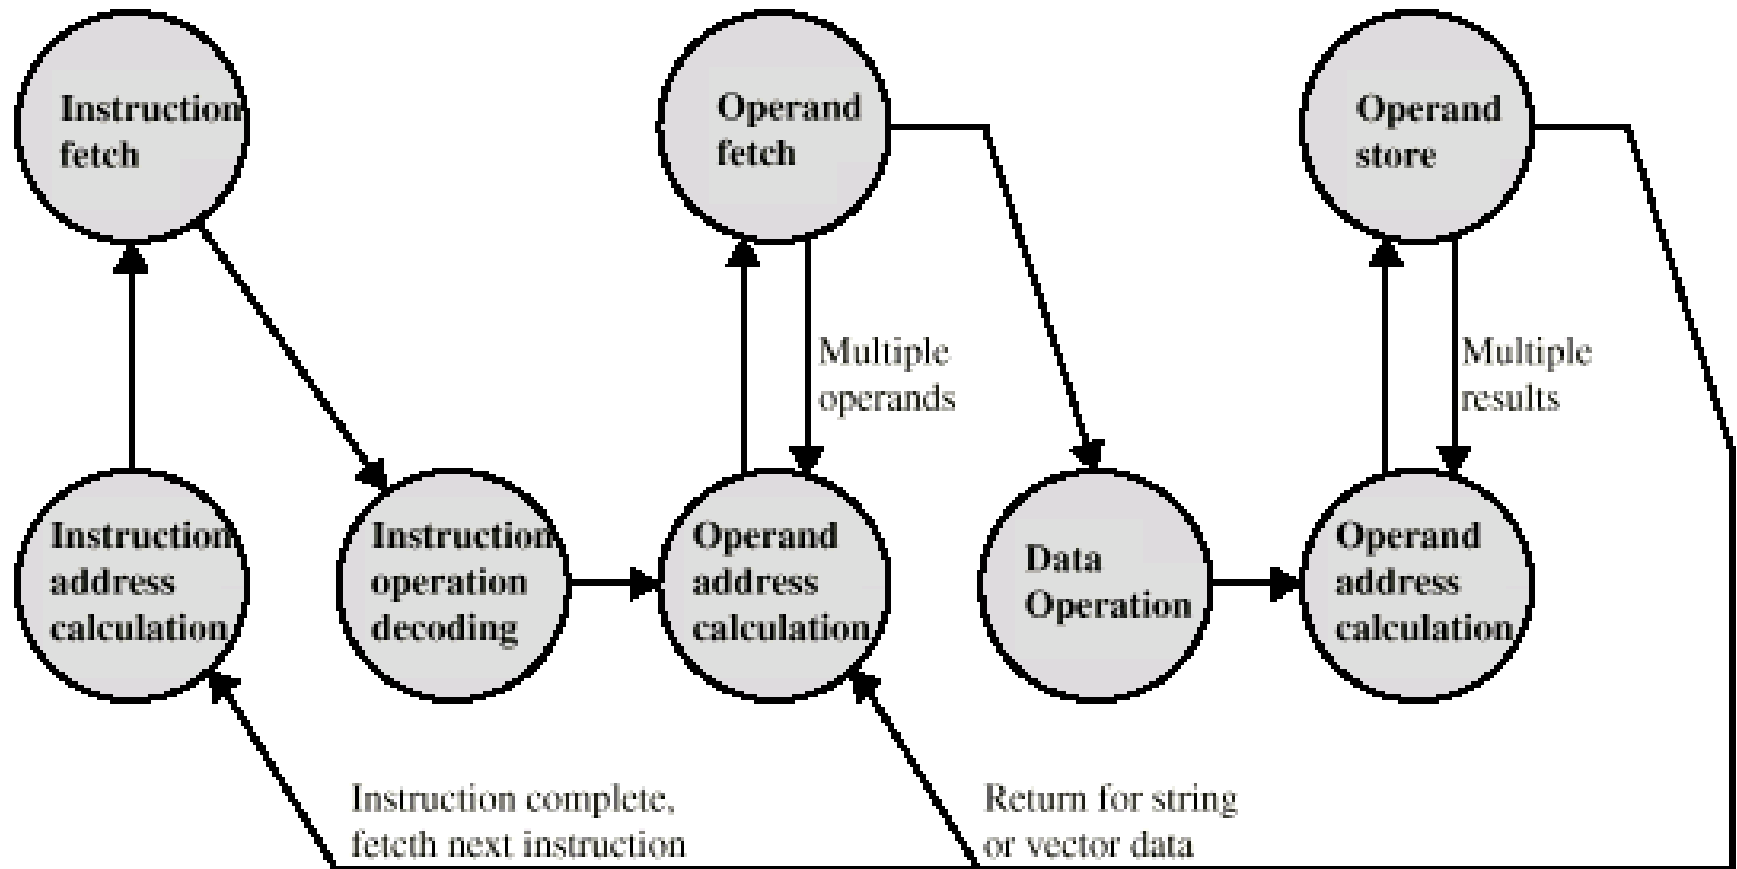
Execute Cycle

- ❑ Processor-memory
 - ❑ data transfer between CPU and main memory
- ❑ Processor I/O
 - ❑ Data transfer between CPU and I/O module
- ❑ Data processing
 - ❑ Some arithmetic or logical operation on data
- ❑ Control
 - ❑ Alteration of sequence of operations
 - ❑ e.g. jump
- ❑ Combination of above

Example of Program Execution




Instruction Cycle - State Diagram



Interrupts

- ❑ Mechanism by which other modules (e.g. I/O) may interrupt normal sequence of processing
- ❑ Program
 - ❑ e.g. overflow, division by zero
- ❑ Timer
 - ❑ Generated by internal processor timer
 - ❑ Used in pre-emptive multi-tasking
- ❑ I/O
 - ❑ from I/O controller
- ❑ Hardware failure
 - ❑ e.g. memory parity error

Program Flow Control



(a) No interrupts


(b) Interrupts; short I/O wait

(c) Interrupts; long I/O wait

Interrupt Cycle

- ❑ Added to instruction cycle
- ❑ Processor checks for interrupt
 - ❑ Indicated by an interrupt signal
- ❑ If no interrupt, fetch next instruction
- ❑ If interrupt pending:
 - ❑ Suspend execution of current program
 - ❑ Save context
 - ❑ Set PC to start address of interrupt handler routine
 - ❑ Process interrupt
 - ❑ Restore context and continue interrupted program

Instruction Cycle (with Interrupts) - State Diagram



Multiple Interrupts

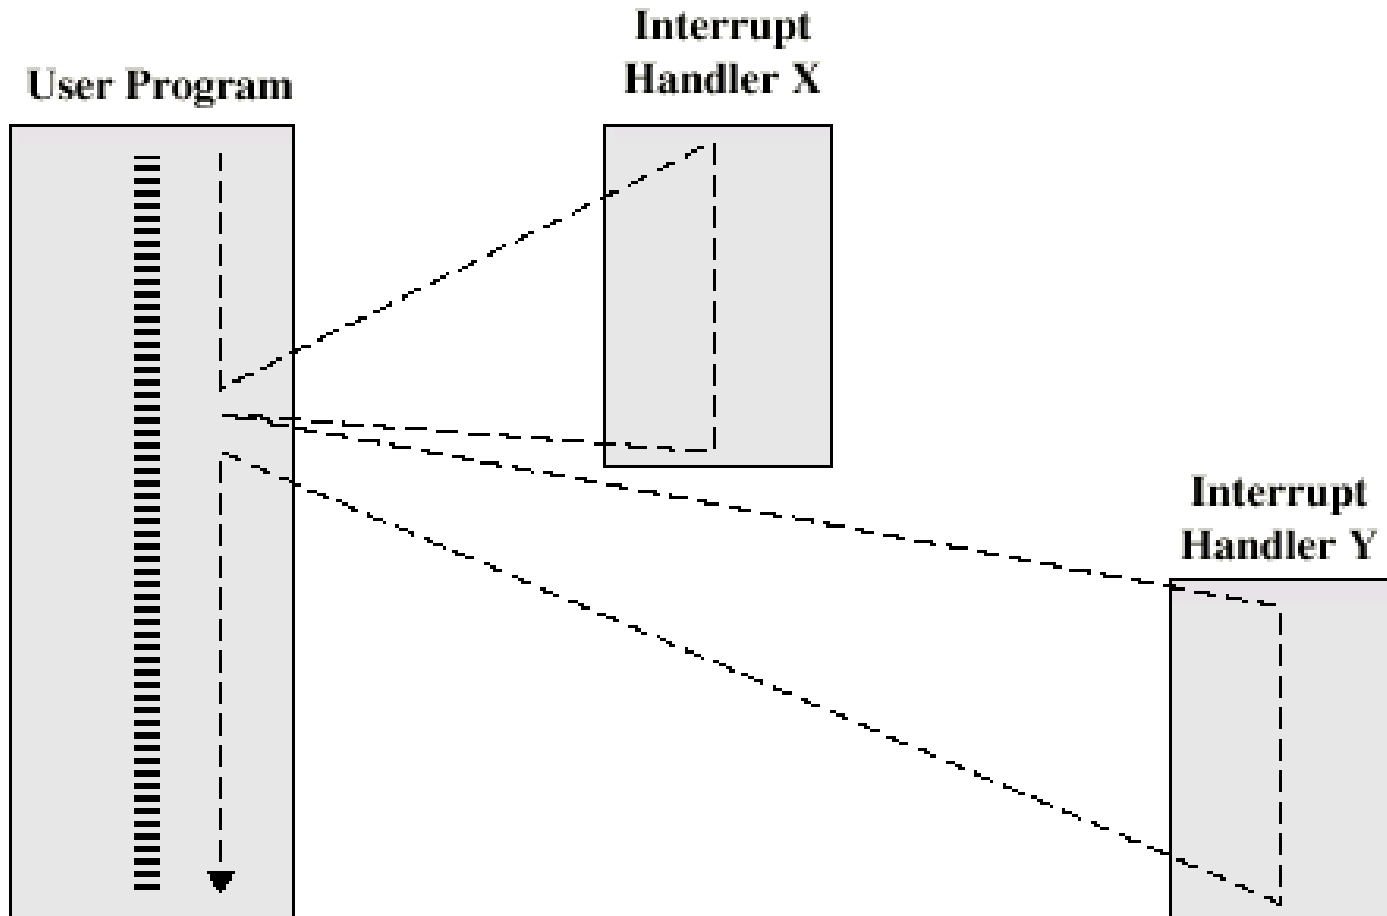
❑ Disable interrupts

- ❑ Processor will ignore further interrupts whilst processing one interrupt
- ❑ Interrupts remain pending and are checked after first interrupt has been processed
- ❑ Interrupts handled in sequence as they occur

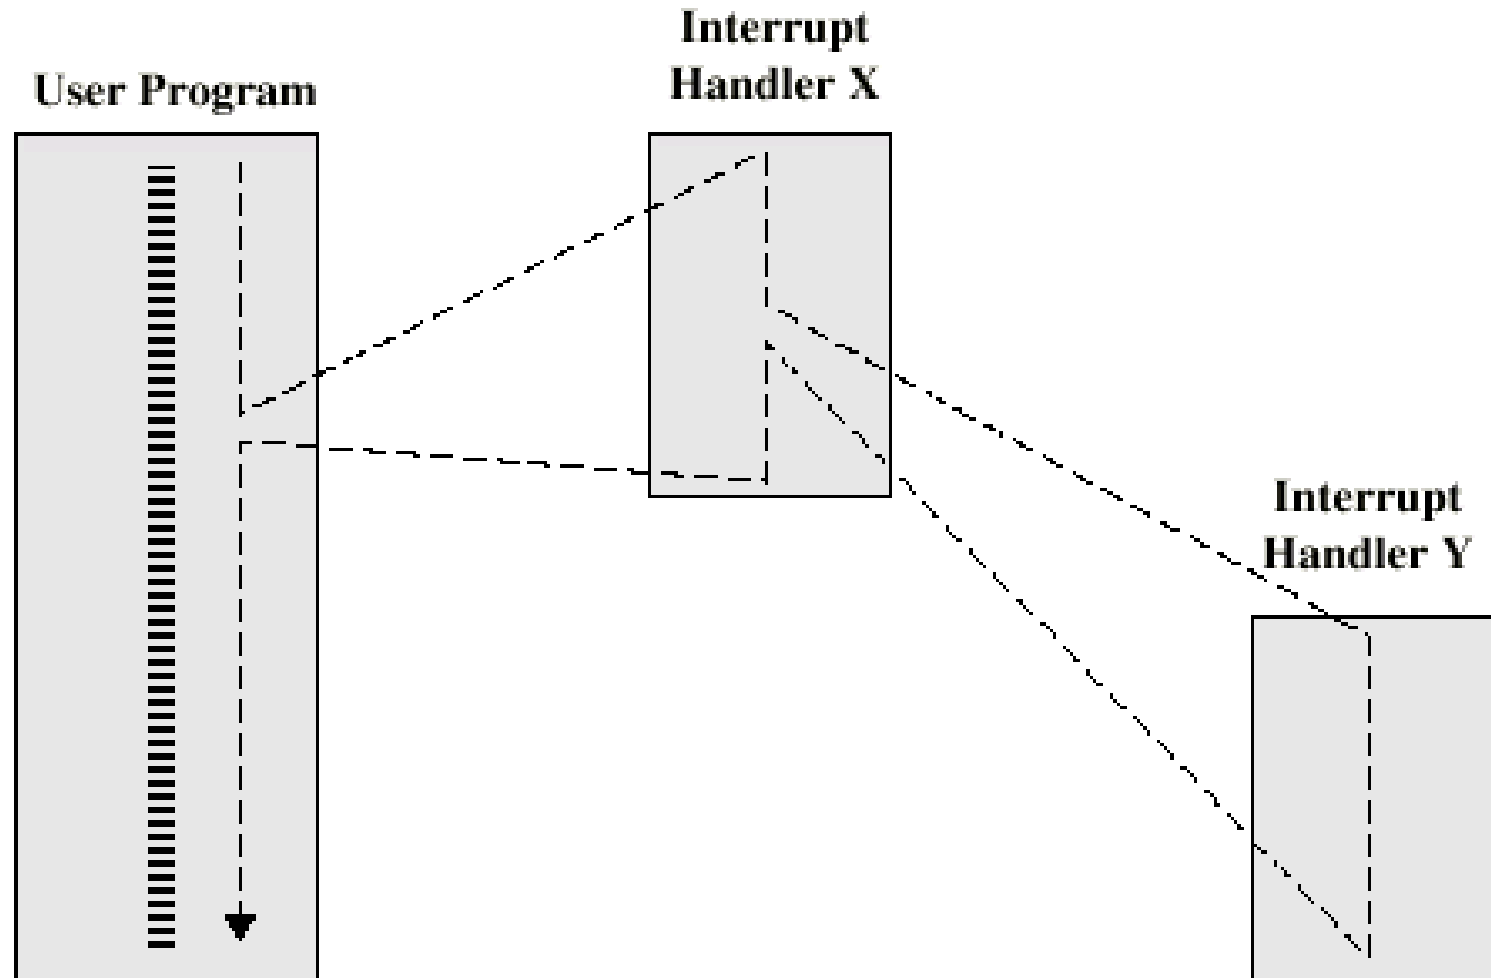
❑ Define priorities

- ❑ Low priority interrupts can be interrupted by higher priority interrupts
- ❑ When higher priority interrupt has been processed, processor returns to previous interrupt

Multiple Interrupts - Sequential



Multiple Interrupts - Nested



Connecting

- ❑ All the units must be connected
- ❑ Different type of connection for different type of unit
 - ❑ Memory
 - ❑ Input/Output
 - ❑ CPU

Memory Connection

- ❑ Receives and sends data
- ❑ Receives addresses (of locations)
- ❑ Receives control signals
 - ❑ Read
 - ❑ Write
 - ❑ Timing

Input/Output Connection(1)

- Similar to memory from computer's viewpoint
- Output
 - Receive data from computer
 - Send data to peripheral
- Input
 - Receive data from peripheral
 - Send data to computer

Input/Output Connection(2)

- Receive control signals from computer
- Send control signals to peripherals
 - e.g. spin disk
- Receive addresses from computer
 - e.g. port number to identify peripheral
- Send interrupt signals (control)

CPU Connection

- ❑ Reads instruction and data
- ❑ Writes out data (after processing)
- ❑ Sends control signals to other units
- ❑ Receives (& acts on) interrupts

Buses

- There are a number of possible interconnection systems
- Single and multiple BUS structures are most common
- e.g. Control/Address/Data bus (PC)
- e.g. Unibus (DEC-PDP)

What is a Bus?

- ❑ A communication pathway connecting two or more devices
- ❑ Usually broadcast
- ❑ Often grouped
 - ❑ A number of channels in one bus
 - ❑ e.g. 32 bit data bus is 32 separate single bit channels
- ❑ Power lines may not be shown

Data Bus

- Carries data
 - Remember that there is no difference between “data” and “instruction” at this level
- Width is a key determinant of performance
 - 8, 16, 32, 64 bit

Address bus

- Identify the source or destination of data
- e.g. CPU needs to read an instruction (data) from a given location in memory
- Bus width determines maximum memory capacity of system
 - e.g. 8080 has 16 bit address bus giving 64k address space

Control Bus


- Control and timing information
 - Memory read/write signal
 - Interrupt request
 - Clock signals

Control Bus

□ Typical control lines include

- Memory write
- Memory read
- I/O write
- I/O read
- Transfer ACK
- Bus request
- Bus grant
- Interrupt request
- Interrupt ACK
- Clock
- Reset

Bus Interconnection Scheme



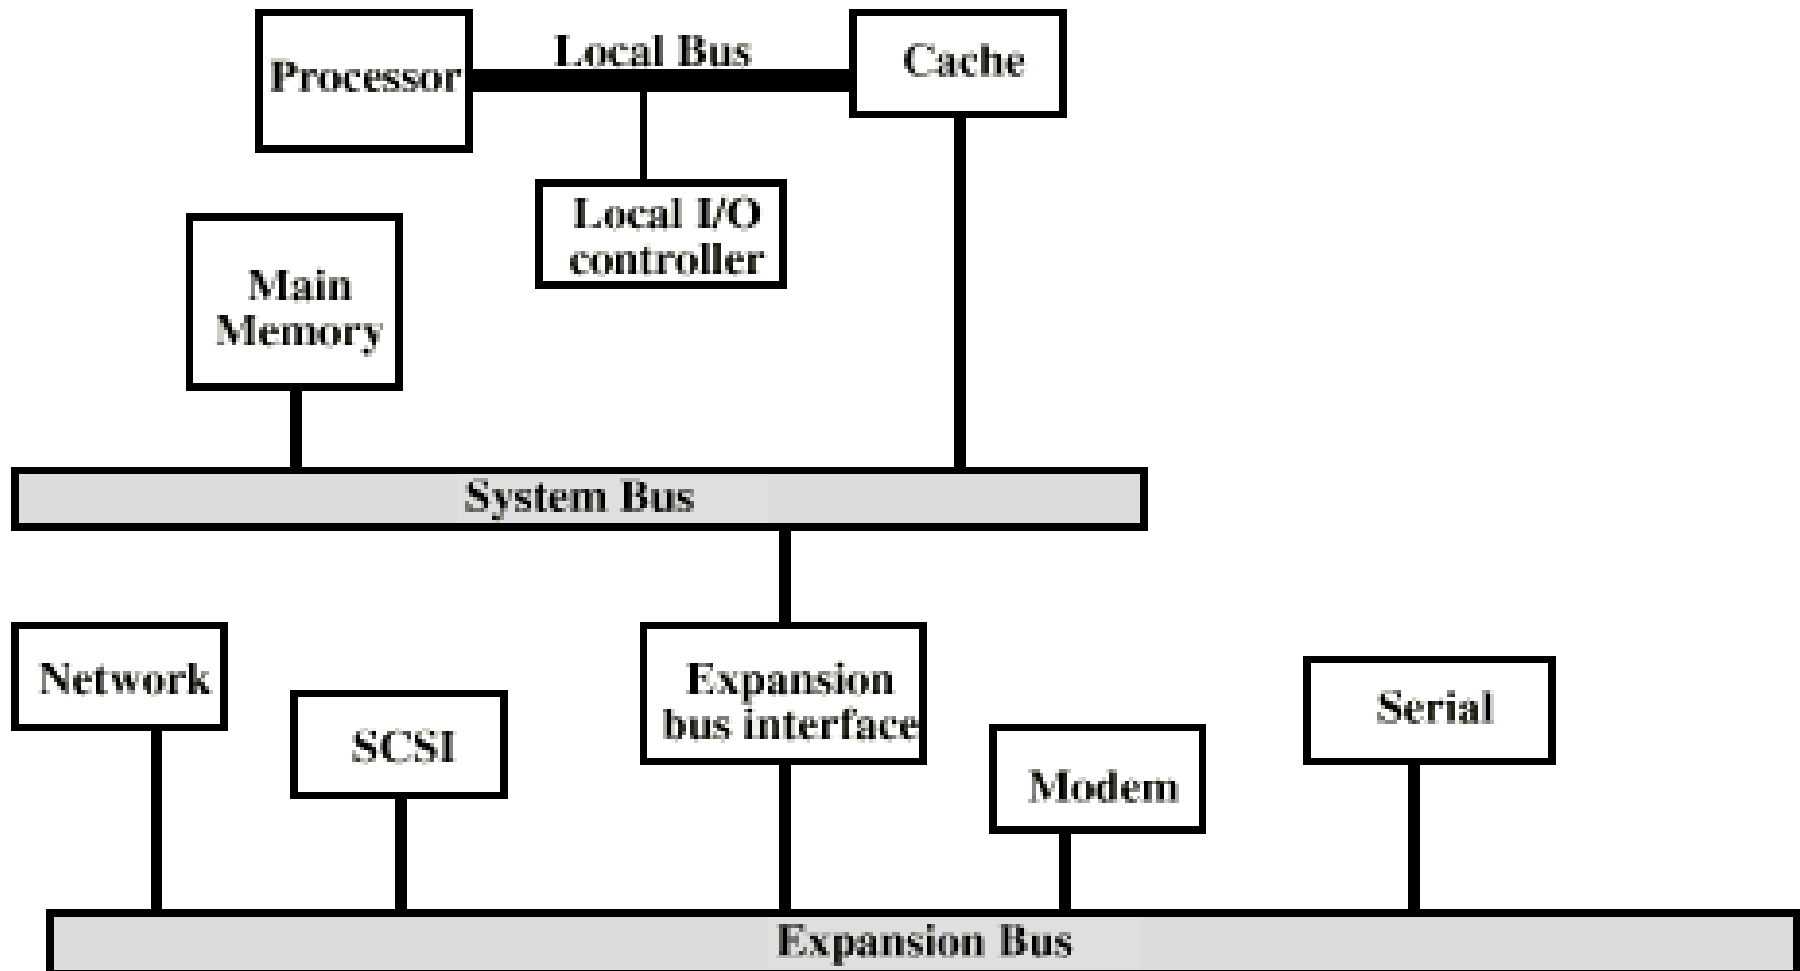
Big and Yellow?

- What do buses look like?
 - Parallel lines on circuit boards
 - Ribbon cables
 - Strip connectors on mother boards
 - e.g. PCI
 - Sets of wires

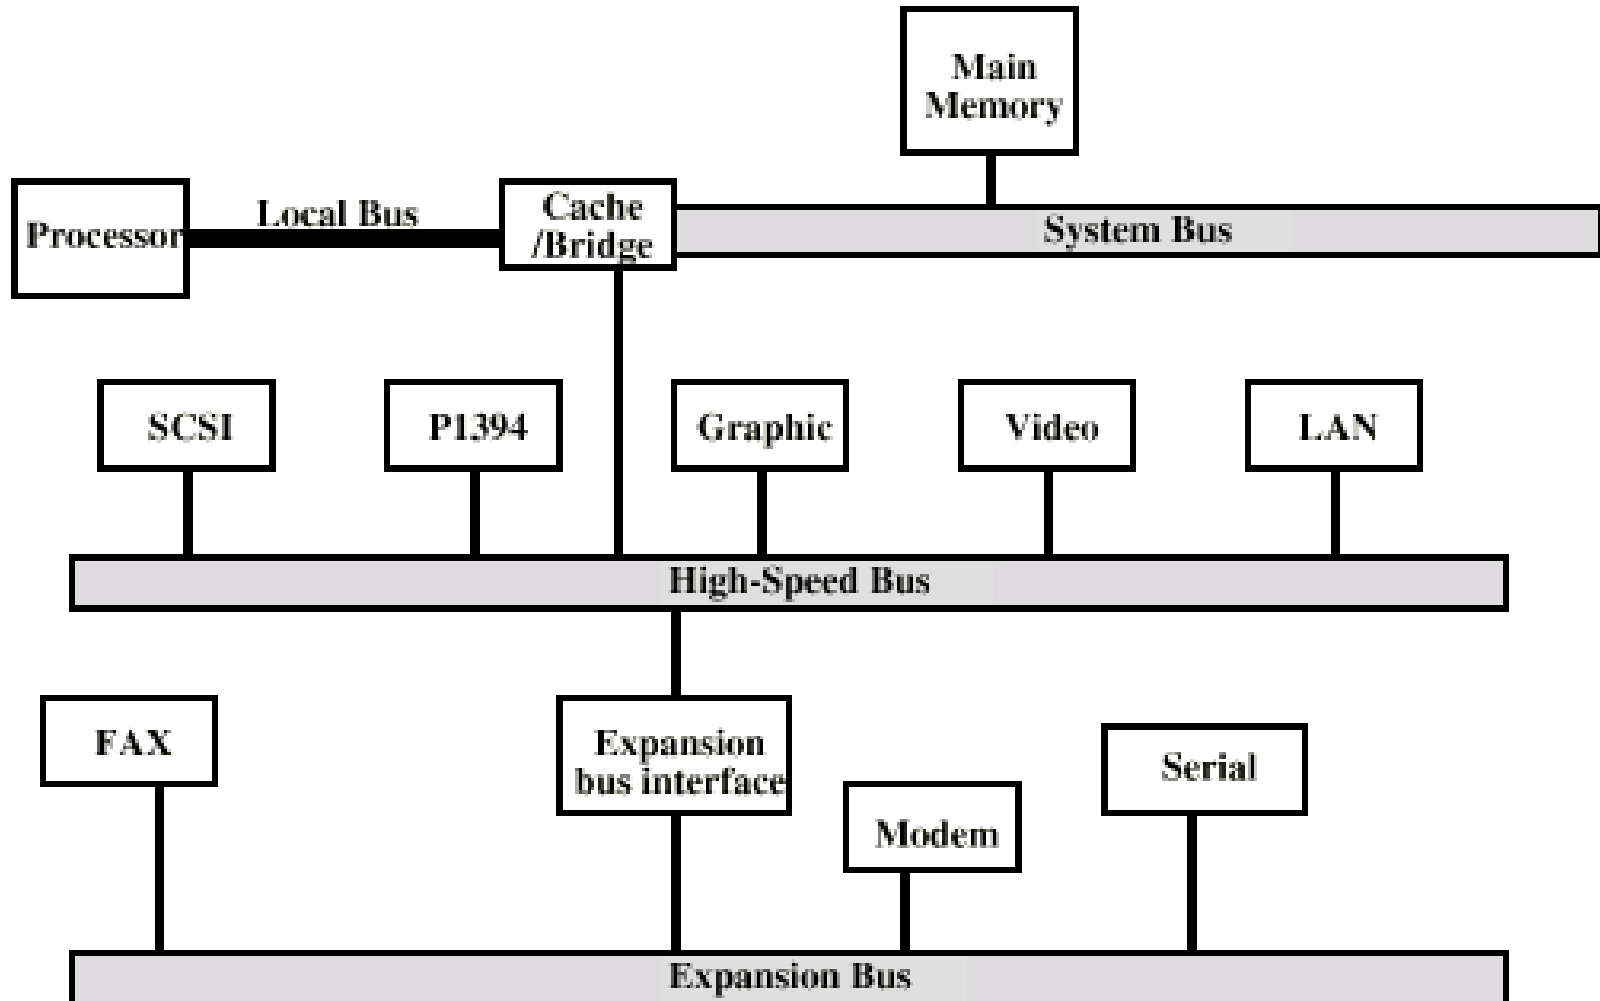
Single Bus Problems

- ❑ Lots of devices on one bus leads to:
 - ❑ Propagation delays
 - ❑ Long data paths mean that co-ordination of bus use can adversely affect performance
 - ❑ If aggregate data transfer approaches bus capacity
- ❑ Most systems use multiple buses to overcome these problems

Traditional (ISA) (with cache)



High Performance Bus



Bus Types

- ❑ Dedicated

- ❑ Separate data & address lines

- ❑ Multiplexed

- ❑ Shared lines
 - ❑ Address valid or data valid control line
 - ❑ Advantage - fewer lines
 - ❑ Disadvantages
 - ❑ More complex control
 - ❑ Ultimate performance

Elements of Bus Design

- Type
 - Dedicated Address
 - Multiplexed Data
- Bus Width
 - Address
 - Data
- Method of Arbitration
 - Centralized Read
 - Distributed Write
- Timing Read-modify-write
 - Synchronous Read-after-write
 - Asynchronous Block

Elements of Bus Design

☐ Data Transfer Type

- ☐ Read
- ☐ Write
- ☐ Read-modify-write
- ☐ Read-after-write
- ☐ Block

Bus Arbitration

- ❑ More than one module controlling the bus
- ❑ e.g. CPU and DMA controller
- ❑ Only one module may control bus at one time
- ❑ Arbitration may be centralised or distributed

Centralised Arbitration

- ❑ Single hardware device controlling bus access
 - ❑ Bus Controller
 - ❑ Arbiter
- ❑ May be part of CPU or separate

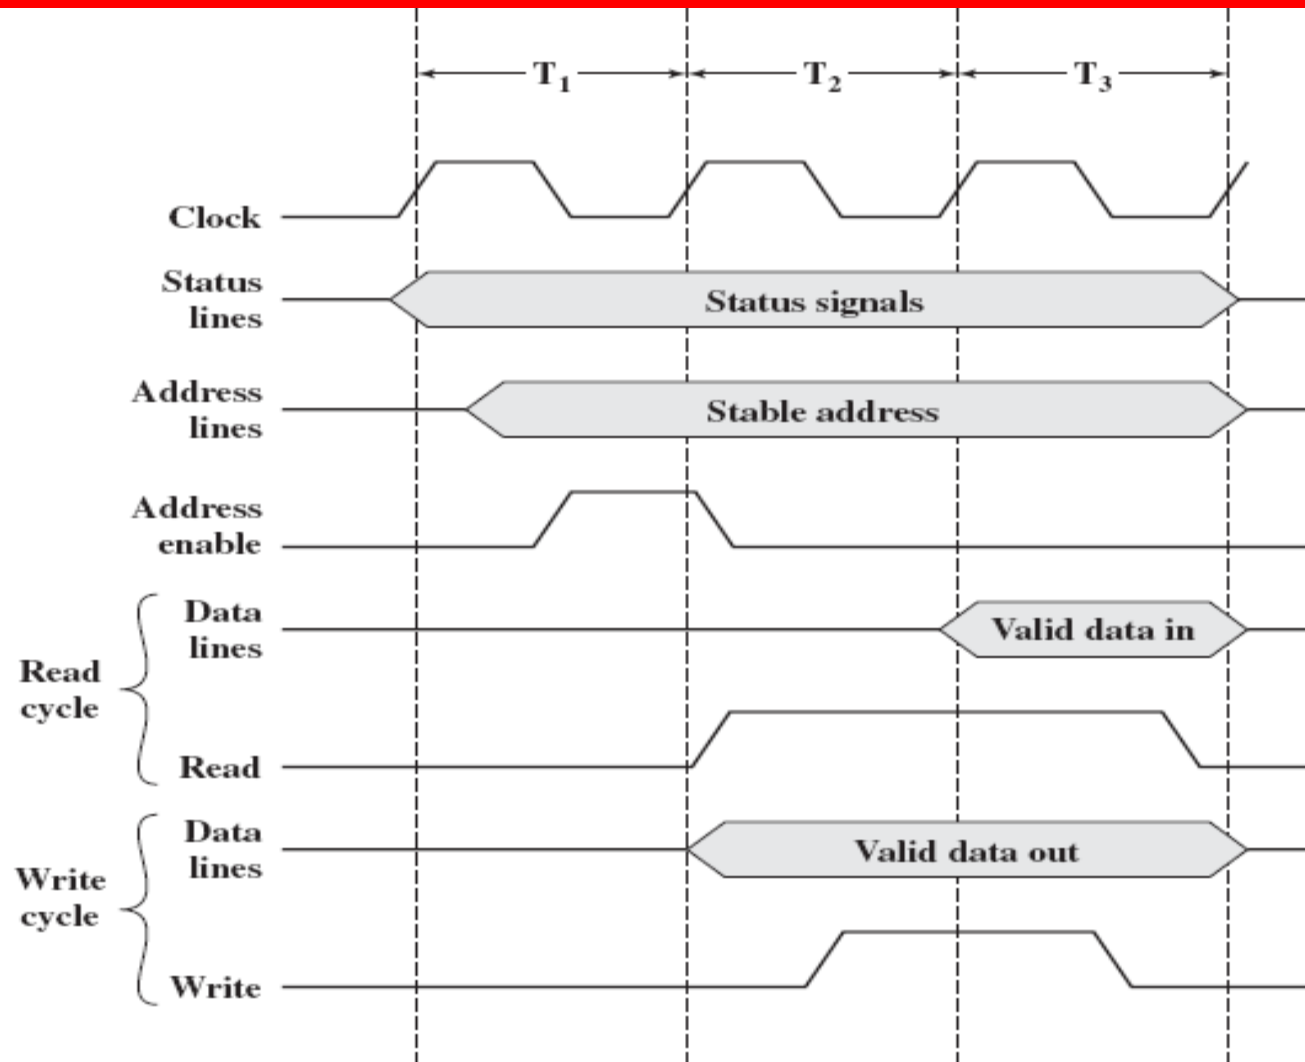
Distributed Arbitration

- Each module may claim the bus
- Control logic on all modules


Timing

- ❑ Co-ordination of events on bus
- ❑ Synchronous
 - ❑ Events determined by clock signals
 - ❑ Control Bus includes clock line
 - ❑ A single 1-0 is a bus cycle
 - ❑ All devices can read clock line
 - ❑ Usually sync on leading edge
 - ❑ Usually a single cycle for an event

Synchronous Timing Diagram



Asynchronous Timing Diagram



(a) System bus read cycle



(b) System bus write cycle

PCI Bus

- ❑ Peripheral Component Interconnection
- ❑ Intel released to public domain
- ❑ 32 or 64 bit
- ❑ 50 lines

PCI Bus Lines (required)

- ❑ Systems lines
 - ❑ Including clock and reset
- ❑ Address & Data
 - ❑ 32 time mux lines for address/data
 - ❑ Interrupt & validate lines
- ❑ Interface Control
- ❑ Arbitration
 - ❑ Not shared
 - ❑ Direct connection to PCI bus arbiter
- ❑ Error lines

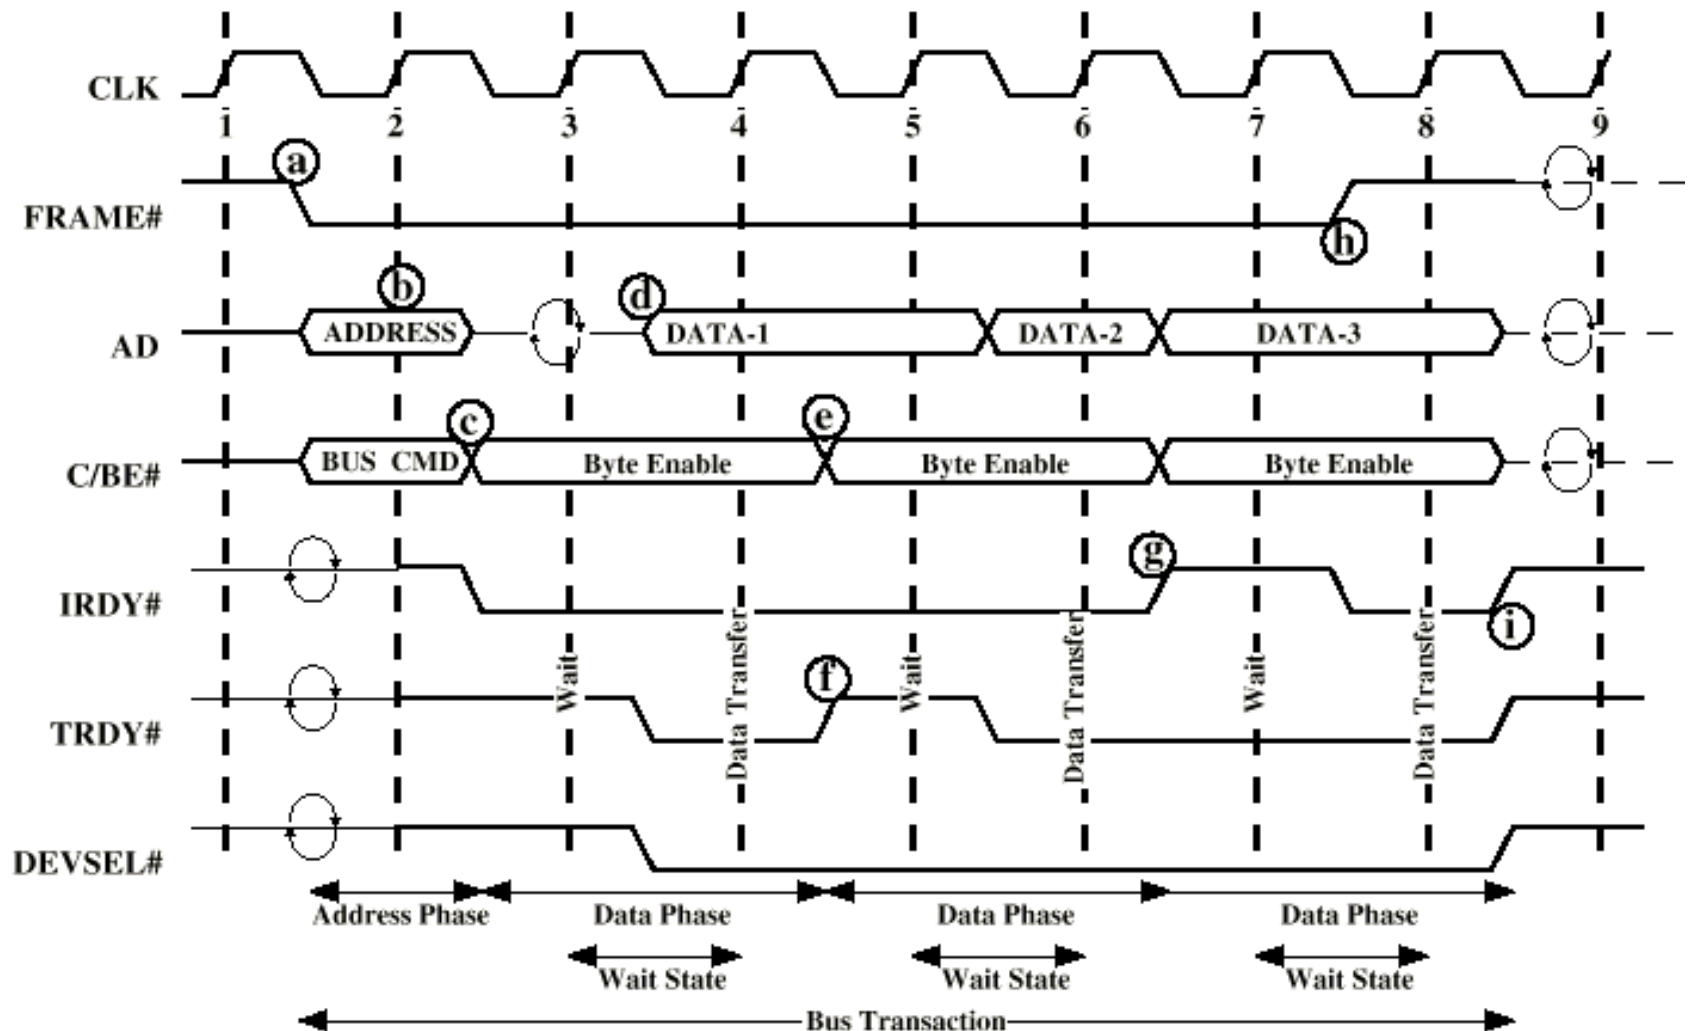
PCI Bus Lines (Optional)

- ❑ Interrupt lines
 - ❑ Not shared
- ❑ Cache support
- ❑ 64-bit Bus Extension
 - ❑ Additional 32 lines
 - ❑ Time multiplexed
 - ❑ 2 lines to enable devices to agree to use 64-bit transfer
- ❑ JTAG/Boundary Scan
 - ❑ For testing procedures

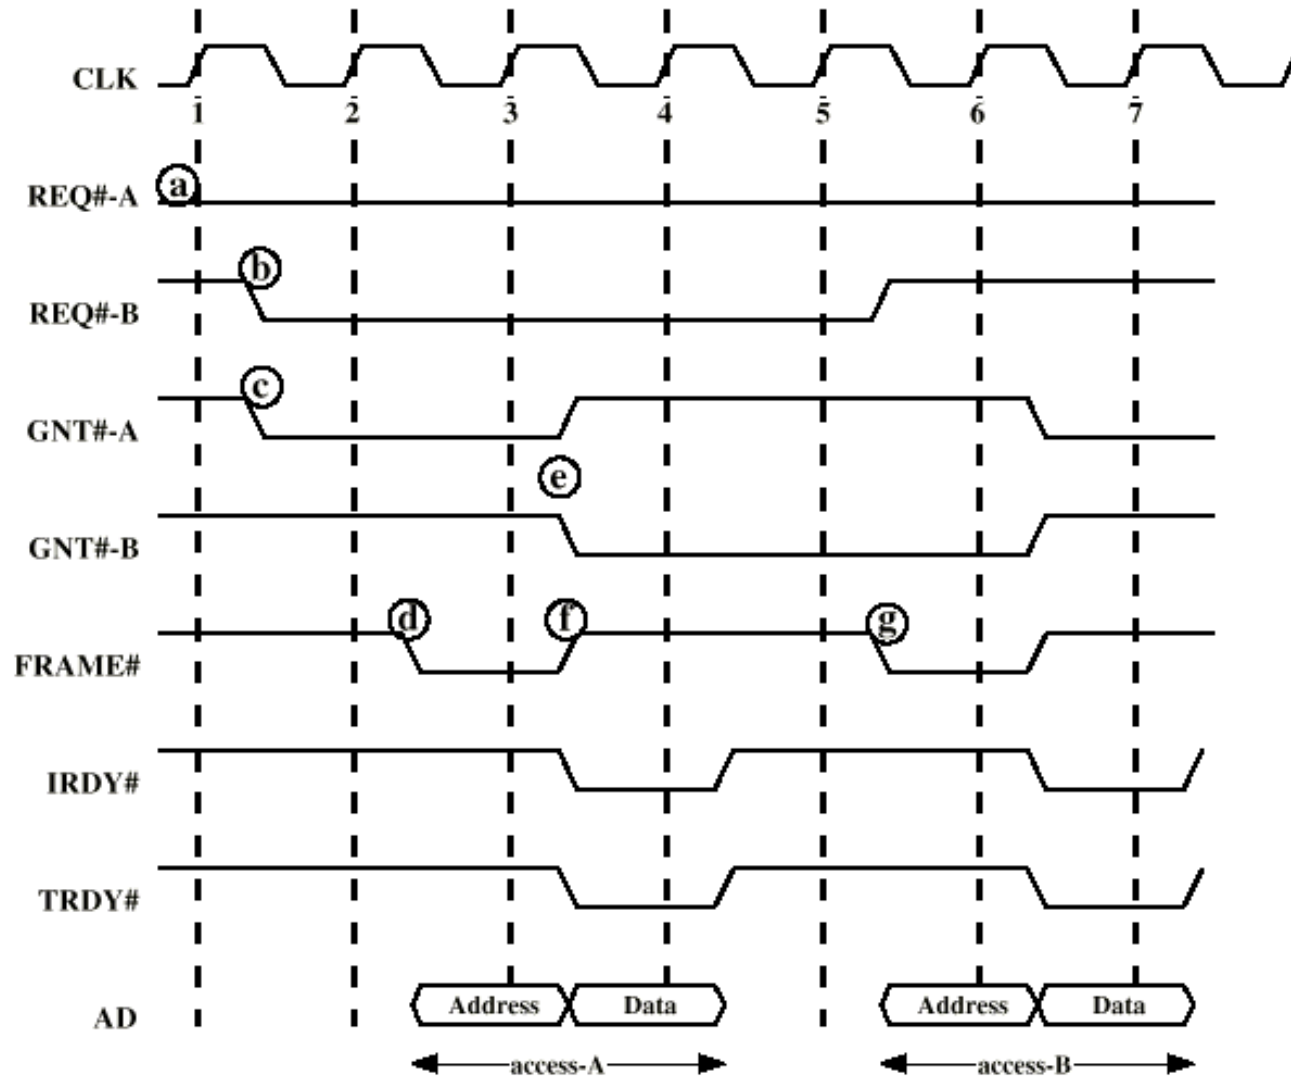
PCI Commands

- ❑ Transaction between initiator (master) and target
- ❑ Master claims bus
- ❑ Determine type of transaction
 - ❑ e.g. I/O read/write
- ❑ Address phase
- ❑ One or more data phases

PCI Read Timing Diagram



PCI Bus Arbitration



Foreground Reading

- ❑ Stallings, chapter 3 (all of it)
- ❑ www.pcguide.com/ref/mbsys/buses/
- ❑ In fact, read the whole site!
- ❑ www.pcguide.com/

William Stallings

Computer Organization

and Architecture

Chapter 4

Internal Memory

Characteristics

- ❑ Location
- ❑ Capacity
- ❑ Unit of transfer
- ❑ Access method
- ❑ Performance
- ❑ Physical type
- ❑ Physical characteristics
- ❑ Organisation

Location

- ☐ CPU
- ☐ Internal
- ☐ External

Capacity

- Word size
 - The natural unit of organisation
- Number of words
 - or Bytes

Unit of Transfer

- ❑ Internal
 - ❑ Usually governed by data bus width
- ❑ External
 - ❑ Usually a block which is much larger than a word
- ❑ Addressable unit
 - ❑ Smallest location which can be uniquely addressed
 - ❑ Word internally
 - ❑ Cluster on M\$ disks

Access Methods (1)

□ Sequential

- Start at the beginning and read through in order
- Access time depends on location of data and previous location
- e.g. tape

□ Direct

- Individual blocks have unique address
- Access is by jumping to vicinity plus sequential search
- Access time depends on location and previous location
- e.g. disk

Access Methods (2)

□ Random

- Individual addresses identify locations exactly
- Access time is independent of location or previous access
- e.g. RAM

□ Associative

- Data is located by a comparison with contents of a portion of the store
- Access time is independent of location or previous access
- e.g. cache

Memory Hierarchy

- Registers
 - In CPU
- Internal or Main memory
 - May include one or more levels of cache
 - “RAM”
- External memory
 - Backing store

Performance

- ❑ Access time
 - ❑ Time between presenting the address and getting the valid data
- ❑ Memory Cycle time
 - ❑ Time may be required for the memory to “recover” before next access
 - ❑ Cycle time is access + recovery
- ❑ Transfer Rate
 - ❑ Rate at which data can be moved

Physical Types

- Semiconductor

 - RAM

- Magnetic

 - Disk & Tape

- Optical

 - CD & DVD

- Others

 - Bubble

 - Hologram

Physical Characteristics

- ❑ Decay
- ❑ Volatility
- ❑ Erasable
- ❑ Power consumption

Organisation

- Physical arrangement of bits into words
- Not always obvious
- e.g. interleaved

The Bottom Line

- How much?
 - Capacity
- How fast?
 - Time is money
- How expensive?

Hierarchy List

- Registers
- L1 Cache
- L2 Cache
- Main memory
- Disk cache
- Disk
- Optical
- Tape

So you want fast?

- ❑ It is possible to build a computer which uses only static RAM (see later)
- ❑ This would be very fast
- ❑ This would need no cache
 - ❑ How can you cache cache?
- ❑ This would cost a very large amount

Locality of Reference

- During the course of the execution of a program, memory references tend to cluster
- e.g. loops

Semiconductor Memory

□ RAM

- Misnamed as all semiconductor memory is random access
- Read/Write
- Volatile
- Temporary storage
- Static or dynamic

Dynamic RAM

- ❑ Bits stored as charge in capacitors
- ❑ Charges leak
- ❑ Need refreshing even when powered
- ❑ Simpler construction
- ❑ Smaller per bit
- ❑ Less expensive
- ❑ Need refresh circuits
- ❑ Slower
- ❑ Main memory

Static RAM

- ❑ Bits stored as on/off switches
- ❑ No charges to leak
- ❑ No refreshing needed when powered
- ❑ More complex construction
- ❑ Larger per bit
- ❑ More expensive
- ❑ Does not need refresh circuits
- ❑ Faster
- ❑ Cache

Read Only Memory (ROM)

- ❑ Permanent storage
- ❑ Microprogramming (see later)
- ❑ Library subroutines
- ❑ Systems programs (BIOS)
- ❑ Function tables

Types of ROM

- ❑ Written during manufacture
 - ❑ Very expensive for small runs
- ❑ Programmable (once)
 - ❑ PROM
 - ❑ Needs special equipment to program
- ❑ Read “mostly”
 - ❑ Erasable Programmable (EPROM)
 - ❑ Erased by UV
 - ❑ Electrically Erasable (EEPROM)
 - ❑ Takes much longer to write than read
 - ❑ Flash memory
 - ❑ Erase whole memory electrically

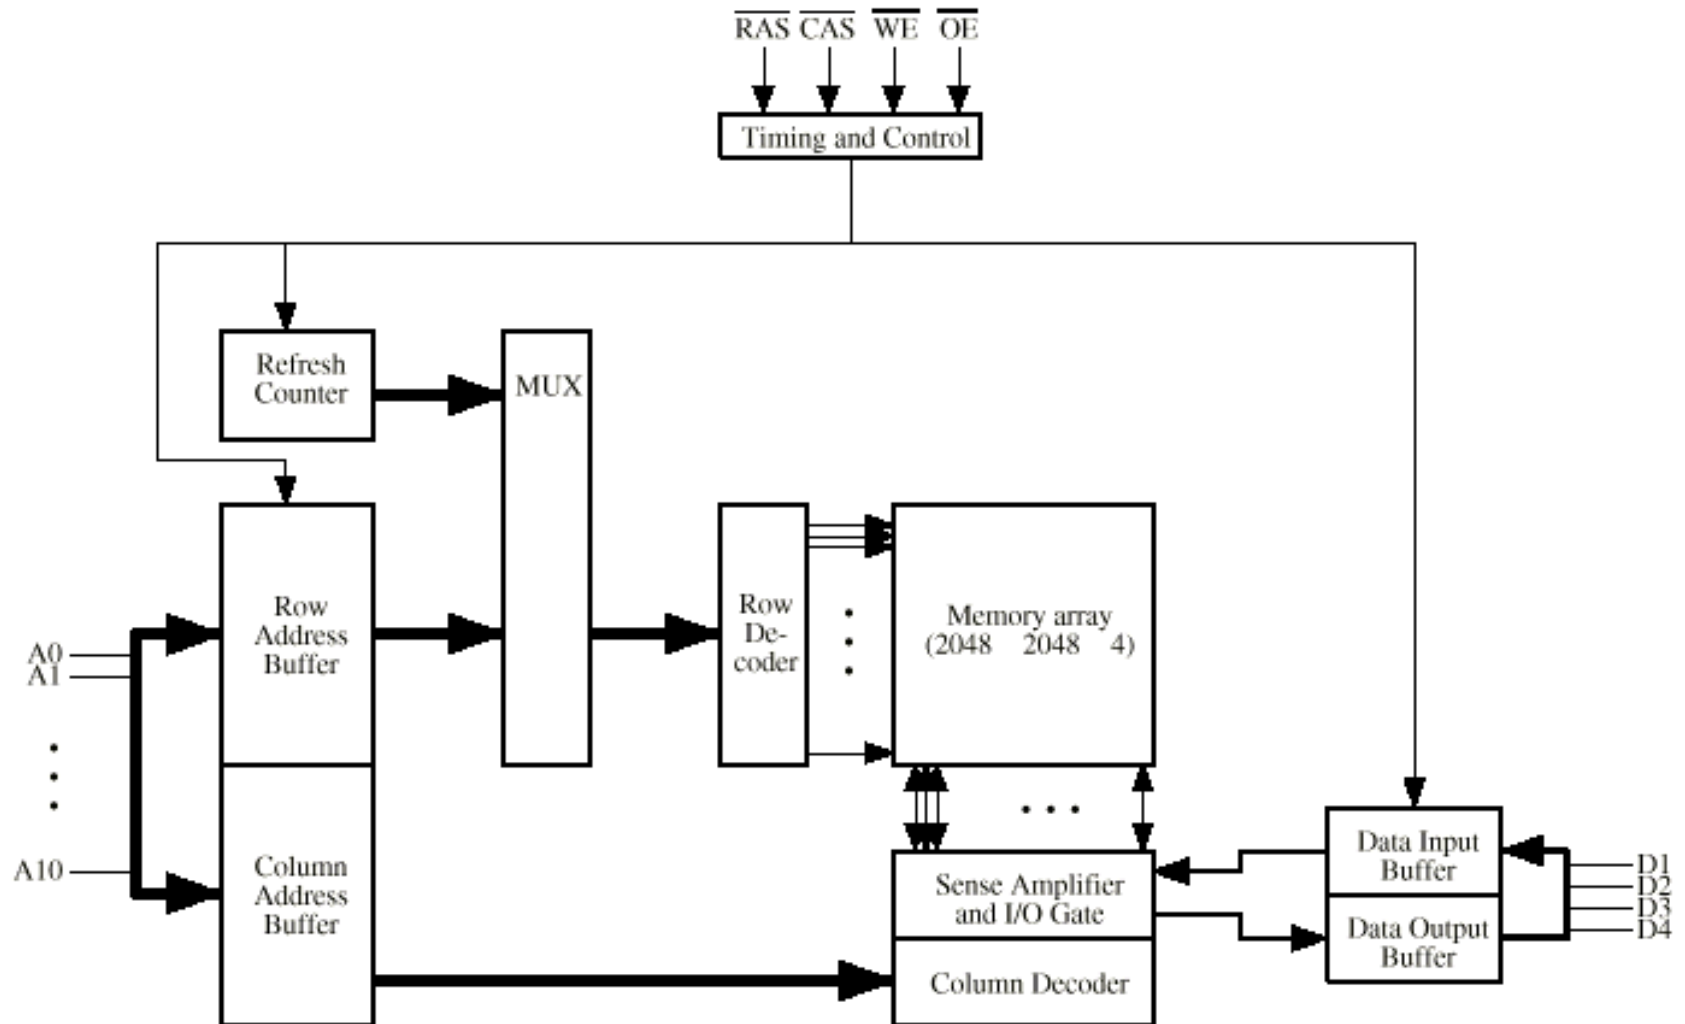
Organisation in detail

- ❑ A 16Mbit chip can be organised as 1M of 16 bit words
- ❑ A bit per chip system has 16 lots of 1Mbit chip with bit 1 of each word in chip 1 and so on
- ❑ A 16Mbit chip can be organised as a 2048 x 2048 x 4bit array
 - ❑ Reduces number of address pins
 - ❑ Multiplex row address and column address
 - ❑ 11 pins to address ($2^{11}=2048$)
 - ❑ Adding one more pin doubles range of values so x4 capacity


Refreshing

- ❑ Refresh circuit included on chip
- ❑ Disable chip
- ❑ Count through rows
- ❑ Read & Write back
- ❑ Takes time
- ❑ Slows down apparent performance


Typical 16 Mb DRAM (4M x 4)



Packaging

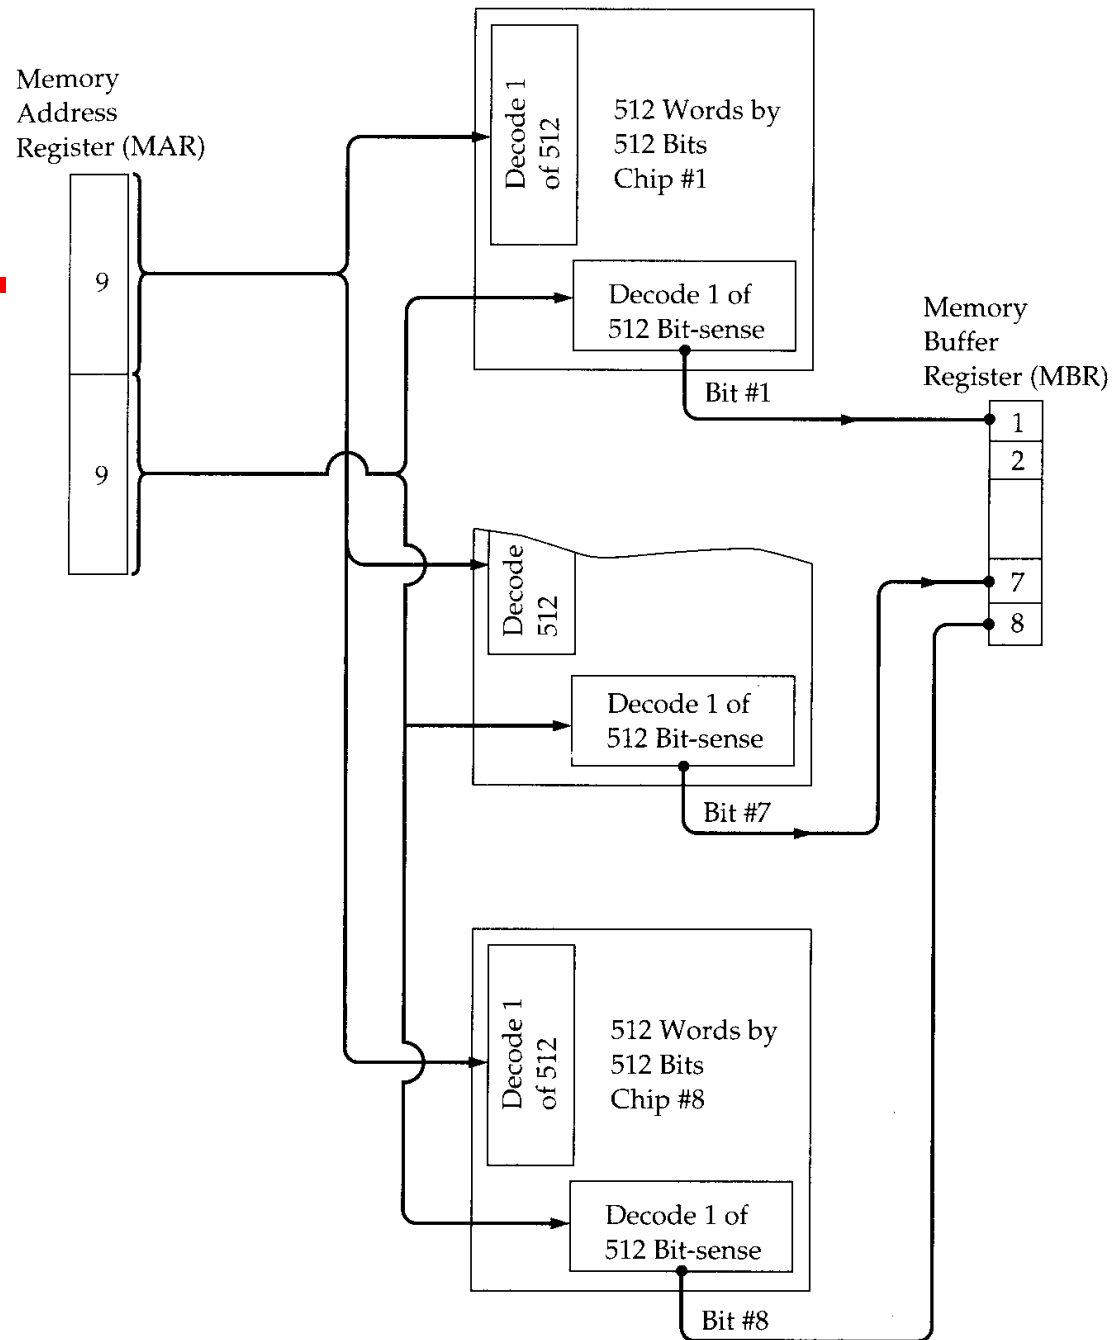


(a) 8 Mbit EPROM

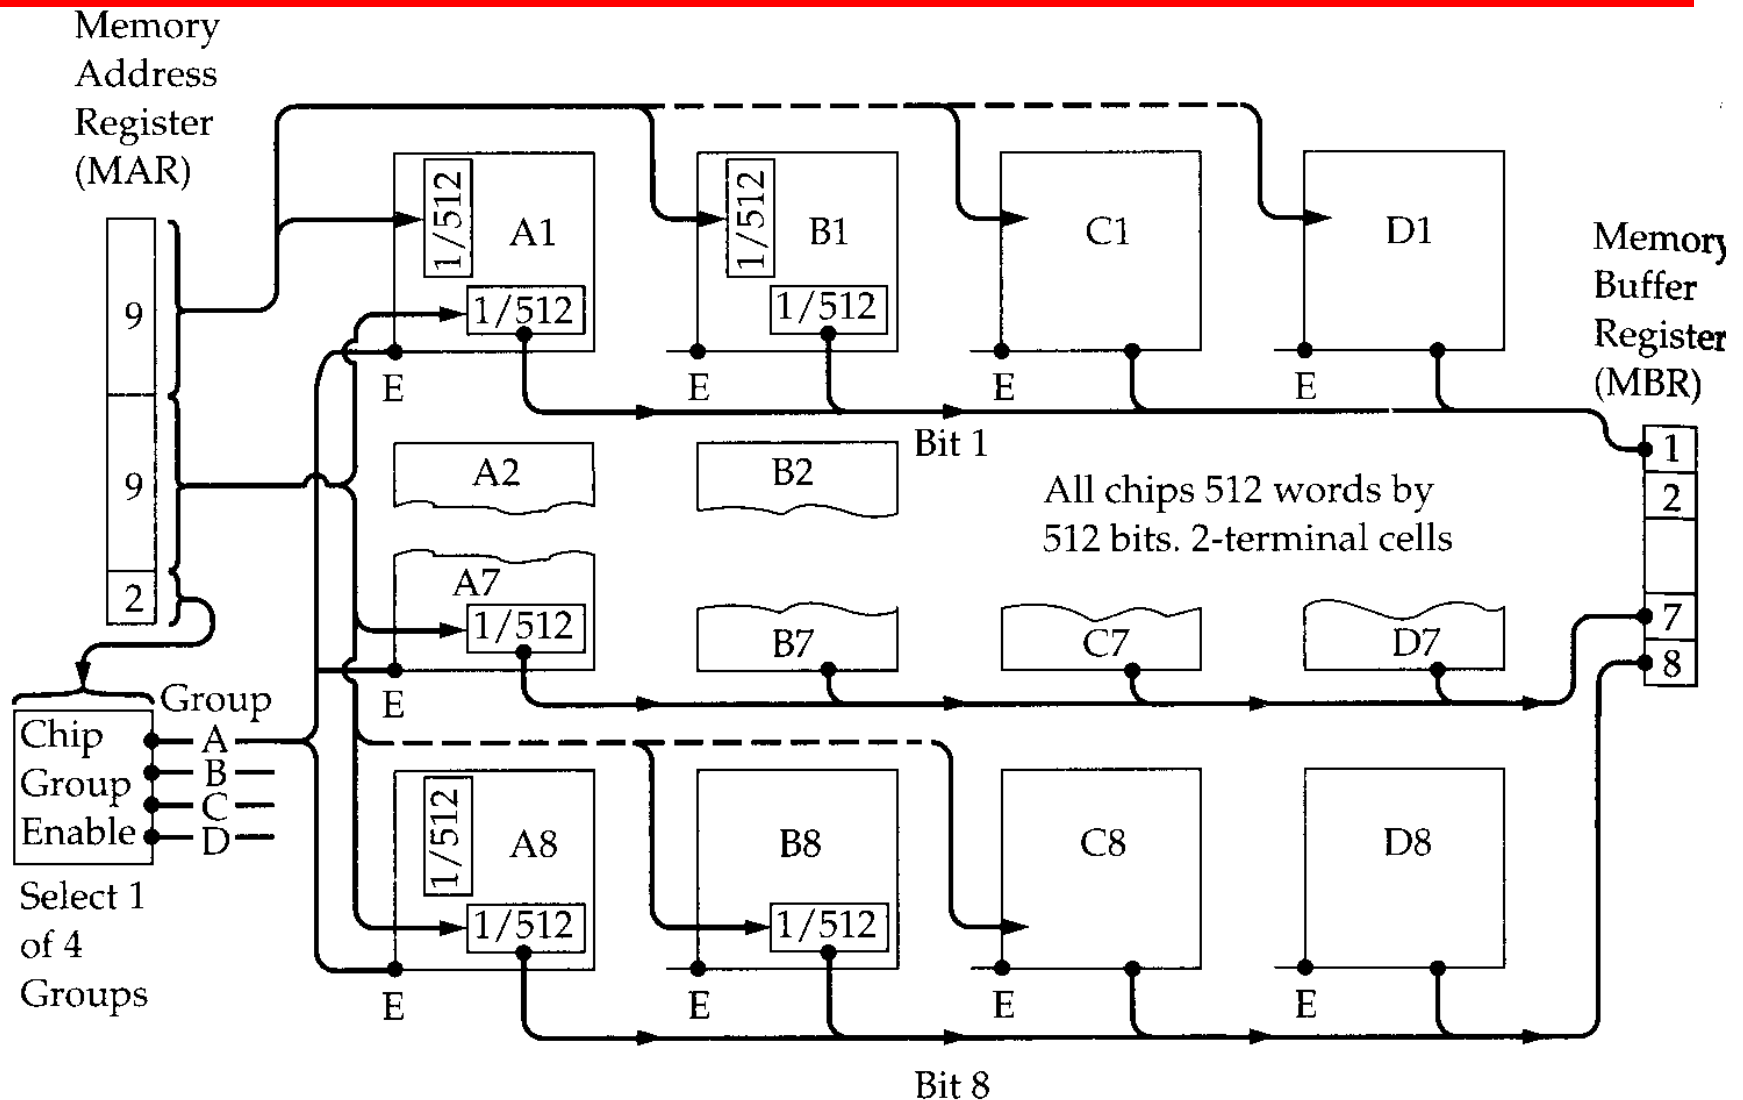


(b) 16 Mbit DRAM

Module Organisation




Module Organisation (2)



Error Correction


- ❑ Hard Failure
 - ❑ Permanent defect
- ❑ Soft Error
 - ❑ Random, non-destructive
 - ❑ No permanent damage to memory
- ❑ Detected using Hamming error correcting code

Error Correcting Code Function



Cache

- ❑ Small amount of fast memory
- ❑ Sits between normal main memory and CPU
- ❑ May be located on CPU chip or module



Cache operation - overview

- ❑ CPU requests contents of memory location
- ❑ Check cache for this data
- ❑ If present, get from cache (fast)
- ❑ If not present, read required block from main memory to cache
- ❑ Then deliver from cache to CPU
- ❑ Cache includes tags to identify which block of main memory is in each cache slot

Cache Design

- Size
- Mapping Function
- Replacement Algorithm
- Write Policy
- Block Size
- Number of Caches

Size does matter

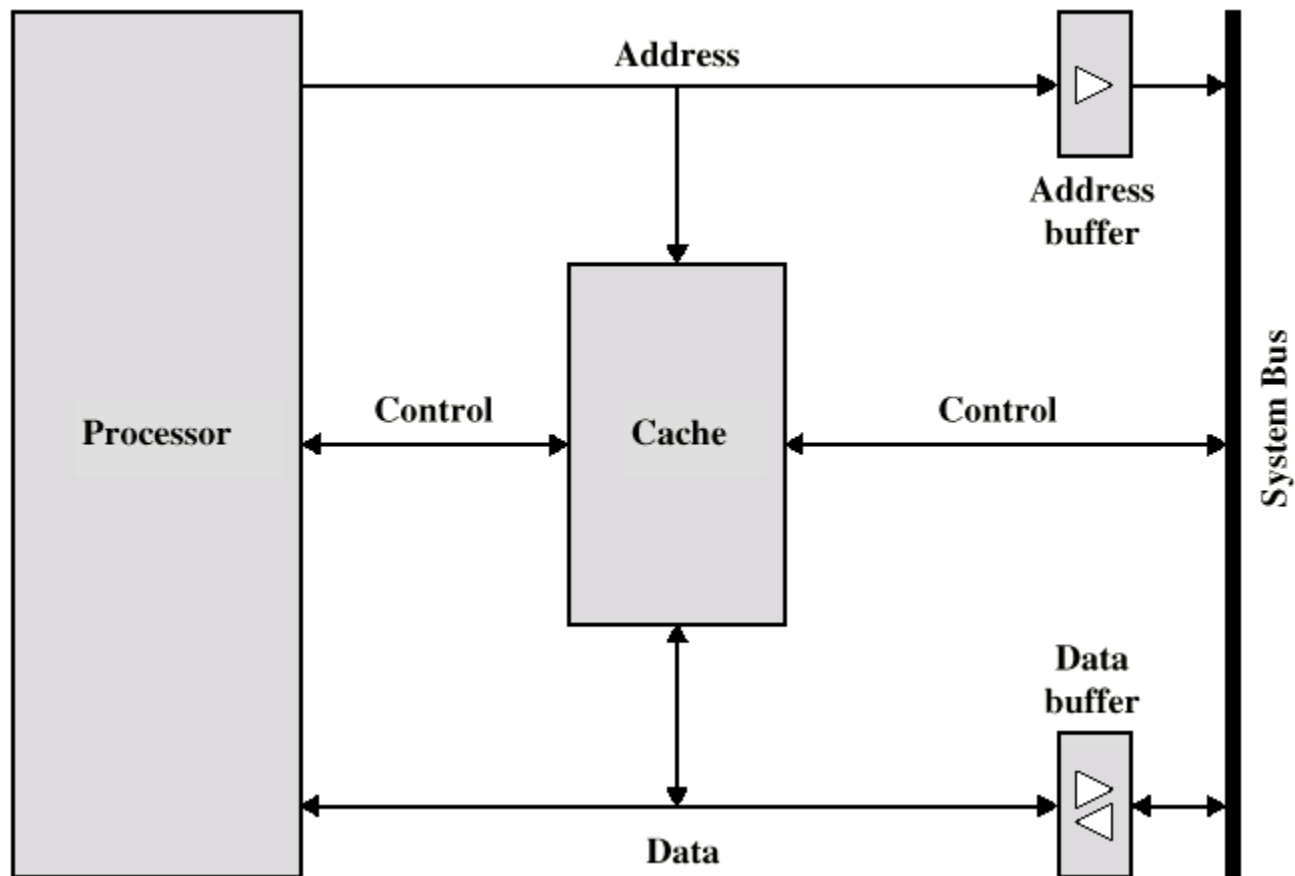
- ❑ Cost

- ❑ More cache is expensive

- ❑ Speed

- ❑ More cache is faster (up to a point)
 - ❑ Checking cache for data takes time

Typical Cache Organization



Mapping Function

- ❑ Cache of 64kByte
- ❑ Cache block of 4 bytes
 - ❑ i.e. cache is 16k (2^{14}) lines of 4 bytes
- ❑ 16MBytes main memory
- ❑ 24 bit address
 - ❑ ($2^{24}=16\text{M}$)

Direct Mapping

- ❑ Each block of main memory maps to only one cache line
 - ❑ i.e. if a block is in cache, it must be in one specific place
- ❑ Address is in two parts
- ❑ Least Significant w bits identify unique word
- ❑ Most Significant s bits specify one memory block
- ❑ The MSBs are split into a cache line field r and a tag of $s-r$ (most significant)

Direct Mapping Address Structure


Tag s-r	Line or Slot r	Word w
8	14	2

- ❑ 24 bit address
- ❑ 2 bit word identifier (4 byte block)
- ❑ 22 bit block identifier
 - ❑ 8 bit tag (=22-14)
 - ❑ 14 bit slot or line
- ❑ No two blocks in the same line have the same Tag field
- ❑ Check contents of cache by finding line and checking Tag

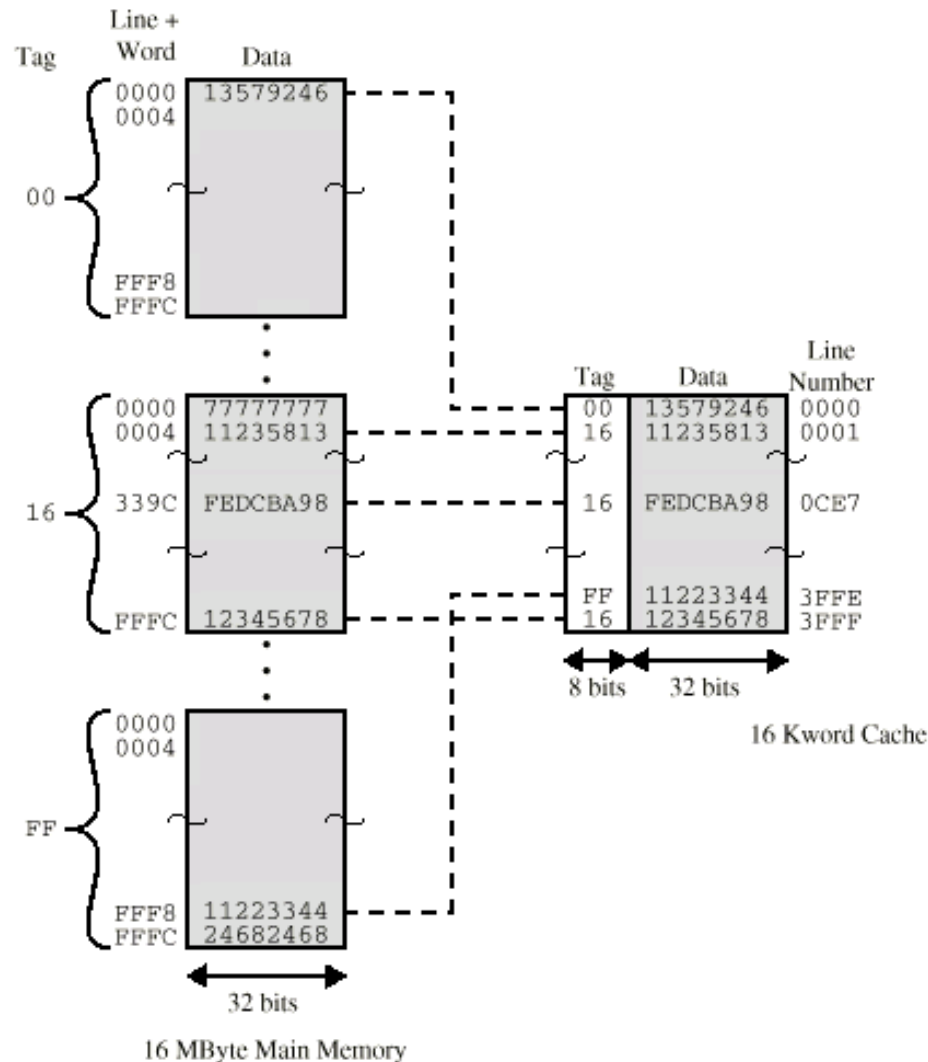
Direct Mapping Cache Line Table

Cache line	Main Memory blocks held
0	$0, m, 2m, 3m \dots 2^s - m$
1	$1, m+1, 2m+1 \dots 2^s - m + 1$
$m-1$	$m-1, 2m-1, 3m-1 \dots 2^s - 1$

Direct Mapping Cache Organization



Direct Mapping Example



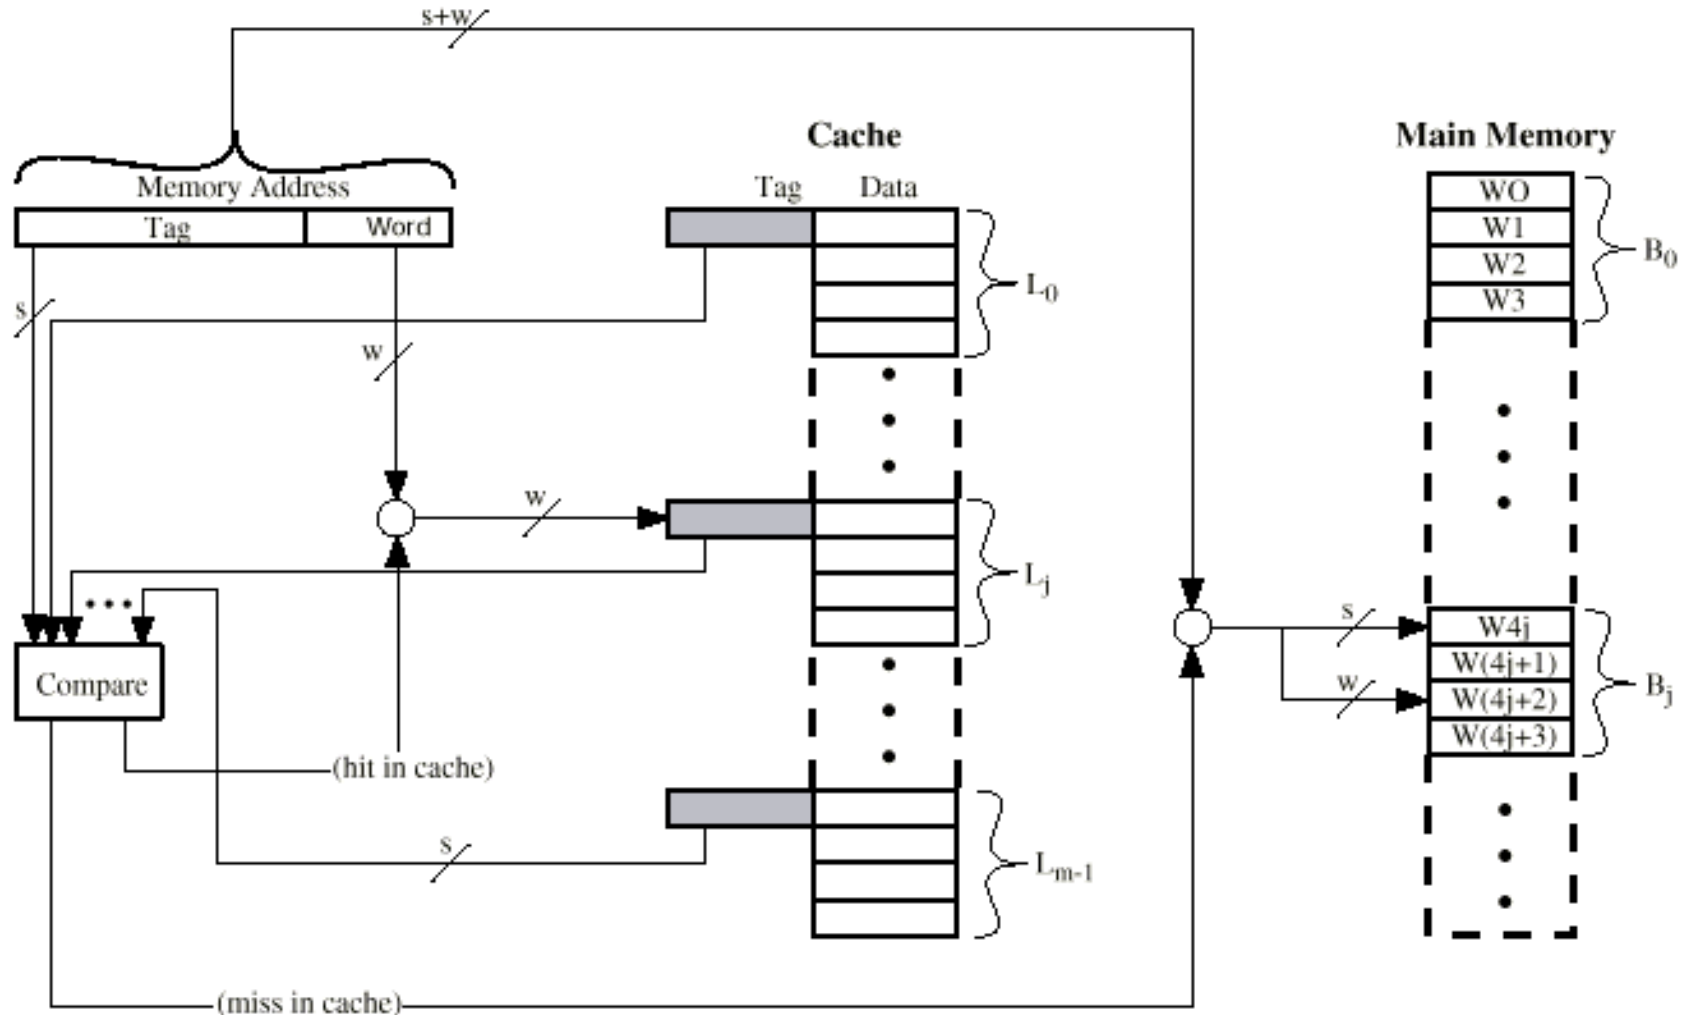
Direct Mapping pros & cons

- ❑ Simple
- ❑ Inexpensive
- ❑ Fixed location for given block
 - ❑ If a program accesses 2 blocks that map to the same line repeatedly, cache misses are very high

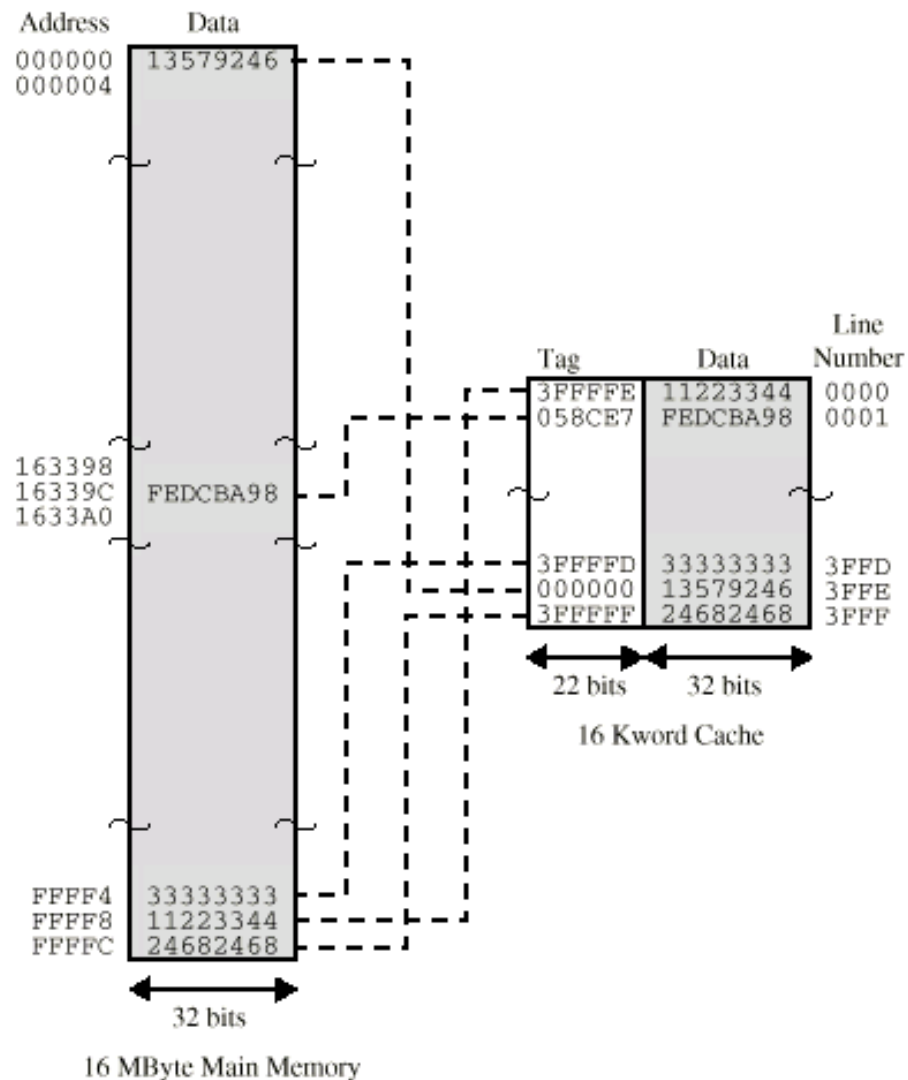
Associative Mapping

- ❑ A main memory block can load into any line of cache
- ❑ Memory address is interpreted as tag and word
- ❑ Tag uniquely identifies block of memory
- ❑ Every line's tag is examined for a match
- ❑ Cache searching gets expensive

Fully Associative Cache Organization



Associative Mapping Example



Associative Mapping Address Structure

Tag 22 bit	Word 2 bit
------------	---------------

- ❑ 22 bit tag stored with each 32 bit block of data
- ❑ Compare tag field with tag entry in cache to check for hit
- ❑ Least significant 2 bits of address identify which 16 bit word is required from 32 bit data block
- ❑ e.g.

❑ Address	Tag	Data	Cache line
❑ FFFFFC	FFFFFC	24682468	3FFF


Set Associative Mapping

- ❑ Cache is divided into a number of sets
- ❑ Each set contains a number of lines
- ❑ A given block maps to any line in a given set
 - ❑ e.g. Block B can be in any line of set i
- ❑ e.g. 2 lines per set
 - ❑ 2 way associative mapping
 - ❑ A given block can be in one of 2 lines in only one set

Set Associative Mapping Example

- 13 bit set number
- Block number in main memory is modulo 2^{13}
- 000000, 00A000, 00B000, 00C000 ... map to same set

Two Way Set Associative Cache Organization




Set Associative Mapping Address Structure

Tag 9 bit	Set 13 bit	Word 2 bit
-----------	------------	---------------

- Use set field to determine cache set to look in
- Compare tag field to see if we have a hit
- e.g

□ Address	Tag	Data	Set number
□ 1FF 7FFC	1FF	12345678	1FFF
□ 001 7FFC	001	11223344	1FFF

Two Way Set Associative Mapping Example



Replacement Algorithms (1)

Direct mapping

- ❑ No choice
- ❑ Each block only maps to one line
- ❑ Replace that line

Replacement Algorithms (2)

Associative & Set Associative

- ❑ Hardware implemented algorithm (speed)
- ❑ Least Recently used (LRU)
 - ❑ e.g. in 2 way set associative
 - ❑ Which of the 2 block is lru?
- ❑ First in first out (FIFO)
 - ❑ replace block that has been in cache longest
- ❑ Least frequently used
 - ❑ replace block which has had fewest hits
- ❑ Random

Write Policy

- ❑ Must not overwrite a cache block unless main memory is up to date
- ❑ Multiple CPUs may have individual caches
- ❑ I/O may address main memory directly

Write through

- ❑ All writes go to main memory as well as cache
 - ❑ Multiple CPUs can monitor main memory traffic to keep local (to CPU) cache up to date
 - ❑ Lots of traffic
 - ❑ Slows down writes
-
- ❑ Remember bogus write through caches!

Write back

- ❑ Updates initially made in cache only
- ❑ Update bit for cache slot is set when update occurs
- ❑ If block is to be replaced, write to main memory only if update bit is set
- ❑ Other caches get out of sync
- ❑ I/O must access main memory through cache
- ❑ N.B. 15% of memory references are writes

Pentium Cache

- ❑ Foreground reading
- ❑ Find out detail of Pentium II cache systems
- ❑ NOT just from Stallings!

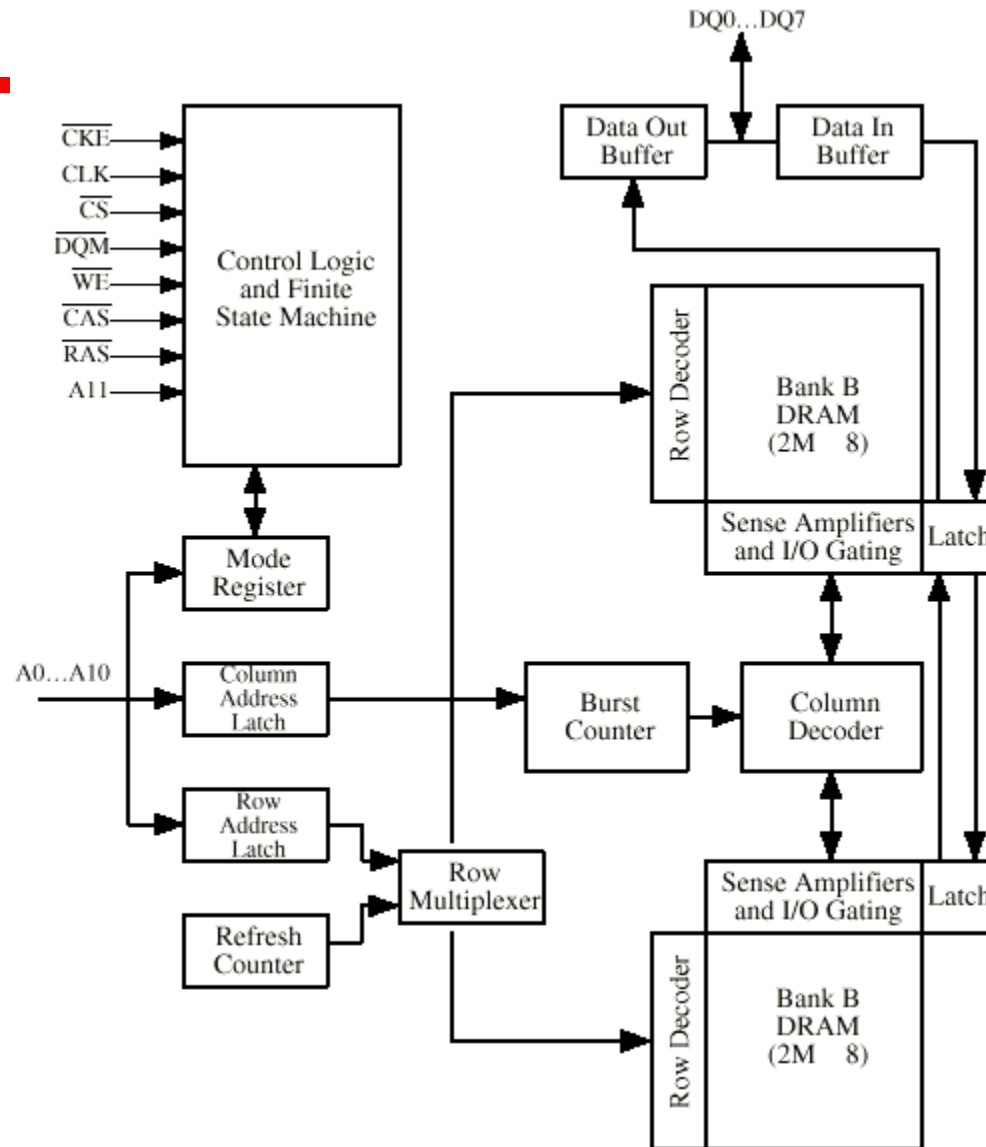
Newer RAM Technology (1)

- ❑ Basic DRAM same since first RAM chips
- ❑ Enhanced DRAM
 - ❑ Contains small SRAM as well
 - ❑ SRAM holds last line read (c.f. Cache!)
- ❑ Cache DRAM
 - ❑ Larger SRAM component
 - ❑ Use as cache or serial buffer

Newer RAM Technology (2)

- ❑ Synchronous DRAM (SDRAM)
 - ❑ currently on DIMMs
 - ❑ Access is synchronized with an external clock
 - ❑ Address is presented to RAM
 - ❑ RAM finds data (CPU waits in conventional DRAM)
 - ❑ Since SDRAM moves data in time with system clock, CPU knows when data will be ready
 - ❑ CPU does not have to wait, it can do something else
 - ❑ Burst mode allows SDRAM to set up stream of data and fire it out in block

SDRAM



Newer RAM Technology (3)

- Foreground reading
- Check out any other RAM you can find
- See Web site:
 - The RAM Guide

William Stallings

Computer Organization

and Architecture

Chapter 5

External Memory

Types of External Memory

- ❑ Magnetic Disk
 - ❑ RAID
 - ❑ Removable
- ❑ Optical
 - ❑ CD-ROM
 - ❑ CD-Writable (WORM)
 - ❑ CD-R/W
 - ❑ DVD
- ❑ Magnetic Tape

Magnetic Disk

- ❑ Metal or plastic disk coated with magnetizable material (iron oxide...rust)
- ❑ Range of packaging
 - ❑ Floppy
 - ❑ Winchester hard disk
 - ❑ Removable hard disk

Data Organization and Formatting

- ❑ Concentric rings or tracks
 - ❑ Gaps between tracks
 - ❑ Reduce gap to increase capacity
 - ❑ Same number of bits per track (variable packing density)
 - ❑ Constant angular velocity
- ❑ Tracks divided into sectors
- ❑ Minimum block size is one sector
- ❑ May have more than one sector per block

Disk Data Layout

Fixed/Movable Head Disk

- ❑ Fixed head
 - ❑ One read write head per track
 - ❑ Heads mounted on fixed ridged arm
- ❑ Movable head
 - ❑ One read write head per side
 - ❑ Mounted on a movable arm

Fixed and Movable Heads

Removable or Not

☐ Removable disk

- ☐ Can be removed from drive and replaced with another disk
- ☐ Provides unlimited storage capacity
- ☐ Easy data transfer between systems

☐ Nonremovable disk

- ☐ Permanently mounted in the drive

Floppy Disk

- ❑ 8", 5.25", 3.5"
- ❑ Small capacity
 - ❑ Up to 1.44Mbyte (2.88M never popular)
- ❑ Slow
- ❑ Universal
- ❑ Cheap

Winchester Hard Disk (1)

- ❑ Developed by IBM in Winchester (USA)
- ❑ Sealed unit
- ❑ One or more platters (disks)
- ❑ Heads fly on boundary layer of air as disk spins
- ❑ Very small head to disk gap
- ❑ Getting more robust

Winchester Hard Disk (2)

- ❑ Universal
- ❑ Cheap
- ❑ Fastest external storage
- ❑ Getting larger all the time
 - ❑ Multiple Gigabyte now usual

Removable Hard Disk

☐ ZIP

- ☐ Cheap
- ☐ Very common
- ☐ Only 100M

☐ JAZ

- ☐ Not cheap
- ☐ 1G

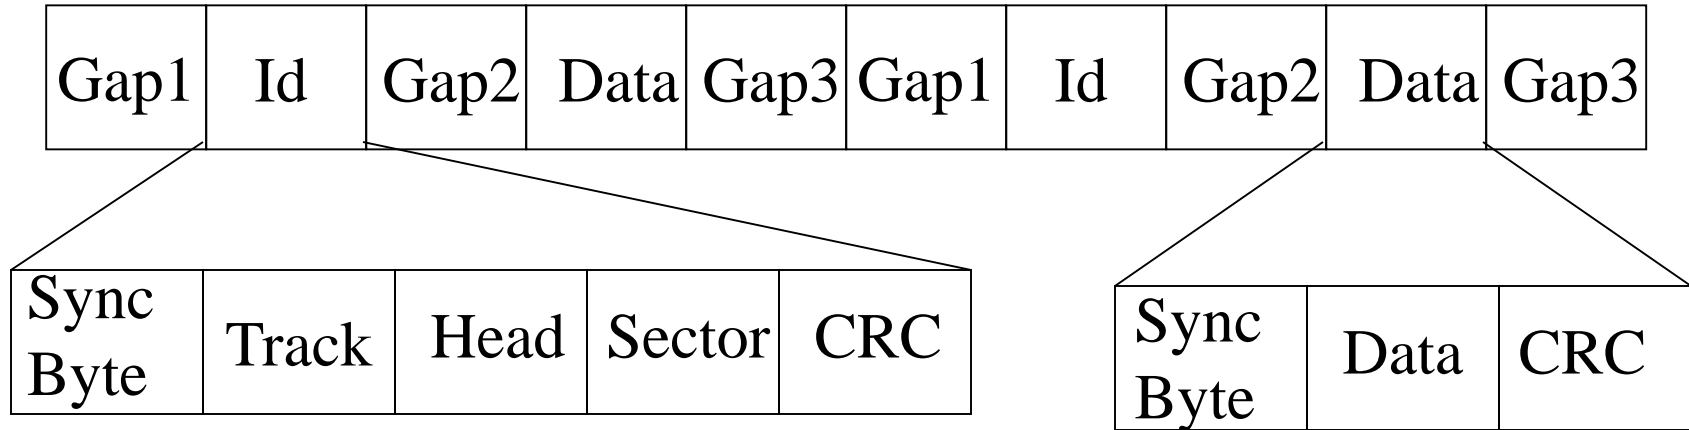
☐ L-120 (a: drive)

- ☐ Also reads 3.5" floppy
- ☐ Becoming more popular?

Finding Sectors

- ❑ Must be able to identify start of track and sector
- ❑ Format disk
 - ❑ Additional information not available to user
 - ❑ Marks tracks and sectors

ST506 format (old!)



❑ Foreground reading

❑ Find others

Characteristics

- ❑ Fixed (rare) or movable head
- ❑ Removable or fixed
- ❑ Single or double (usually) sided
- ❑ Single or multiple platter
- ❑ Head mechanism
 - ❑ Contact (Floppy)
 - ❑ Fixed gap
 - ❑ Flying (Winchester)

Multiple Platter

- ❑ One head per side
- ❑ Heads are joined and aligned
- ❑ Aligned tracks on each platter form cylinders
- ❑ Data is striped by cylinder
 - ❑ reduces head movement
 - ❑ Increases speed (transfer rate)

Speed

- Seek time
 - Moving head to correct track
- (Rotational) latency
 - Waiting for data to rotate under head
- Access time = Seek + Latency
- Transfer rate

RAID

- ❑ Redundant Array of Independent Disks
- ❑ Redundant Array of Inexpensive Disks
- ❑ 6 levels in common use
- ❑ Not a hierarchy
- ❑ Set of physical disks viewed as single logical drive by O/S
- ❑ Data distributed across physical drives
- ❑ Can use redundant capacity to store parity information

RAID 0

- ❑ No redundancy
- ❑ Data striped across all disks
- ❑ Round Robin striping
- ❑ Increase speed
 - ❑ Multiple data requests probably not on same disk
 - ❑ Disks seek in parallel
 - ❑ A set of data is likely to be striped across multiple disks

RAID 1

- ❑ Mirrored Disks
- ❑ Data is striped across disks
- ❑ 2 copies of each stripe on separate disks
- ❑ Read from either
- ❑ Write to both
- ❑ Recovery is simple
 - ❑ Swap faulty disk & re-mirror
 - ❑ No down time
- ❑ Expensive

RAID 2

- ❑ Disks are synchronized
- ❑ Very small stripes
 - ❑ Often single byte/word
- ❑ Error correction calculated across corresponding bits on disks
- ❑ Multiple parity disks store Hamming code error correction in corresponding positions
- ❑ Lots of redundancy
 - ❑ Expensive
 - ❑ Not used

RAID 3

- ❑ Similar to RAID 2
- ❑ Only one redundant disk, no matter how large the array
- ❑ Simple parity bit for each set of corresponding bits
- ❑ Data on failed drive can be reconstructed from surviving data and parity info
- ❑ Very high transfer rates

RAID 4

- ❑ Each disk operates independently
- ❑ Good for high I/O request rate
- ❑ Large stripes
- ❑ Bit by bit parity calculated across stripes on each disk
- ❑ Parity stored on parity disk

RAID 5

- ❑ Like RAID 4
- ❑ Parity striped across all disks
- ❑ Round robin allocation for parity stripe
- ❑ Avoids RAID 4 bottleneck at parity disk
- ❑ Commonly used in network servers

- ❑ N.B. DOES NOT MEAN 5 DISKS!!!!

Optical Storage CD-ROM


- ❑ Originally for audio
- ❑ 650Mbytes giving over 70 minutes audio
- ❑ Polycarbonate coated with highly reflective coat, usually aluminum
- ❑ Data stored as pits
- ❑ Read by reflecting laser
- ❑ Constant packing density
- ❑ Constant linear velocity

CD-ROM Drive Speeds

- ❑ Audio is single speed
 - ❑ Constant linear velocity
 - ❑ 1.2 ms^{-1}
 - ❑ Track (spiral) is 5.27km long
 - ❑ Gives 4391 seconds = 73.2 minutes
- ❑ Other speeds are quoted as multiples
- ❑ e.g. 24x
- ❑ The quoted figure is the maximum the drive can achieve

Comparison of Disk Layouts

CD-ROM Format



- ❑ Mode 0=blank data field
- ❑ Mode 1=2048 byte data+error correction
- ❑ Mode 2=2336 byte data

Random Access on CD-ROM

- ❑ Difficult
- ❑ Move head to rough position
- ❑ Set correct speed
- ❑ Read address
- ❑ Adjust to required location
- ❑ (Yawn!)

CD-ROM for & against

- ❑ Large capacity (?)
 - ❑ Easy to mass produce
 - ❑ Removable
 - ❑ Robust
-
- ❑ Expensive for small runs
 - ❑ Slow
 - ❑ Read only

Other Optical Storage

- ❑ CD-Writable
 - ❑ WORM
 - ❑ Now affordable
 - ❑ Compatible with CD-ROM drives
- ❑ CD-RW
 - ❑ Erasable
 - ❑ Getting cheaper
 - ❑ Mostly CD-ROM drive compatible

DVD - what's in a name?

- ❑ Digital Video Disk
 - ❑ Used to indicate a player for movies
 - ❑ Only plays video disks
- ❑ Digital Versatile Disk
 - ❑ Used to indicate a computer drive
 - ❑ Will read computer disks and play video disks
- ❑ Dogs Veritable Dinner
- ❑ Officially - nothing!!!

DVD - technology

- ❑ Multi-layer
- ❑ Very high capacity (4.7G per layer)
- ❑ Full length movie on single disk
 - ❑ Using MPEG compression
- ❑ Finally standardized (honest!)
- ❑ Movies carry regional coding
- ❑ Players only play correct region films
- ❑ Can be “fixed”

DVD - Writable

- ❑ Loads of trouble with standards
- ❑ First generation DVD drives may not read first generation DVD-W disks
- ❑ First generation DVD drives may not read CD-RW disks
- ❑ Wait for it to settle down before buying!

Foreground Reading

- Check out optical disk storage options
- Check out Mini Disk

Magnetic Tape

- ❑ Serial access
- ❑ Slow
- ❑ Very cheap
- ❑ Backup and archive

Digital Audio Tape (DAT)

- Uses rotating head (like video)
- High capacity on small tape
 - 4Gbyte uncompressed
 - 8Gbyte compressed
- Backup of PC/network servers

William Stallings

Computer Organization

and Architecture

Chapter 6

Input/Output

Input/Output Problems

- ❑ Wide variety of peripherals
 - ❑ Delivering different amounts of data
 - ❑ At different speeds
 - ❑ In different formats
- ❑ All slower than CPU and RAM
- ❑ Need I/O modules

Input/Output Module

- ❑ Interface to CPU and Memory
- ❑ Interface to one or more peripherals
- ❑ GENERIC MODEL OF I/O DIAGRAM 6.1

External Devices

- ❑ Human readable
 - ❑ Screen, printer, keyboard
- ❑ Machine readable
 - ❑ Monitoring and control
- ❑ Communication
 - ❑ Modem
 - ❑ Network Interface Card (NIC)

I/O Module Function

- ❑ Control & Timing
- ❑ CPU Communication
- ❑ Device Communication
- ❑ Data Buffering
- ❑ Error Detection

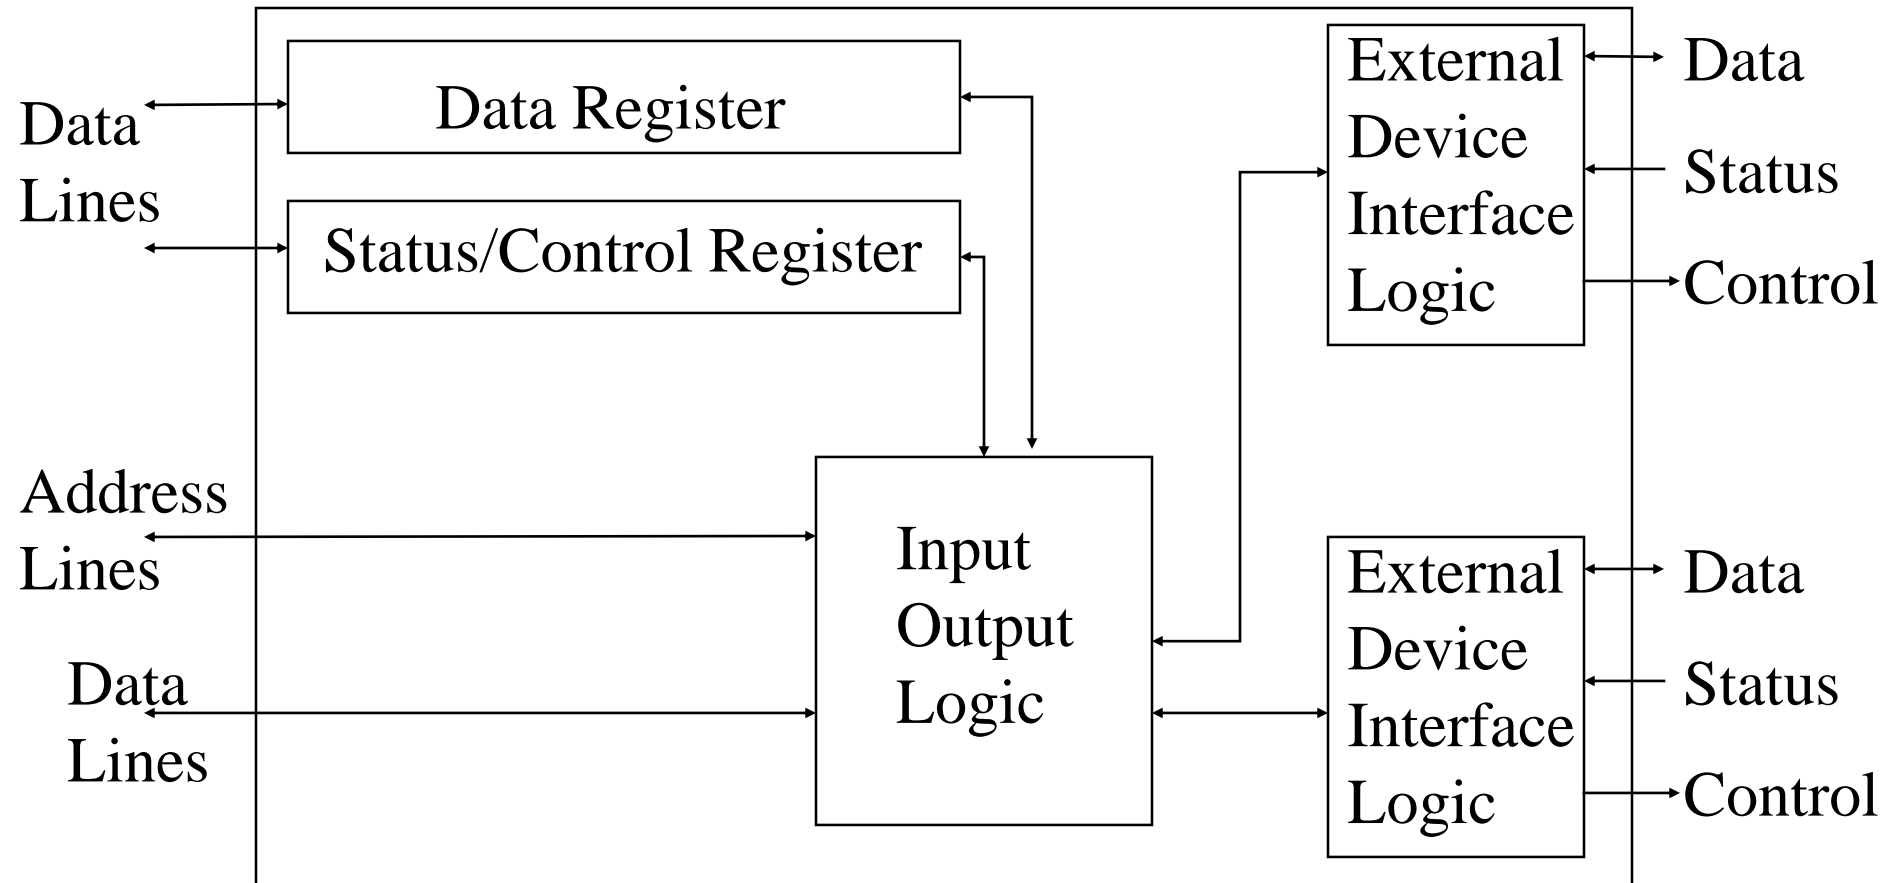
I/O Steps

- ❑ CPU checks I/O module device status
- ❑ I/O module returns status
- ❑ If ready, CPU requests data transfer
- ❑ I/O module gets data from device
- ❑ I/O module transfers data to CPU
- ❑ Variations for output, DMA, etc.

I/O Module Diagram

Systems Bus Interface

External Device Interface



I/O Module Decisions

- ❑ Hide or reveal device properties to CPU
- ❑ Support multiple or single device
- ❑ Control device functions or leave for CPU
- ❑ Also O/S decisions
 - ❑ e.g. Unix treats everything it can as a file

Input Output Techniques

- ❑ Programmed
- ❑ Interrupt driven
- ❑ Direct Memory Access (DMA)

Programmed I/O

- ❑ CPU has direct control over I/O
 - ❑ Sensing status
 - ❑ Read/write commands
 - ❑ Transferring data
- ❑ CPU waits for I/O module to complete operation
- ❑ Wastes CPU time

Programmed I/O - detail

- ❑ CPU requests I/O operation
- ❑ I/O module performs operation
- ❑ I/O module sets status bits
- ❑ CPU checks status bits periodically
- ❑ I/O module does not inform CPU directly
- ❑ I/O module does not interrupt CPU
- ❑ CPU may wait or come back later

I/O Commands

- ❑ CPU issues address
 - ❑ Identifies module (& device if >1 per module)
- ❑ CPU issues command
 - ❑ Control - telling module what to do
 - ❑ e.g. spin up disk
 - ❑ Test - check status
 - ❑ e.g. power? Error?
 - ❑ Read/Write
 - ❑ Module transfers data via buffer from/to device

Addressing I/O Devices

- ❑ Under programmed I/O data transfer is very like memory access (CPU viewpoint)
- ❑ Each device given unique identifier
- ❑ CPU commands contain identifier (address)

I/O Mapping

❑ Memory mapped I/O

- ❑ Devices and memory share an address space
- ❑ I/O looks just like memory read/write
- ❑ No special commands for I/O
 - ❑ Large selection of memory access commands available

❑ Isolated I/O

- ❑ Separate address spaces
- ❑ Need I/O or memory select lines
- ❑ Special commands for I/O
 - ❑ Limited set

Interrupt Driven I/O

- ❑ Overcomes CPU waiting
- ❑ No repeated CPU checking of device
- ❑ I/O module interrupts when ready

Interrupt Driven I/O

Basic Operation

- ❑ CPU issues read command
- ❑ I/O module gets data from peripheral whilst CPU does other work
- ❑ I/O module interrupts CPU
- ❑ CPU requests data
- ❑ I/O module transfers data

CPU Viewpoint

- ❑ Issue read command
- ❑ Do other work
- ❑ Check for interrupt at end of each instruction cycle
- ❑ If interrupted:-
 - ❑ Save context (registers)
 - ❑ Process interrupt
 - ❑ Fetch data & store
- ❑ See Operating Systems notes

Design Issues

- How do you identify the module issuing the interrupt?
- How do you deal with multiple interrupts?
 - i.e. an interrupt handler being interrupted

Identifying Interrupting Module (1)

- ❑ Different line for each module
 - ❑ PC
 - ❑ Limits number of devices
- ❑ Software poll
 - ❑ CPU asks each module in turn
 - ❑ Slow

Identifying Interrupting Module (2)

- ❑ Daisy Chain or Hardware poll
 - ❑ Interrupt Acknowledge sent down a chain
 - ❑ Module responsible places vector on bus
 - ❑ CPU uses vector to identify handler routine
- ❑ Bus Master
 - ❑ Module must claim the bus before it can raise interrupt
 - ❑ e.g. PCI & SCSI

Multiple Interrupts

- ❑ Each interrupt line has a priority
- ❑ Higher priority lines can interrupt lower priority lines
- ❑ If bus mastering only current master can interrupt


Example - PC Bus

- ❑ 80x86 has one interrupt line
- ❑ 8086 based systems use one 8259A interrupt controller
- ❑ 8259A has 8 interrupt lines

Sequence of Events

- ❑ 8259A accepts interrupts
- ❑ 8259A determines priority
- ❑ 8259A signals 8086 (raises INTR line)
- ❑ CPU Acknowledges
- ❑ 8259A puts correct vector on data bus
- ❑ CPU processes interrupt


PC Interrupt Layout



ISA Bus Interrupt System

- ❑ ISA bus chains two 8259As together
- ❑ Link is via interrupt 2
- ❑ Gives 15 lines
 - ❑ 16 lines less one for link
- ❑ IRQ 9 is used to re-route anything trying to use IRQ 2
 - ❑ Backwards compatibility
- ❑ Incorporated in chip set

ISA Interrupt Layout



Foreground Reading

- ❑ <http://www.pcguide.com/ref/mbsys/res/irq/func.htm>
- ❑ In fact look at <http://www.pcguide.com/>

Direct Memory Access

- ❑ Interrupt driven and programmed I/O require active CPU intervention
 - ❑ Transfer rate is limited
 - ❑ CPU is tied up
- ❑ DMA is the answer

DMA Function

- Additional Module (hardware) on bus
- DMA controller takes over from CPU for I/O

DMA Operation

- ❑ CPU tells DMA controller:-
 - ❑ Read/Write
 - ❑ Device address
 - ❑ Starting address of memory block for data
 - ❑ Amount of data to be transferred
- ❑ CPU carries on with other work
- ❑ DMA controller deals with transfer
- ❑ DMA controller sends interrupt when finished

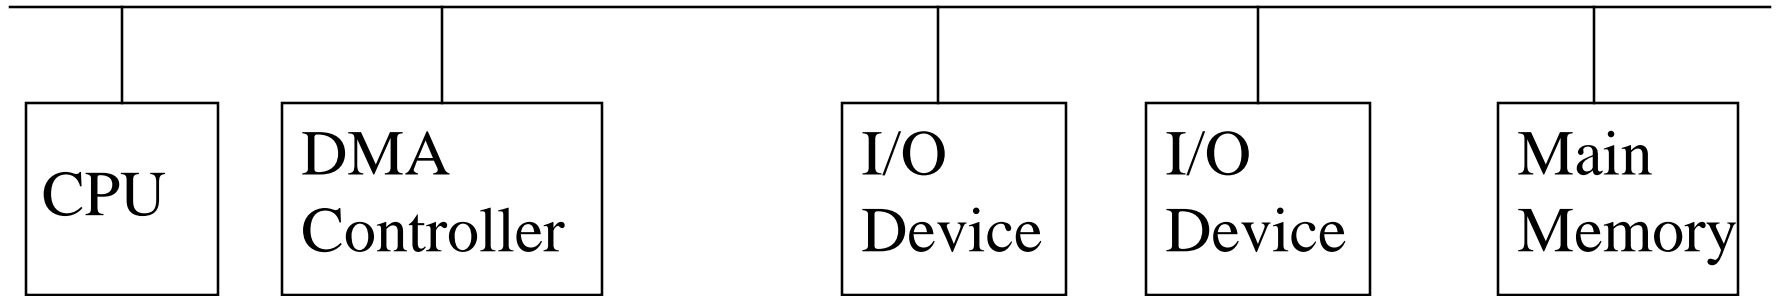
DMA Transfer Cycle Stealing

- ❑ DMA controller takes over bus for a cycle
- ❑ Transfer of one word of data
- ❑ Not an interrupt
 - ❑ CPU does not switch context
- ❑ CPU suspended just before it accesses bus
 - ❑ i.e. before an operand or data fetch or a data write
- ❑ Slows down CPU but not as much as CPU doing transfer

Aside

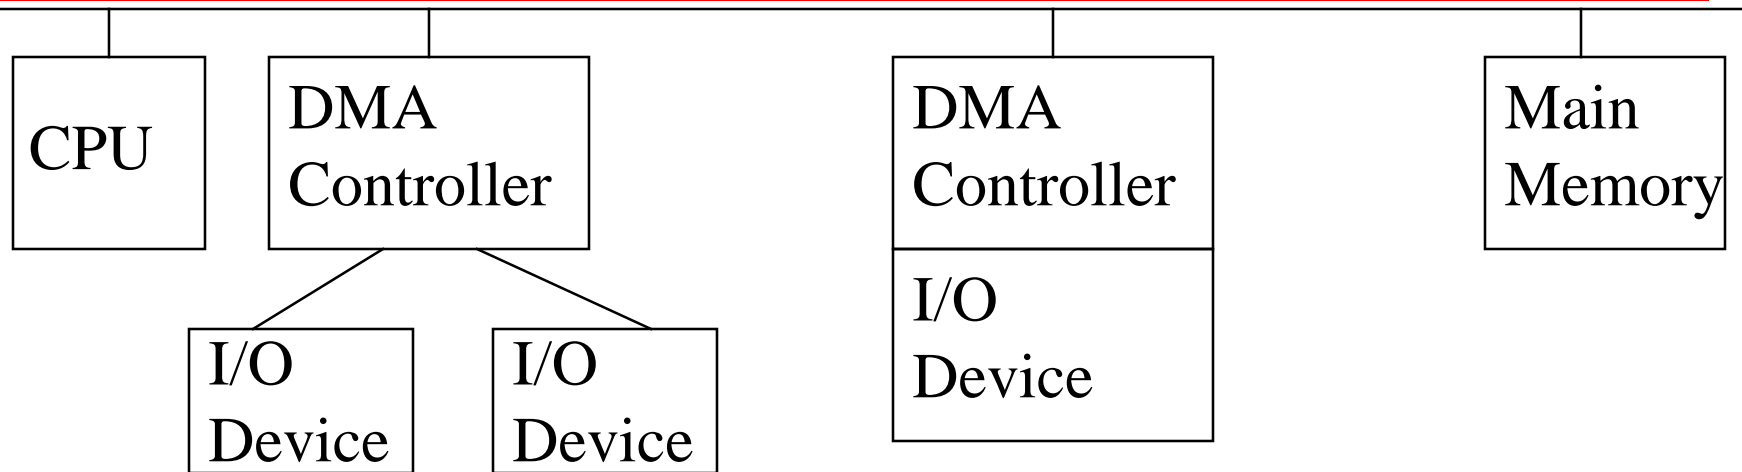
- What effect does caching memory have on DMA?
- Hint: how much are the system buses available?

DMA Configurations (1)



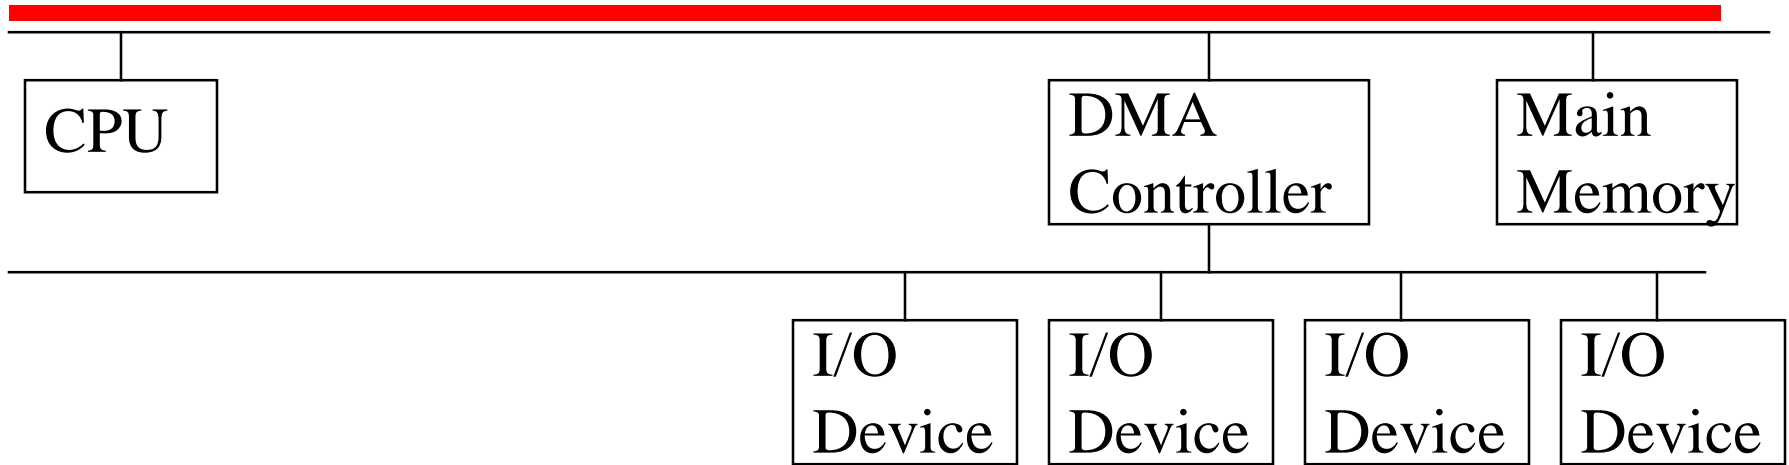
- ❑ Single Bus, Detached DMA controller
- ❑ Each transfer uses bus twice
 - ❑ I/O to DMA then DMA to memory
- ❑ CPU is suspended twice

DMA Configurations (2)



- ❑ Single Bus, Integrated DMA controller
- ❑ Controller may support >1 device
- ❑ Each transfer uses bus once
 - ❑ DMA to memory
- ❑ CPU is suspended once

DMA Configurations (3)



- ❑ Separate I/O Bus
- ❑ Bus supports all DMA enabled devices
- ❑ Each transfer uses bus once
 - ❑ DMA to memory
- ❑ CPU is suspended once

I/O Channels

- ❑ I/O devices getting more sophisticated
- ❑ e.g. 3D graphics cards
- ❑ CPU instructs I/O controller to do transfer
- ❑ I/O controller does entire transfer
- ❑ Improves speed
 - ❑ Takes load off CPU
 - ❑ Dedicated processor is faster

Interfacing

- ❑ Connecting devices together
- ❑ Bit of wire?
- ❑ Dedicated processor/memory/buses?
- ❑ E.g. SCSI, FireWire

Small Computer Systems Interface (SCSI)

- ❑ Parallel interface
- ❑ 8, 16, 32 bit data lines
- ❑ Daisy chained
- ❑ Devices are independent
- ❑ Devices can communicate with each other as well as host

SCSI - 1

- Early 1980s
- 8 bit
- 5MHz
- Data rate 5MBytes.s⁻¹
- Seven devices
 - Eight including host interface

SCSI - 2

- 1991
- 16 and 32 bit
- 10MHz
- Data rate 20 or 40 Mbytes.s⁻¹
- (Check out Ultra/Wide SCSI)

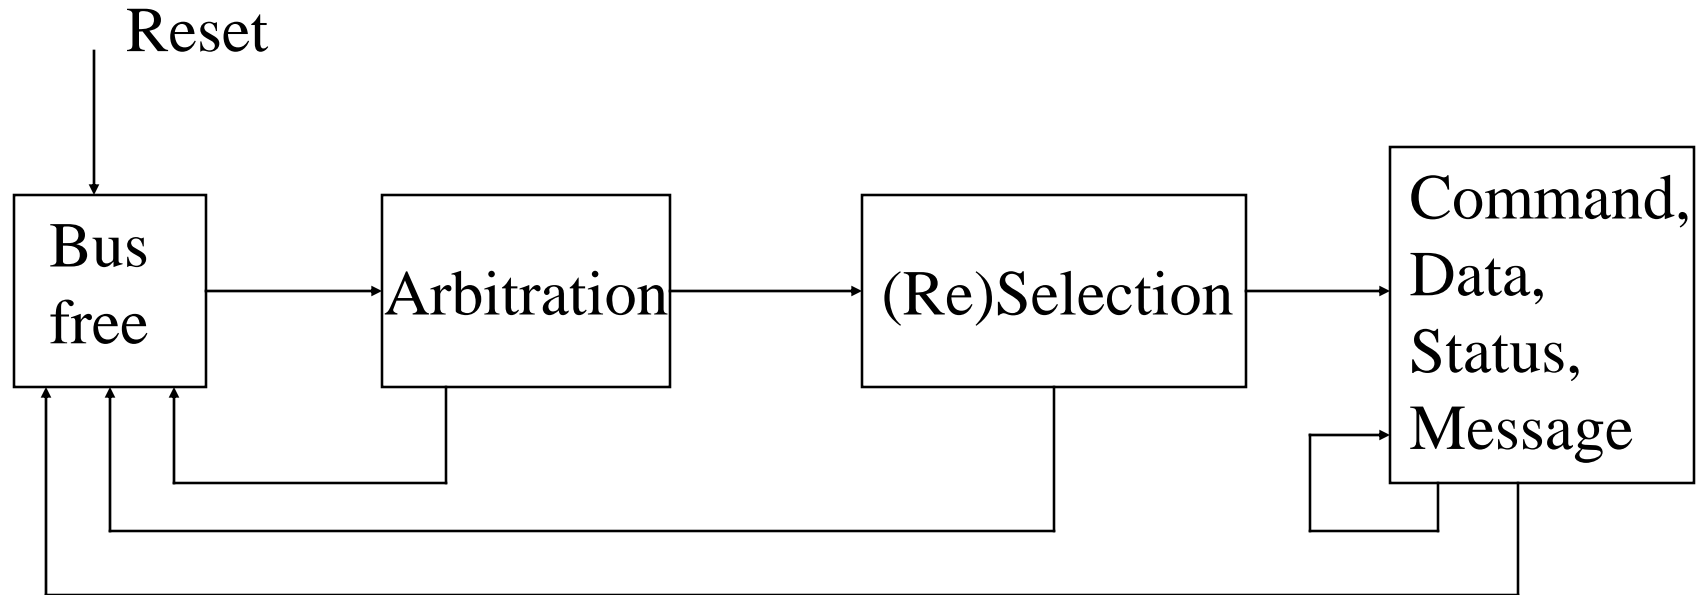
SCSI Signaling (1)

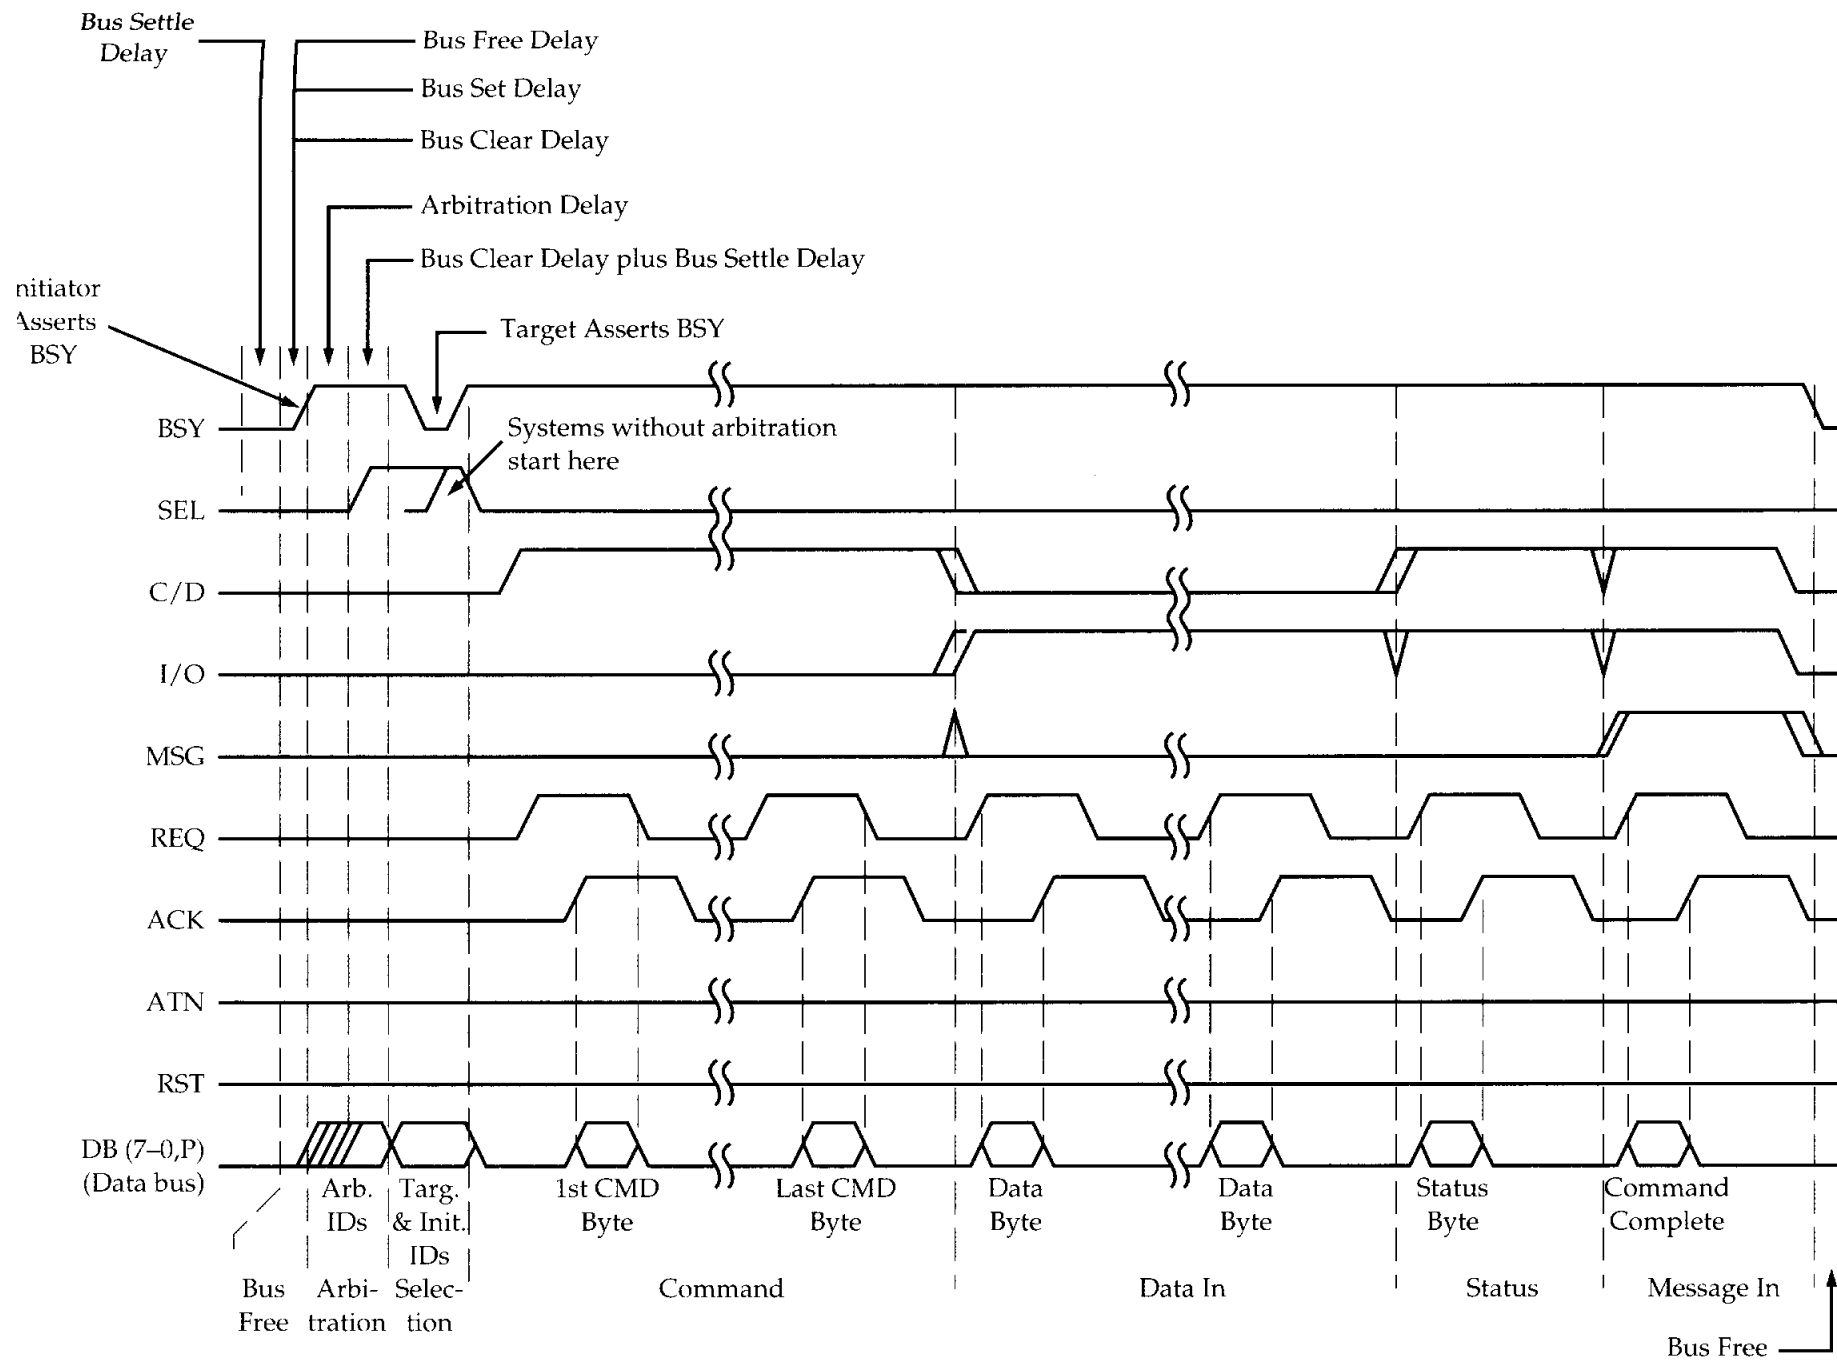
- ❑ Between initiator and target
 - ❑ Usually host & device
- ❑ Bus free? (c.f. Ethernet)
- ❑ Arbitration - take control of bus (c.f. PCI)
- ❑ Select target
- ❑ Reselection
 - ❑ Allows reconnection after suspension
 - ❑ e.g. if request takes time to execute, bus can be released

SCSI Signaling (2)

- ❑ Command - target requesting from initiator
- ❑ Data request
- ❑ Status request
- ❑ Message request (both ways)

SCSI Bus Phases





Configuring SCSI

- ❑ Bus must be terminated at each end
 - ❑ Usually one end is host adapter
 - ❑ Plug in terminator or switch(es)
- ❑ SCSI Id must be set
 - ❑ Jumpers or switches
 - ❑ Unique on chain
 - ❑ 0 (zero) for boot device
 - ❑ Higher number is higher priority in arbitration

IEEE 1394 FireWire

- ❑ High performance serial bus
- ❑ Fast
- ❑ Low cost
- ❑ Easy to implement
- ❑ Also being used in digital cameras, VCRs and TV

FireWire Configuration

- ❑ Daisy chain
- ❑ Up to 63 devices on single port
 - ❑ Really 64 of which one is the interface itself
- ❑ Up to 1022 buses can be connected with bridges
- ❑ Automatic configuration
- ❑ No bus terminators
- ❑ May be tree structure

FireWire v SCSI

FireWire 3 Layer Stack

- Physical

- Transmission medium, electrical and signaling characteristics

- Link

- Transmission of data in packets

- Transaction

- Request-response protocol

FireWire - Physical Layer

- ❑ Data rates from 25 to 400Mbps
- ❑ Two forms of arbitration
 - ❑ Based on tree structure
 - ❑ Root acts as arbiter
 - ❑ First come first served
 - ❑ Natural priority controls simultaneous requests
 - ❑ i.e. who is nearest to root
 - ❑ Fair arbitration
 - ❑ Urgent arbitration

FireWire - Link Layer

- ❑ Two transmission types
 - ❑ Asynchronous
 - ❑ Variable amount of data and several bytes of transaction data transferred as a packet
 - ❑ To explicit address
 - ❑ Acknowledgement returned
 - ❑ Isochronous
 - ❑ Variable amount of data in sequence of fixed size packets at regular intervals
 - ❑ Simplified addressing
 - ❑ No acknowledgement

FireWire Subactions

Foreground Reading

- Check out Universal Serial Bus (USB)
- Compare with other communication standards
e.g. Ethernet

William Stallings

Computer Organization and Architecture

Chapter 7

Operating System Support

Objectives and Functions


- Convenience

- Making the computer easier to use

- Efficiency

- Allowing better use of computer resources

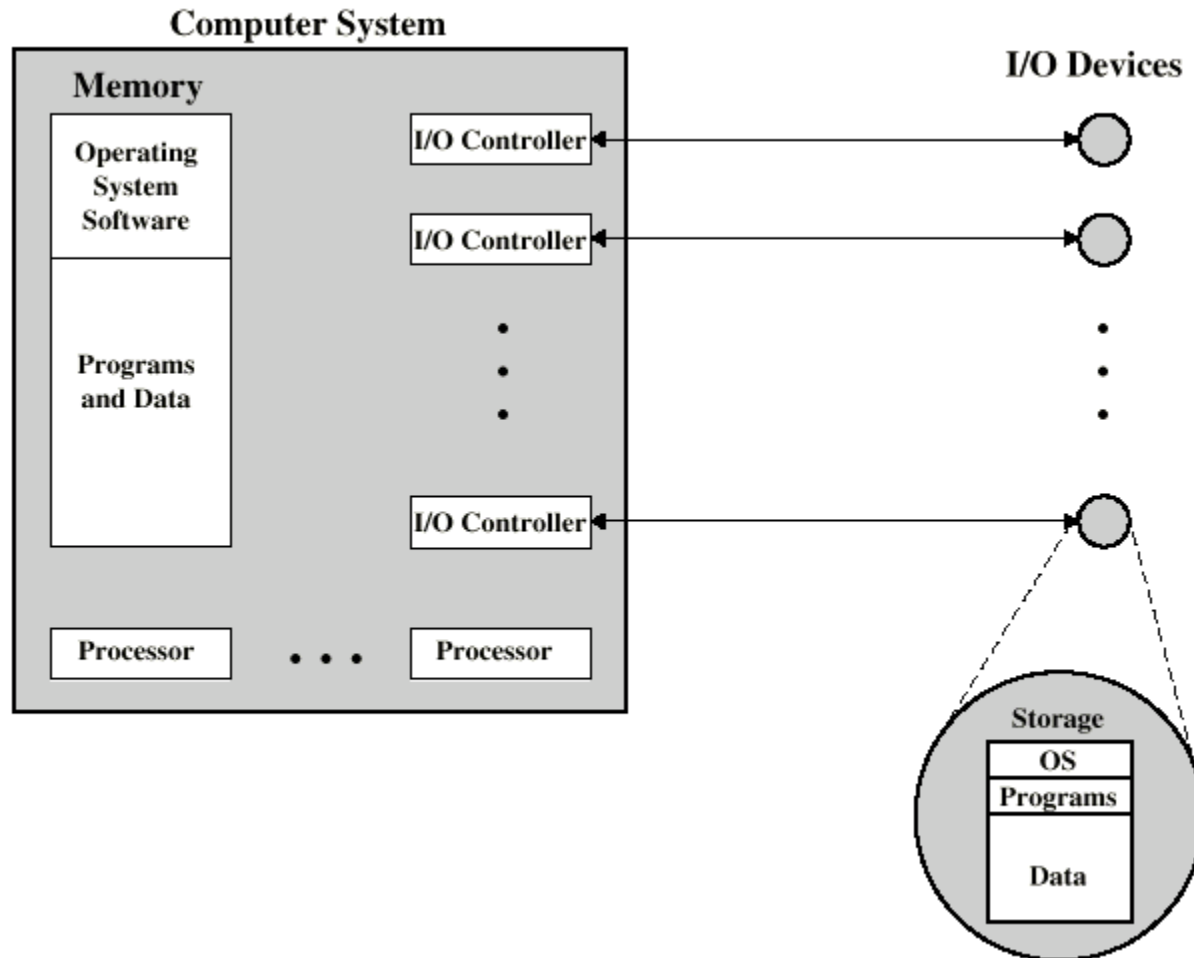
Layers and Views of a Computer System



Operating System Services

- ❑ Program creation
- ❑ Program execution
- ❑ Access to I/O devices
- ❑ Controlled access to files
- ❑ System access
- ❑ Error detection and response
- ❑ Accounting

O/S as a Resource Manager



Types of Operating System

- ❑ Interactive
- ❑ Batch
- ❑ Single program (Uni-programming)
- ❑ Multi-programming (Multi-tasking)

Early Systems

- ❑ Late 1940s to mid 1950s
- ❑ No Operating System
- ❑ Programs interact directly with hardware
- ❑ Two main problems:
 - ❑ Scheduling
 - ❑ Setup time

Simple Batch Systems

- ❑ Resident Monitor program
- ❑ Users submit jobs to operator
- ❑ Operator batches jobs
- ❑ Monitor controls sequence of events to process batch
- ❑ When one job is finished, control returns to Monitor which reads next job
- ❑ Monitor handles scheduling

Job Control Language

- ❑ Instructions to Monitor
- ❑ Usually denoted by \$
- ❑ e.g.
 - ❑ \$JOB
 - ❑ \$FTN
 - ❑ ... Some Fortran instructions
 - ❑ \$LOAD
 - ❑ \$RUN
 - ❑ ... Some data
 - ❑ \$END


Desirable Hardware Features

- ❑ Memory protection
 - ❑ To protect the Monitor
- ❑ Timer
 - ❑ To prevent a job monopolizing the system
- ❑ Privileged instructions
 - ❑ Only executed by Monitor
 - ❑ e.g. I/O
- ❑ Interrupts
 - ❑ Allows for relinquishing and regaining control

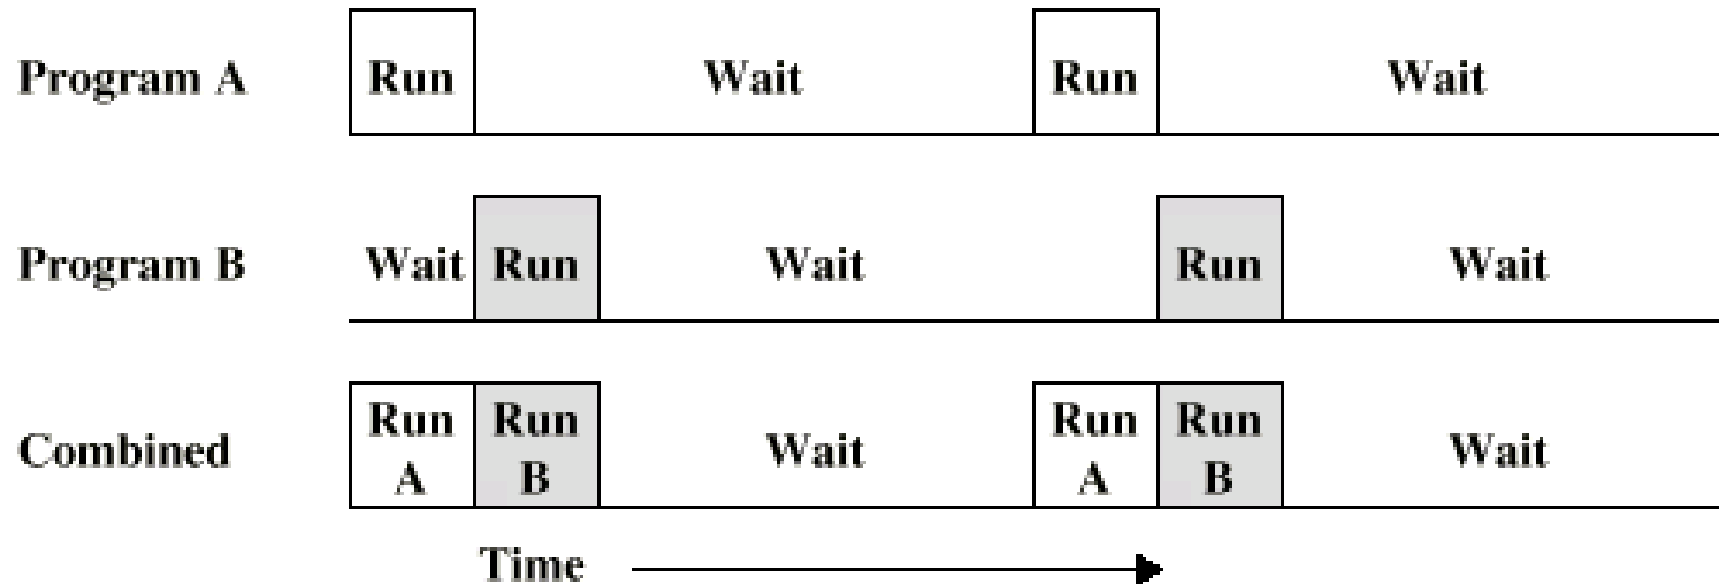
Multi-programmed Batch Systems

- ❑ I/O devices very slow
- ❑ When one program is waiting for I/O, another can use the CPU

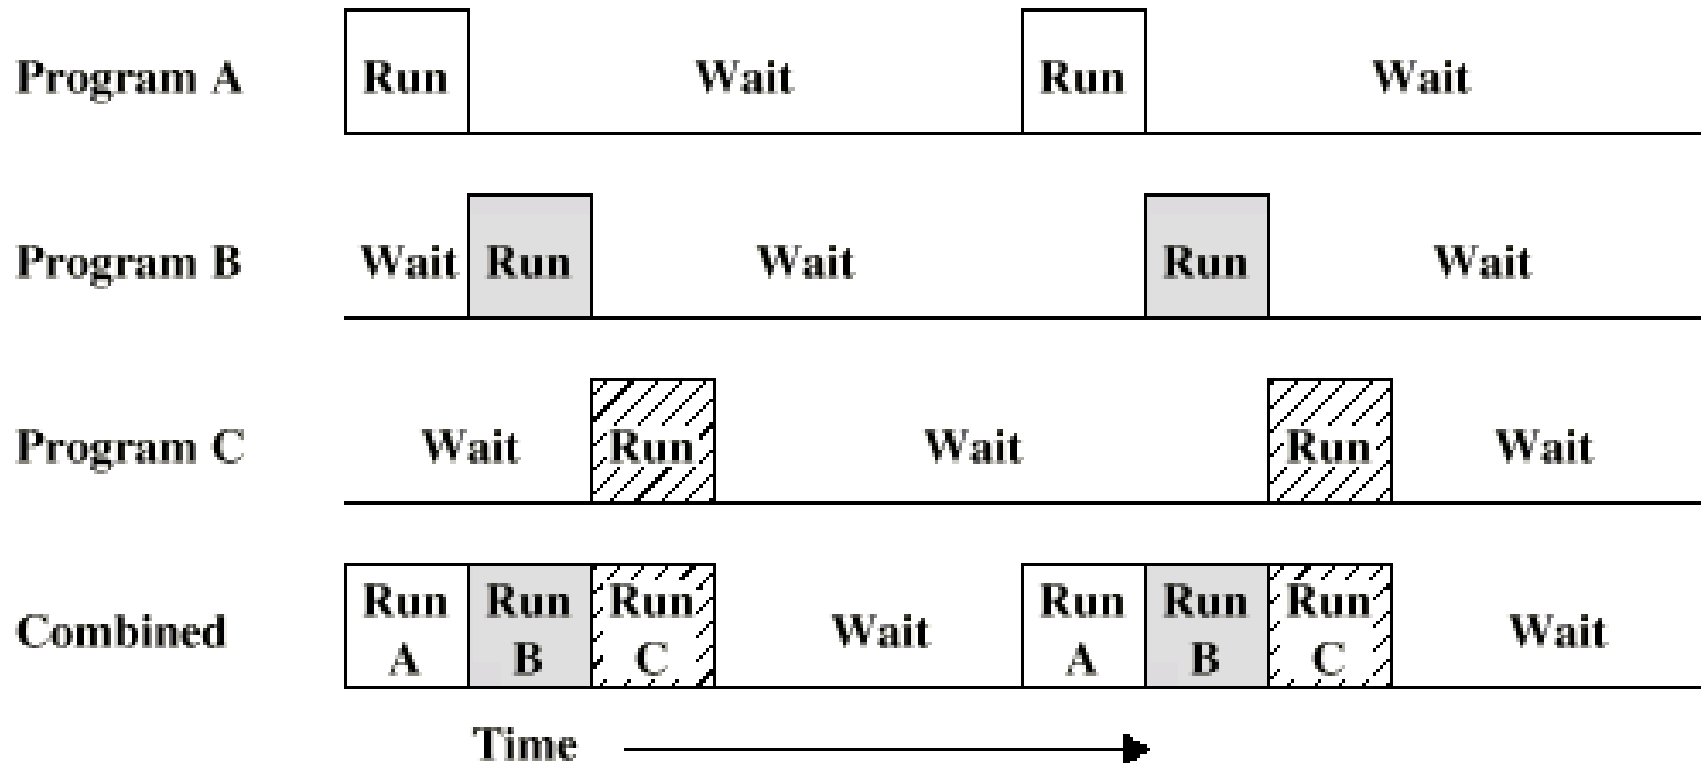
Single Program



Multi-Programming with Two Programs



Multi-Programming with Three Programs



Time Sharing Systems

- Allow users to interact directly with the computer
 - i.e. Interactive
- Multi-programming allows a number of users to interact with the computer

Scheduling

- ❑ Key to multi-programming
- ❑ Long term
- ❑ Medium term
- ❑ Short term
- ❑ I/O

Long Term Scheduling

- ❑ Determines which programs are submitted for processing
- ❑ i.e. controls the degree of multi-programming
- ❑ Once submitted, a job becomes a process for the short term scheduler
- ❑ (or it becomes a swapped out job for the medium term scheduler)


Medium Term Scheduling

- ❑ Part of the swapping function (later...)
- ❑ Usually based on the need to manage multi-programming
- ❑ If no virtual memory, memory management is also an issue

Short Term Scheduler

- ❑ Dispatcher
- ❑ Fine grained decisions of which job to execute next
- ❑ i.e. which job actually gets to use the processor in the next time slot

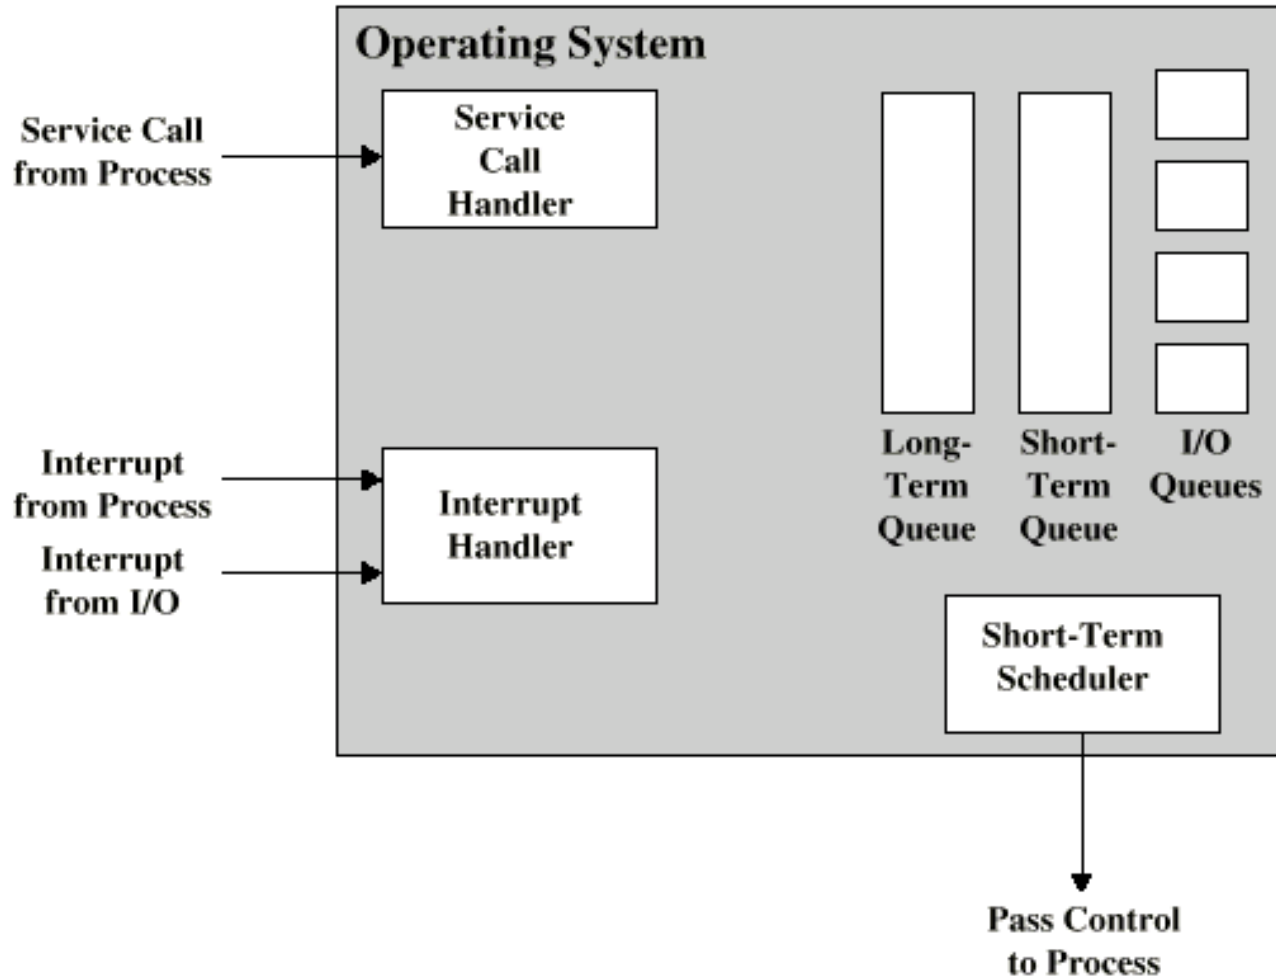
Process States



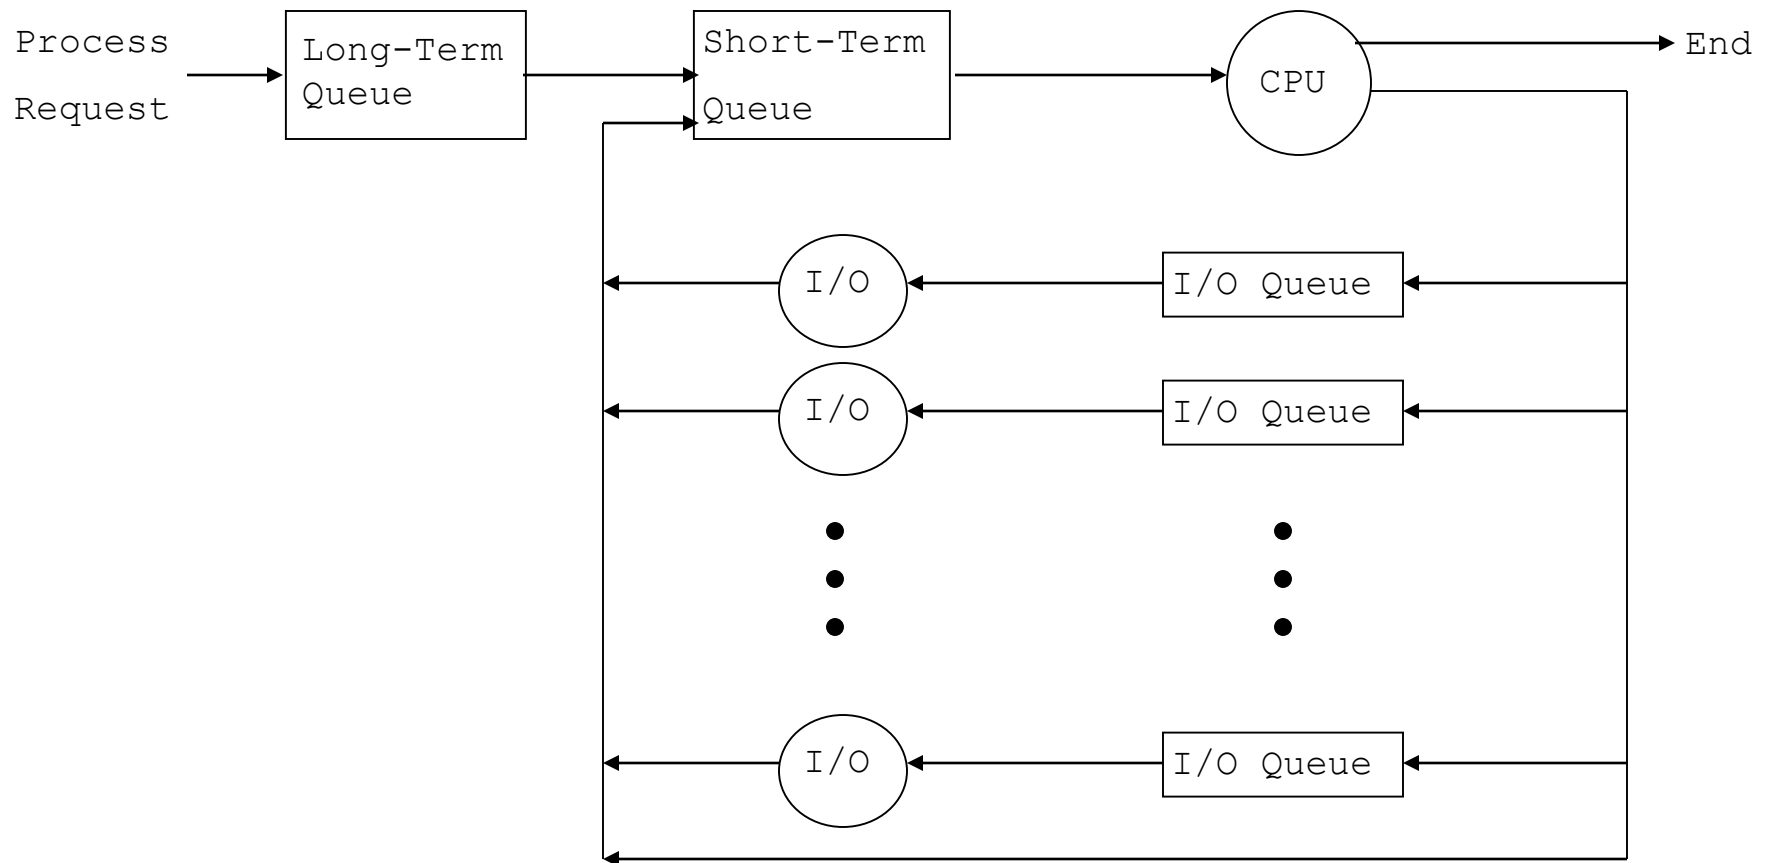
Process Control Block

- ❑ Identifier
- ❑ State
- ❑ Priority
- ❑ Program counter
- ❑ Memory pointers
- ❑ Context data
- ❑ I/O status
- ❑ Accounting information

Key Elements of O/S



Process Scheduling



Memory Management

- Uni-program

- Memory split into two
 - One for Operating System (monitor)
 - One for currently executing program

- Multi-program

- “User” part is sub-divided and shared among active processes

Swapping

- ❑ Problem: I/O is so slow compared with CPU that even in multi-programming system, CPU can be idle most of the time
- ❑ Solutions:
 - ❑ Increase main memory
 - ❑ Expensive
 - ❑ Leads to larger programs
 - ❑ Swapping


What is Swapping?

- ❑ Long term queue of processes stored on disk
- ❑ Processes “swapped” in as space becomes available
- ❑ As a process completes it is moved out of main memory
- ❑ If none of the processes in memory are ready (i.e. all I/O blocked)
 - ❑ Swap out a blocked process to intermediate queue
 - ❑ Swap in a ready process or a new process
 - ❑ But swapping is an I/O process...


Partitioning

- ❑ Splitting memory into sections to allocate to processes (including Operating System)
- ❑ Fixed-sized partitions
 - ❑ May not be equal size
 - ❑ Process is fitted into smallest hole that will take it (best fit)
 - ❑ Some wasted memory
 - ❑ Leads to variable sized partitions

Fixed Partitioning



(a) Equal-size partitions



(b) Unequal-size partitions


Variable Sized Partitions (1)

- ❑ Allocate exactly the required memory to a process
- ❑ This leads to a hole at the end of memory, too small to use
 - ❑ Only one small hole - less waste
- ❑ When all processes are blocked, swap out a process and bring in another
- ❑ New process may be smaller than swapped out process
- ❑ Another hole


Variable Sized Partitions (2)

- Eventually have lots of holes (fragmentation)
- Solutions:
 - Coalesce - Join adjacent holes into one large hole
 - Compaction - From time to time go through memory and move all hole into one free block (c.f. disk de-fragmentation)

Effect of Dynamic Partitioning




(a)




(b)




(c)




(d)




(e)



(f)



(g)



(h)

Relocation


- ❑ No guarantee that process will load into the same place in memory
- ❑ Instructions contain addresses
 - ❑ Locations of data
 - ❑ Addresses for instructions (branching)
- ❑ Logical address - relative to beginning of program
- ❑ Physical address - actual location in memory (this time)
- ❑ Automatic conversion using base address

Paging

- ❑ Split memory into equal sized, small chunks - page frames
- ❑ Split programs (processes) into equal sized small chunks - pages
- ❑ Allocate the required number page frames to a process
- ❑ Operating System maintains list of free frames
- ❑ A process does not require contiguous page frames
- ❑ Use page table to keep track

Logical and Physical Addresses

- Paging



Virtual Memory

- ❑ Demand paging
 - ❑ Do not require all pages of a process in memory
 - ❑ Bring in pages as required
- ❑ Page fault
 - ❑ Required page is not in memory
 - ❑ Operating System must swap in required page
 - ❑ May need to swap out a page to make space
 - ❑ Select page to throw out based on recent history

Thrashing

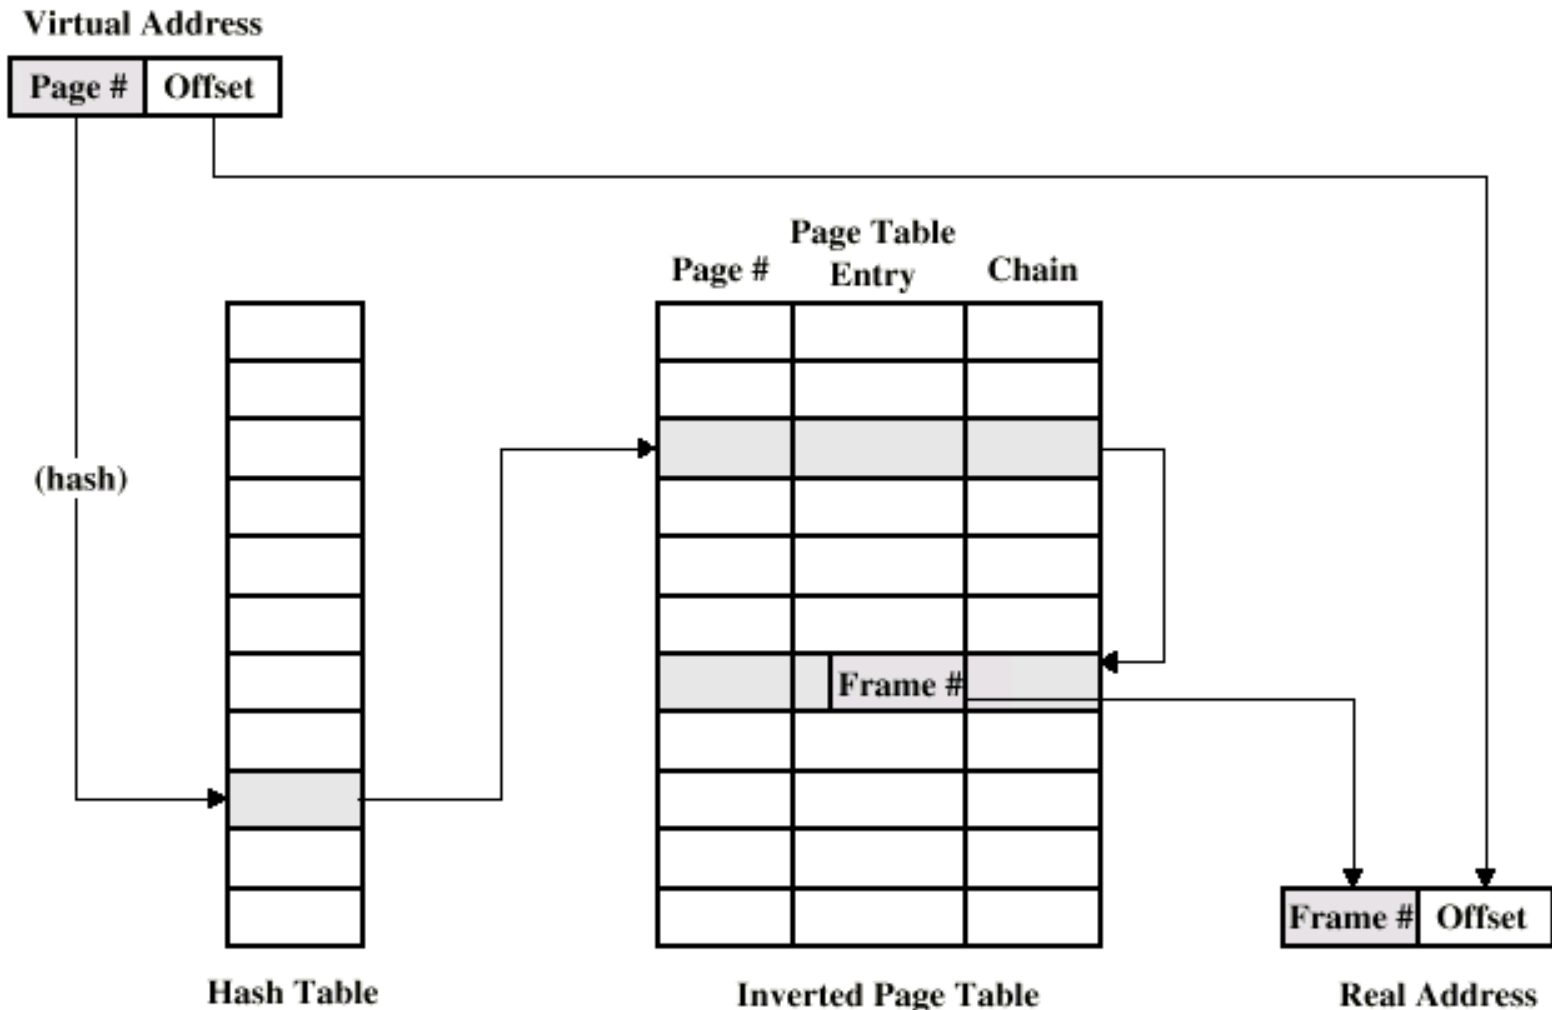
- ❑ Too many processes in too little memory
- ❑ Operating System spends all its time swapping
- ❑ Little or no real work is done
- ❑ Disk light is on all the time

- ❑ Solutions
 - ❑ Good page replacement algorithms
 - ❑ Reduce number of processes running
 - ❑ Fit more memory

Bonus

- ❑ We do not need all of a process in memory for it to run
 - ❑ We can swap in pages as required
 - ❑ So - we can now run processes that are bigger than total memory available!
-
- ❑ Main memory is called real memory
 - ❑ User/programmer sees much bigger memory - virtual memory

Page Table Structure



Segmentation

- ❑ Paging is not (usually) visible to the programmer
- ❑ Segmentation is visible to the programmer
- ❑ Usually different segments allocated to program and data
- ❑ May be a number of program and data segments

Advantages of Segmentation

- ❑ Simplifies handling of growing data structures
- ❑ Allows programs to be altered and recompiled independently, without re-linking and re-loading
- ❑ Lends itself to sharing among processes
- ❑ Lends itself to protection
- ❑ Some systems combine segmentation with paging

Required Reading

- ❑ Stallings chapter 7
- ❑ Stallings, W. Operating Systems, Internals and Design Principles, Prentice Hall 1998
- ❑ Loads of Web sites on Operating Systems

William Stallings

Computer Organization

and Architecture

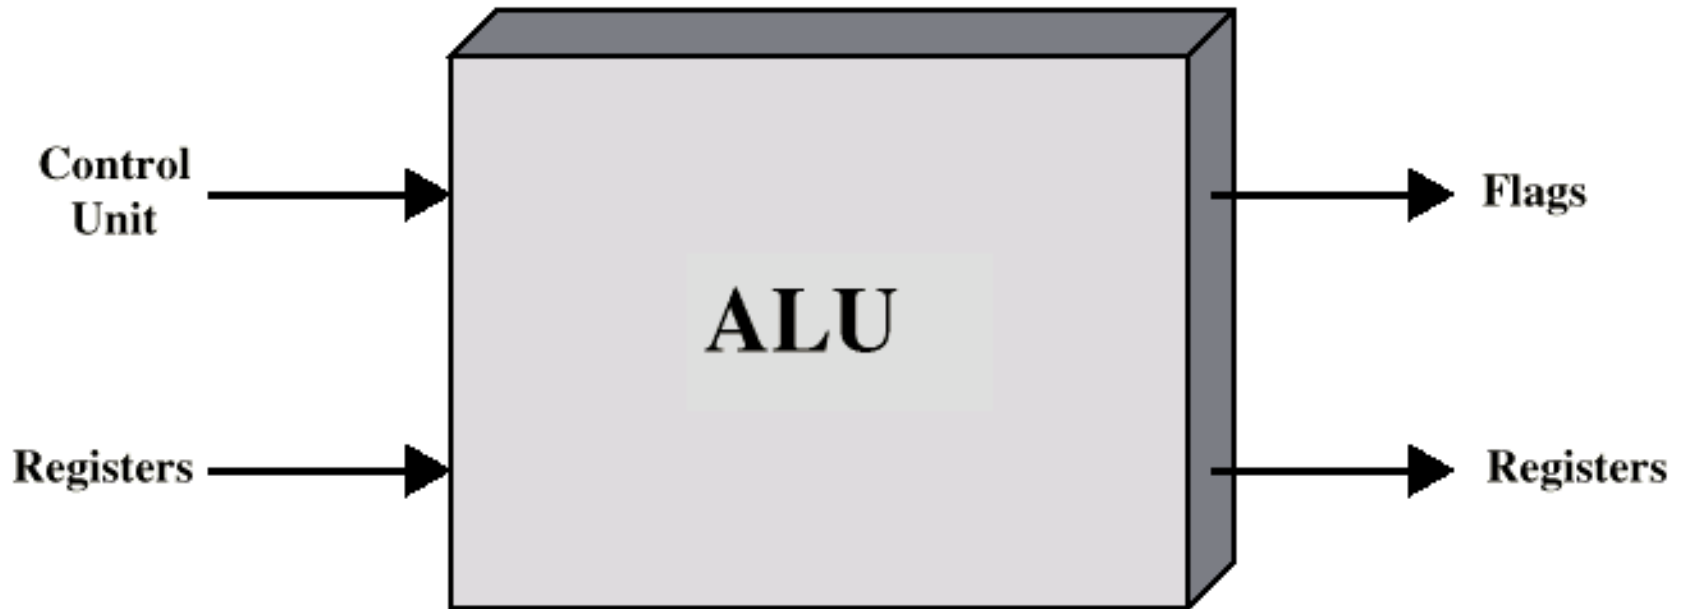
Chapter 8

Computer Arithmetic

Arithmetic & Logic Unit

- ❑ Does the calculations
- ❑ Everything else in the computer is there to service this unit
- ❑ Handles integers
- ❑ May handle floating point (real) numbers
- ❑ May be separate FPU (maths co-processor)
- ❑ May be on chip separate FPU (486DX +)

ALU Inputs and Outputs



Integer Representation

- ❑ Only have 0 & 1 to represent everything
- ❑ Positive numbers stored in binary
 - ❑ e.g. 41=00101001
- ❑ No minus sign
- ❑ No period
- ❑ Sign-Magnitude
- ❑ Two's compliment

Sign-Magnitude

- ❑ Left most bit is sign bit
- ❑ 0 means positive
- ❑ 1 means negative
- ❑ $+18 = 00010010$
- ❑ $-18 = 10010010$
- ❑ Problems
 - ❑ Need to consider both sign and magnitude in arithmetic
 - ❑ Two representations of zero (+0 and -0)

Two's Complement

$$\square +3 = 00000011$$

$$\square +2 = 00000010$$

$$\square +1 = 00000001$$

$$\square +0 = 00000000$$

$$\square -1 = 11111111$$


$$\square -2 = 11111110$$

$$\square -3 = 11111101$$


Benefits

- ❑ One representation of zero
- ❑ Arithmetic works easily (see later)
- ❑ Negating is fairly easy
 - ❑ $3 = 00000011$
 - ❑ Boolean complement gives 11111100
 - ❑ Add 1 to LSB 11111101

Geometric Depiction of Twos Complement Integers



(a) 4-bit numbers



(b) n-bit numbers

Negation Special Case 1

□ 0 = 00000000

□ Bitwise not 11111111

□ Add 1 to LSB +1

□ Result 1 00000000

□ Overflow is ignored, so:

□ - 0 = 0 ✓

Negation Special Case 2

- $-128 = 10000000$
- bitwise not 01111111
- Add 1 to LSB $+1$
- Result 10000000
- So:
- $-(-128) = -128 \quad X$
- Monitor MSB (sign bit)
- It should change during negation

Range of Numbers

□ 8 bit 2s compliment

□ $+127 = 01111111 = 2^7 - 1$

□ $-128 = 10000000 = -2^7$

□ 16 bit 2s compliment

□ $+32767 = 01111111 11111111 = 2^{15} - 1$

□ $-32768 = 10000000 00000000 = -2^{15}$

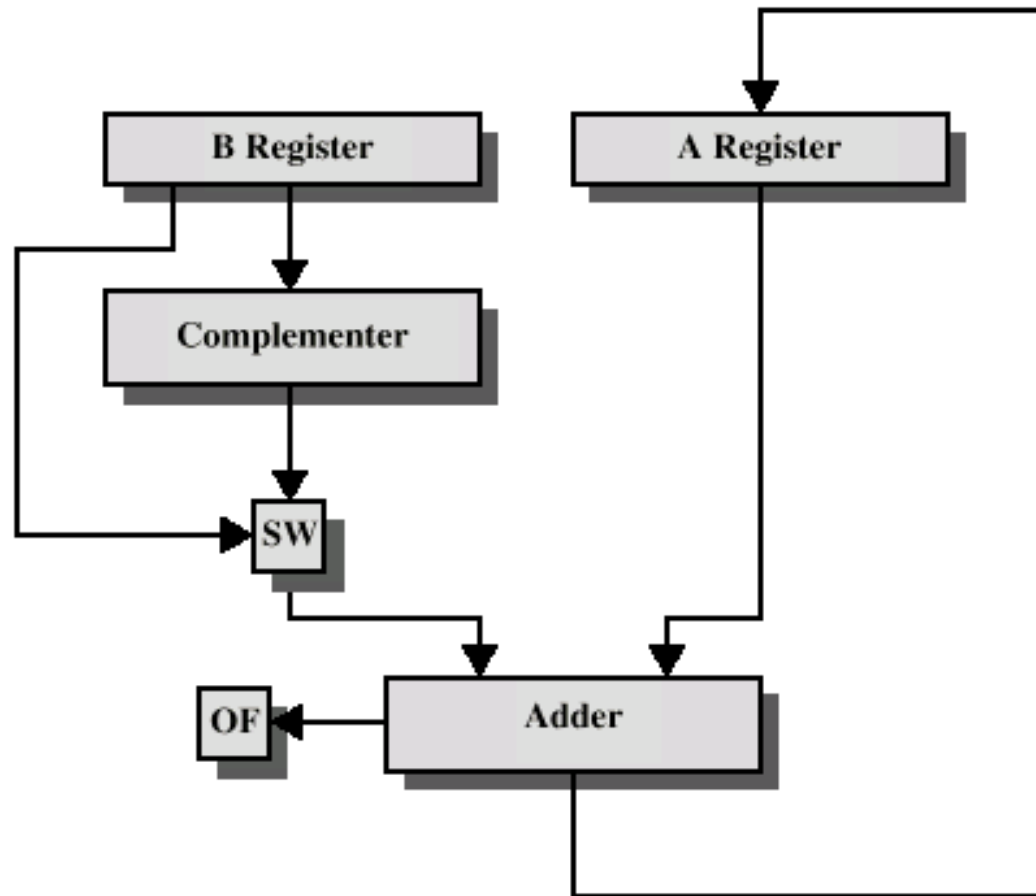
Conversion Between Lengths

- Positive number pack with leading zeros
- $+18 =$ 00010010
- $+18 =$ 00000000 00010010
- Negative numbers pack with leading ones
- $-18 =$ 10010010
- $-18 =$ 11111111 10010010
- i.e. pack with MSB (sign bit)

Addition and Subtraction

- Normal binary addition
- Monitor sign bit for overflow
- Take two's complement of subtrahend and add to minuend
 - i.e. $a - b = a + (-b)$
- So we only need addition and complement circuits

Hardware for Addition and Subtraction



OF = overflow bit

SW = Switch (select addition or subtraction)

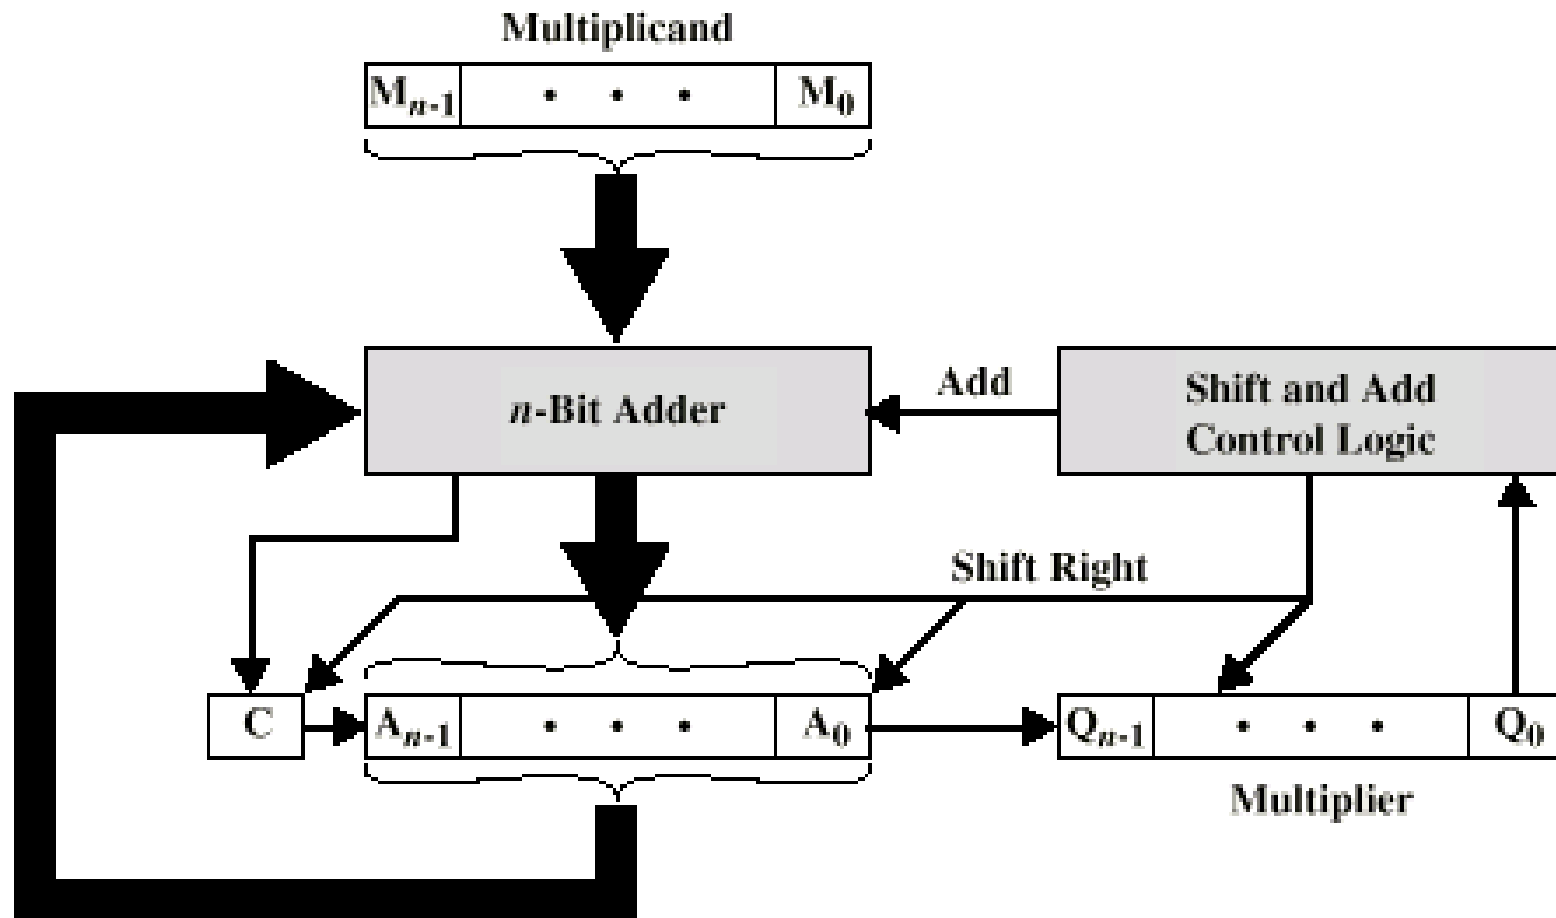
Multiplication

- ❑ Complex
- ❑ Work out partial product for each digit
- ❑ Take care with place value (column)
- ❑ Add partial products

Multiplication Example

- 1011 Multiplicand (11 dec)
- x 1101 Multiplier (13 dec)
- 1011 Partial products
- 0000 Note: if multiplier bit is 1 copy
- 1011 multiplicand (place value)
- 1011 otherwise zero
- 10001111 Product (143 dec)
- Note: need double length result

Unsigned Binary Multiplication

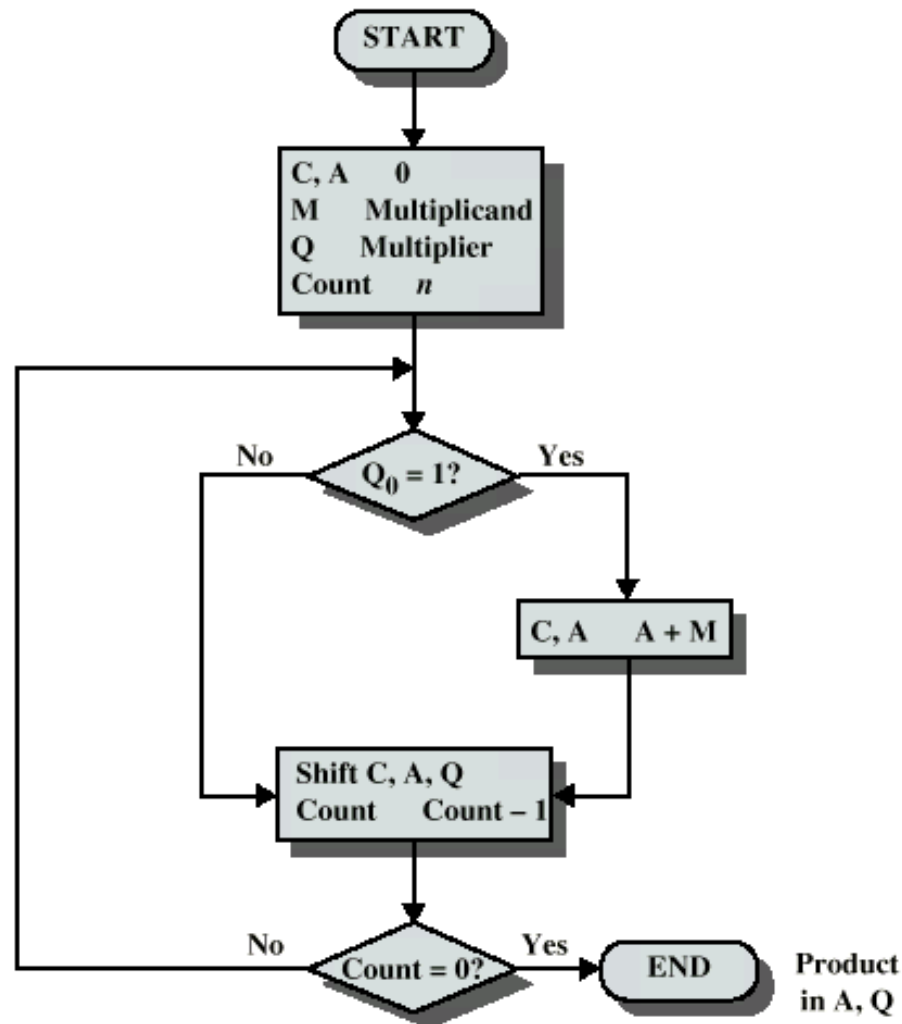


(a) Block Diagram

Execution of Example

C	A	Q	M		
0	0000	1101	1011	Initial Values	
0	1011	1101	1011	Add	} First Cycle
0	0101	1110	1011	Shift	
0	0010	1111	1011	Shift	} Second Cycle
0	1101	1111	1011	Add	
0	0110	1111	1011	Shift	} Third Cycle
1	0001	1111	1011	Add	
0	1000	1111	1011	Shift	} Fourth Cycle


Flowchart for Unsigned Binary Multiplication



Multiplying Negative Numbers

- ❑ This does not work!
- ❑ Solution 1
 - ❑ Convert to positive if required
 - ❑ Multiply as above
 - ❑ If signs were different, negate answer
- ❑ Solution 2
 - ❑ Booth's algorithm

Booth's Algorithm



Example of Booth's Algorithm

A	Q	Q ₋₁	M	Initial Values	
0000	0011	0	0111		
1001	0011	0	0111	A A - M	} First Cycle
1100	1001	1	0111	Shift	
1110	0100	1	0111	Shift	} Second Cycle
0101	0100	1	0111	A A + M	
0010	1010	0	0111	Shift	} Third Cycle
0001	0101	0	0111	Shift	
					} Fourth Cycle

Division

- ❑ More complex than multiplication
- ❑ Negative numbers are really bad!
- ❑ Based on long division

Division of Unsigned Binary Integers

Diagram illustrating the division of unsigned binary integers:

Divisor \rightarrow 1011

Dividend \rightarrow 10010011

Quotient \leftarrow 00001101

Partial Remainders \rightarrow 001110, 001111

Remainder \leftarrow 100

The diagram shows the long division process:

$$\begin{array}{r} 00001101 \\ 1011 \overline{) 10010011} \\ \underline{1011} \\ 001110 \\ \underline{1011} \\ 001111 \\ \underline{1011} \\ 100 \end{array}$$

Real Numbers


- Numbers with fractions
- Could be done in pure binary
 - $1001.1010 = 2^4 + 2^0 + 2^{-1} + 2^{-3} = 9.625$
- Where is the binary point?
- Fixed?
 - Very limited
- Moving?
 - How do you show where it is?

Floating Point

Sign bit	Biased Exponent	Significand or Mantissa
----------	-----------------	-------------------------

- ❑ $\pm \text{.significand} \times 2^{\text{exponent}}$
- ❑ Misnomer
- ❑ Point is actually fixed between sign bit and body of mantissa
- ❑ Exponent indicates place value (point position)

Floating Point Examples



(a) Format

0.11010001	2^{10100}	=	0	10010011	101000100000000000000000
-0.11010001	2^{10100}	=	1	10010011	101000100000000000000000
0.11010001	2^{-10100}	=	0	01101011	101000100000000000000000
-0.11010001	2^{-10100}	=	1	01101011	101000100000000000000000

(b) Examples

Signs for Floating Point

- ❑ Mantissa is stored in 2s compliment
- ❑ Exponent is in excess or biased notation
 - ❑ e.g. Excess (bias) 128 means
 - ❑ 8 bit exponent field
 - ❑ Pure value range 0-255
 - ❑ Subtract 128 to get correct value
 - ❑ Range -128 to +127

Normalization

- FP numbers are usually normalized
- i.e. exponent is adjusted so that leading bit (MSB) of mantissa is 1
- Since it is always 1 there is no need to store it
- (c.f. Scientific notation where numbers are normalized to give a single digit before the decimal point
- e.g. 3.123×10^3)

FP Ranges

- For a 32 bit number

 - 8 bit exponent

 - $\pm 2^{256} \approx 1.5 \times 10^{77}$


- Accuracy

 - The effect of changing lsb of mantissa


 - 23 bit mantissa $2^{-23} \approx 1.2 \times 10^{-7}$

 - About 6 decimal places

Expressible Numbers



(a) Twos Complement Integers



(b) Floating-Point Numbers

IEEE 754

- Standard for floating point storage
- 32 and 64 bit standards
- 8 and 11 bit exponent respectively
- Extended formats (both mantissa and exponent) for intermediate results

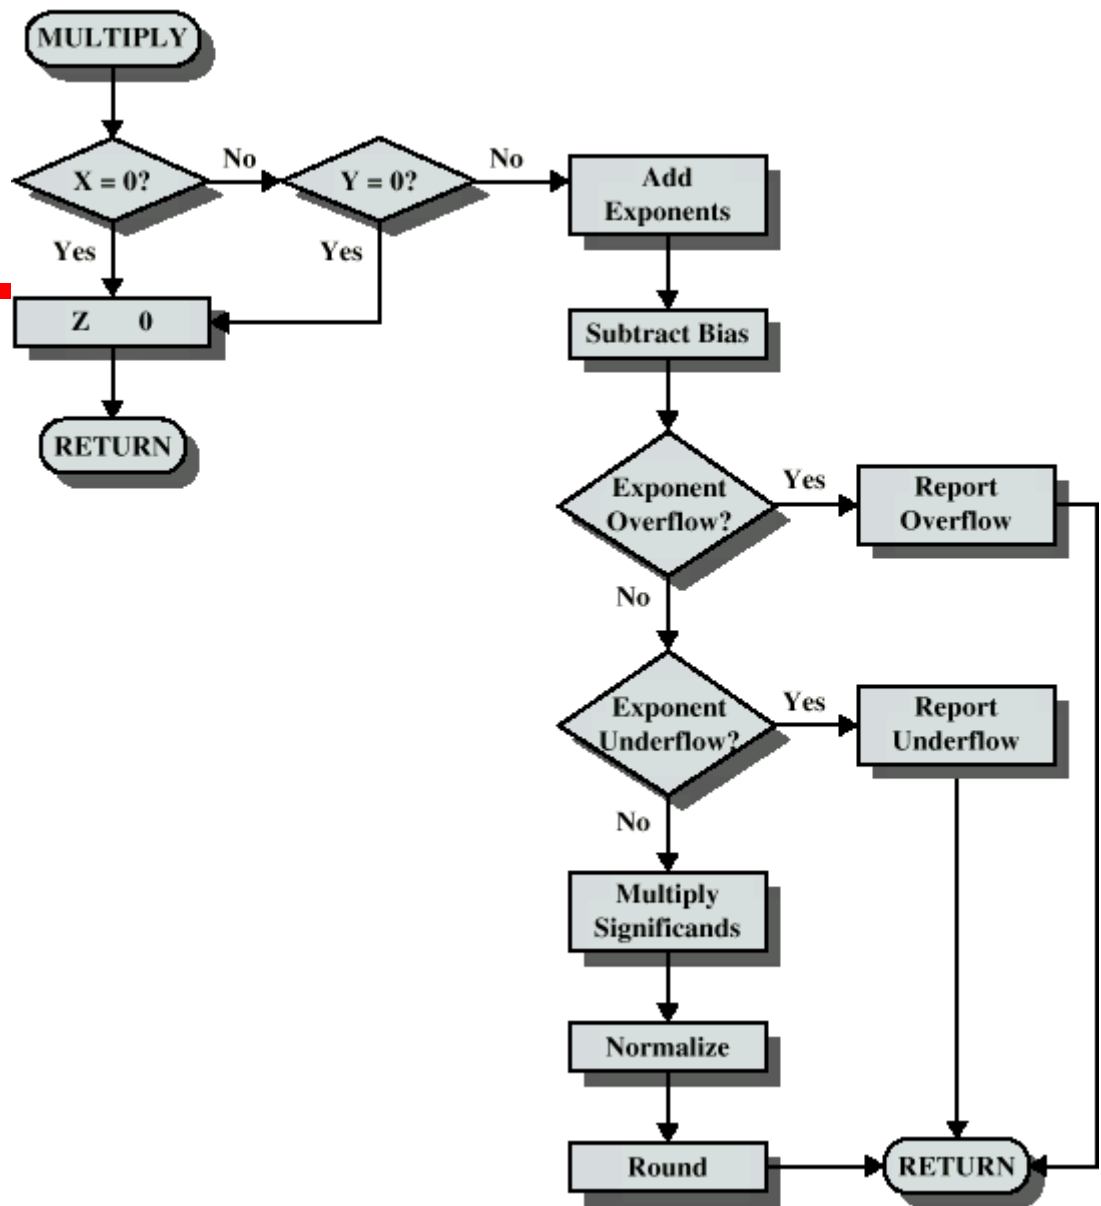
FP Arithmetic +/-

- Check for zeros
- Align significands (adjusting exponents)
- Add or subtract significands
- Normalize result

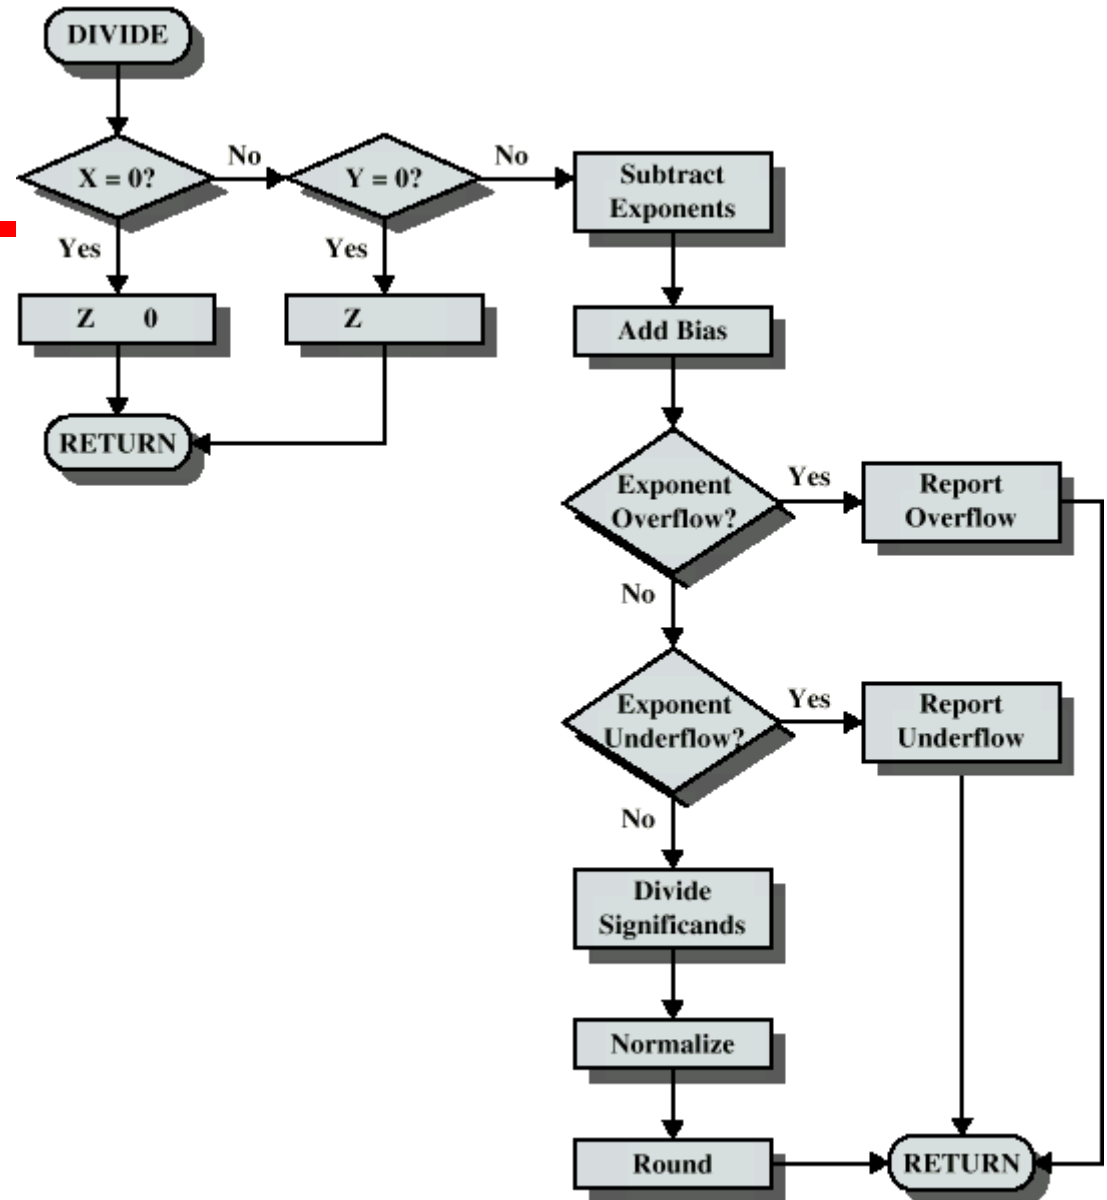
FP Arithmetic \times/\div

- ❑ Check for zero
- ❑ Add/subtract exponents
- ❑ Multiply/divide significands (watch sign)
- ❑ Normalize
- ❑ Round
- ❑ All intermediate results should be in double length storage

Floating Point Multiplication



Floating Point Division



Required Reading

- Stallings Chapter 8
- IEEE 754 on IEEE Web site

William Stallings

Computer Organization

and Architecture

Chapter 9

Instruction Sets:

Characteristics

and Functions

What is an instruction set?

- ❑ The complete collection of instructions that are understood by a CPU
- ❑ Machine Code
- ❑ Binary
- ❑ Usually represented by assembly codes

Elements of an Instruction

- ❑ Operation code (Op code)
 - ❑ Do this
- ❑ Source Operand reference
 - ❑ To this
- ❑ Result Operand reference
 - ❑ Put the answer here
- ❑ Next Instruction Reference
 - ❑ When you have done that, do this...

Where have all the Operands gone?

- ❑ Long time passing....
- ❑ (If you don't understand, you're too young!)
- ❑ Main memory (or virtual memory or cache)
- ❑ CPU register
- ❑ I/O device

Instruction Representation

- ❑ In machine code each instruction has a unique bit pattern
- ❑ For human consumption (well, programmers anyway) a symbolic representation is used
 - ❑ e.g. ADD, SUB, LOAD
- ❑ Operands can also be represented in this way
 - ❑ ADD A,B

Instruction Types

- Data processing
- Data storage (main memory)
- Data movement (I/O)
- Program flow control

Number of Addresses (a)

- 3 addresses
 - Operand 1, Operand 2, Result
 - $a = b + c;$
 - May be a forth - next instruction (usually implicit)
 - Not common
 - Needs very long words to hold everything

Number of Addresses (b)

- 2 addresses
 - One address doubles as operand and result
 - $a = a + b$
 - Reduces length of instruction
 - Requires some extra work
 - Temporary storage to hold some results

Number of Addresses (c)

- 1 address
 - Implicit second address
 - Usually a register (accumulator)
 - Common on early machines

Number of Addresses (d)

- 0 (zero) addresses
 - All addresses implicit
 - Uses a stack
 - e.g. push a
 - push b
 - add
 - pop c

- $c = a + b$

How Many Addresses

□ More addresses

- More complex (powerful?) instructions
- More registers
 - Inter-register operations are quicker
- Fewer instructions per program

□ Fewer addresses

- Less complex (powerful?) instructions
- More instructions per program
- Faster fetch/execution of instructions

Design Decisions (1)

- ❑ Operation repertoire
 - ❑ How many ops?
 - ❑ What can they do?
 - ❑ How complex are they?
- ❑ Data types
- ❑ Instruction formats
 - ❑ Length of op code field
 - ❑ Number of addresses

Design Decisions (2)

- ❑ Registers
 - ❑ Number of CPU registers available
 - ❑ Which operations can be performed on which registers?
- ❑ Addressing modes (later...)
- ❑ RISC v CISC

Types of Operand

- ❑ Addresses
- ❑ Numbers
 - ❑ Integer/floating point
- ❑ Characters
 - ❑ ASCII etc.
- ❑ Logical Data
 - ❑ Bits or flags
- ❑ (Aside: Is there any difference between numbers and characters?
Ask a C programmer!)

Pentium Data Types

- 8 bit Byte
- 16 bit word
- 32 bit double word
- 64 bit quad word
- Addressing is by 8 bit unit
- A 32 bit double word is read at addresses divisible by 4

Specific Data Types

- ❑ General - arbitrary binary contents
- ❑ Integer - single binary value
- ❑ Ordinal - unsigned integer
- ❑ Unpacked BCD - One digit per byte
- ❑ Packed BCD - 2 BCD digits per byte
- ❑ Near Pointer - 32 bit offset within segment
- ❑ Bit field
- ❑ Byte String
- ❑ Floating Point

Pentium Floating Point Data Types

□ See Stallings p324

Types of Operation

- Data Transfer
- Arithmetic
- Logical
- Conversion
- I/O
- System Control
- Transfer of Control

Data Transfer

- Specify
 - Source
 - Destination
 - Amount of data
- May be different instructions for different movements
 - e.g. IBM 370
- Or one instruction and different addresses
 - e.g. VAX

Arithmetic

- Add, Subtract, Multiply, Divide
- Signed Integer
- Floating point ?
- May include
 - Increment ($a++$)
 - Decrement ($a--$)
 - Negate ($-a$)

Logical

- Bitwise operations
- AND, OR, NOT

Conversion

□ E.g. Binary to Decimal

Input/Output

- ❑ May be specific instructions
- ❑ May be done using data movement instructions (memory mapped)
- ❑ May be done by a separate controller (DMA)

Systems Control

- ❑ Privileged instructions
- ❑ CPU needs to be in specific state
 - ❑ Ring 0 on 80386+
 - ❑ Kernel mode
- ❑ For operating systems use

Transfer of Control

- Branch

- e.g. branch to x if result is zero

- Skip

- e.g. increment and skip if zero

- ISZ Register1

- Branch xxxx

- ADD A

- Subroutine call

- c.f. interrupt call

Foreground Reading

- Pentium and PowerPC operation types
- Stallings p338 et. Seq.

Byte Order

(A portion of chips?)

- What order do we read numbers that occupy more than one byte
- e.g. (numbers in hex to make it easy to read)
- 12345678 can be stored in 4x8bit locations as follows
-

Byte Order (example)

□ Address	Value (1)	Value(2)
□ 184	12	78
□ 185	34	56
□ 186	56	34
□ 186	78	12

□ i.e. read top down or bottom up?

Byte Order Names

- ❑ The problem is called Endian
- ❑ The system on the left has the least significant byte in the lowest address
- ❑ This is called big-endian
- ❑ The system on the right has the least significant byte in the highest address
- ❑ This is called little-endian

Standard...What Standard?

- Pentium (80x86), VAX are little-endian
- IBM 370, Motorola 680x0 (Mac), and most RISC are big-endian
- Internet is big-endian
 - Makes writing Internet programs on PC more awkward!
 - WinSock provides htoi and itoh (Host to Internet & Internet to Host) functions to convert